

Advanced Programming in the UNIX Environment

Week 02, Segment 2: open(2) and close(2)

**Department of Computer Science
Stevens Institute of Technology**

Jan Schaumann

jschauma@stevens.edu

<https://stevens.netmeister.org/631/>

Standard I/O

Basic File I/O: almost all UNIX file I/O can be performed using these five functions:

- `open(2)`
- `close(2)`
- `read(2)`
- `write(2)`
- `lseek(2)`

creat(2)

```
#include <fcntl.h>
```

```
int creat(const char *pathname, mode_t mode);
```

Returns: file descriptor if OK, -1 on error

`creat(2)` returns a file handle in write-only mode. To get a read-write file handle:

```
1 if ((fd = creat(path, mode) < 0 ) {  
2     /* error */  
3 }  
4 (void)close(fd);  
5 if ((fd = open(path, O_RDWR) < 0) {  
6     /* error */  
7 }  
8 /* do stuff with 'fd' ... */
```

creat(2)

```
#include <fcntl.h>
```

```
int creat(const char *pathname, mode_t mode);
```

Returns: file descriptor if OK, -1 on error

This interface is made obsolete by `open(2)`.

`creat()` is the same as:

```
open(path, O_CREAT | O_TRUNC | O_WRONLY, mode);
```

open(2)

```
#include <fcntl.h>
```

```
int open(const char *pathname, int oflag, ... /* mode_t mode */);
```

Returns: file descriptor if OK, -1 on error

oflag must be one (and only one) of:

- O_RDONLY - open for reading only
- O_WRONLY - open for writing only
- O_RDWR - open for reading and writing

and may be OR'd with any of these:

- O_APPEND – append on each write
- O_CREAT – create the file if it doesn't exist; requires *mode* argument
- O_EXCL – error if O_CREAT and file already exists. (atomic)
- O_TRUNC – truncate size to 0
- O_NONBLOCK – do not block on open or for data to become available
- O_SYNC – wait for physical I/O to complete

open(2)

```
#include <fcntl.h>
```

```
int open(const char *pathname, int oflag, ... /* mode_t mode */);
```

Returns: file descriptor if OK, -1 on error

Additional *oflags* may be supported on some platforms:

- O_DIRECTORY – if path resolves to a non-directory file, fail and set errno to ENOTDIR
- O_DSYNC – wait for physical I/O for data, except file attributes
- O_EXEC – open file for execute only, fail if it is a directory
- O_NOFOLLOW – do not follow symlinks
- O_PATH – obtain a file descriptor purely for fd-level operations. (Linux >2.6.36 only)
- O_RSYNC – block read operations on any pending writes
- O_SEARCH – open for search only, fail if it is a regular file
- ...

openat(2)

```
#include <fcntl.h>
```

```
int open(const char *pathname, int oflag, ... /* mode_t mode */);
```

```
int openat(int dirfd, const char *pathname, int oflag, ... /* mode_t mode */);
```

Returns: file descriptor if OK, -1 on error

`openat(2)` is used to handle relative pathnames from different working directories in an atomic fashion.

Here, *pathname* is determined relative to the directory associated with the file descriptor *fd* instead of the current working directory.

open(2)

```
#include <fcntl.h>
```

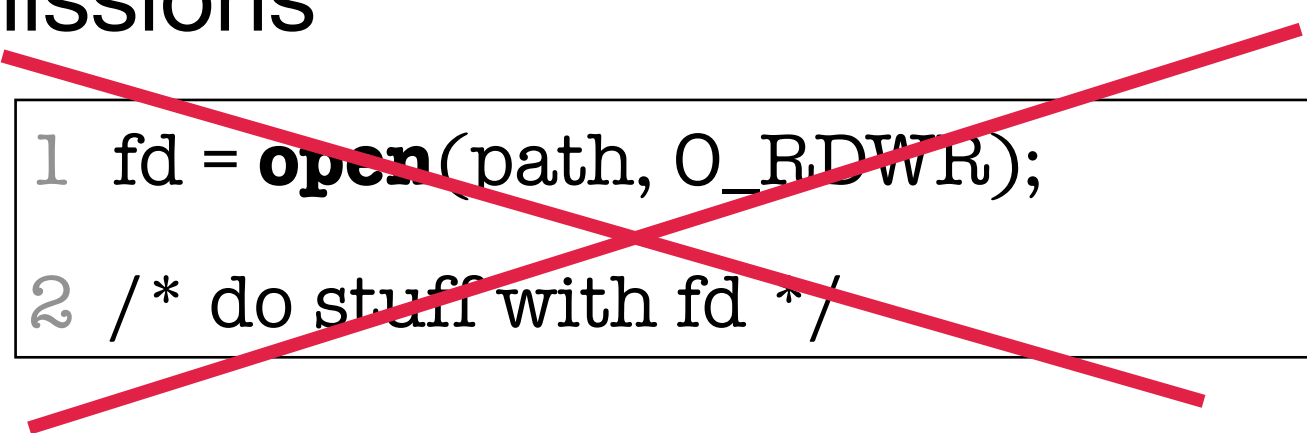
```
int open(const char *pathname, int oflag, ... /* mode_t mode */);
```

Returns: file descriptor if OK, -1 on error

open(2) may fail for a surprising number of reasons. Some of the more common ones include:

- **EEXIST**: **O_CREAT** | **O_EXCL** was specified, but the file exists
- **EMFILE**: process has already reached max number of open file descriptors
- **ENOENT**: file does not exist
- **EPERM**: lack of permissions
- ...

```
1 fd = open(path, O_RDWR);
2 /* do stuff with fd */
```



```
1 if ((fd = open(path, O_RDWR) < 0) {
2     /* error */
3 }
4 /* do stuff with fd */
```


close(2)

```
#include <unistd.h>

int close(int fd);
```

Returns: 0 if OK, -1 on error

- closing a filedescriptor releases any record locks on that file (more on that in future lectures)
- file descriptors not explicitly closed are closed by the kernel when the process terminates.
- to avoid leaking file descriptors, always close(2) them within the same scope

```
#include <fcntl.h>
#include <stdio.h>
#include <string.h>
```

```
int
main() {
    /* imagine a few dozen lines of code here */
}
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
"f.c" 8L, 120C written
```

```
6,8
```

```
71
```

```
75%
```

```
B
```

close(2)

```
#include <unistd.h>

int close(int fd);
```

Returns: 0 if OK, -1 on error

- closing a filedescriptor releases any record locks on that file (more on that in future lectures)
- file descriptors not explicitly closed are closed by the kernel when the process terminates.
- to avoid leaking file descriptors, always `close(2)` them within the same scope

```
1 (void)close(fd);
2 /* you can't do stuff with fd here either way */
```

[\$ ssh apue]

Last login: Wed Sep 2 01:22:10 2020 from 10.0.2.2

NetBSD 9.0 (GENERIC) #0: Fri Feb 14 00:06:28 UTC 2020

Welcome to NetBSD!

apue\$ █



In our next segment...

- `read(2)`
- `write(2)`
- `lseek(2)`

Can you go backwards on a pipe?

What happens when you try to write data way beyond the end of a file?

How efficient is our `simple-cat.c` program?