

# **Advanced Programming in the UNIX Environment**

## **Week 04, Segment 1: The Unix Filesystem**

**Department of Computer Science  
Stevens Institute of Technology**

**Jan Schaumann**

`jschauma@stevens.edu`

`https://stevens.netmeister.org/631/`

hard disk



- a disk can be divided into logical *partitions*

```
[apue$ sudo disklabel wd0]
```

```
# /dev/rwd0:
```

```
type: unknown
```

```
disk: wd
```

```
label: fictitious
```

```
flags:
```

```
bytes/sector: 512
```

```
sectors/track: 63
```

```
tracks/cylinder: 16
```

```
sectors/cylinder: 1008
```

```
cylinders: 33288
```

```
total sectors: 33554432
```

```
rpm: 3600
```

```
interleave: 1
```

```
trackskew: 0
```

```
cylinderskew: 0
```

```
headswitch: 0 # microseconds
```

```
track-to-track seek: 0 # microseconds
```

```
drivedata: 0
```

```
4 partitions:
```

#	size	offset	fstype	[fsize	bsize	cpb/sgs]
a:	31457216	64	4.2BSD	0	0	0 #
b:	2097088	31457344	swap			#
c:	33554368	64	unused	0	0	#
d:	33554432	0	unused	0	0	#

```
apue$
```

physical blocksize

first NetBSD partition  
second NetBSD partition  
NetBSD portion of disk  
entire disk

hard disk

partition

partition

partition

file system

boot  
blocks

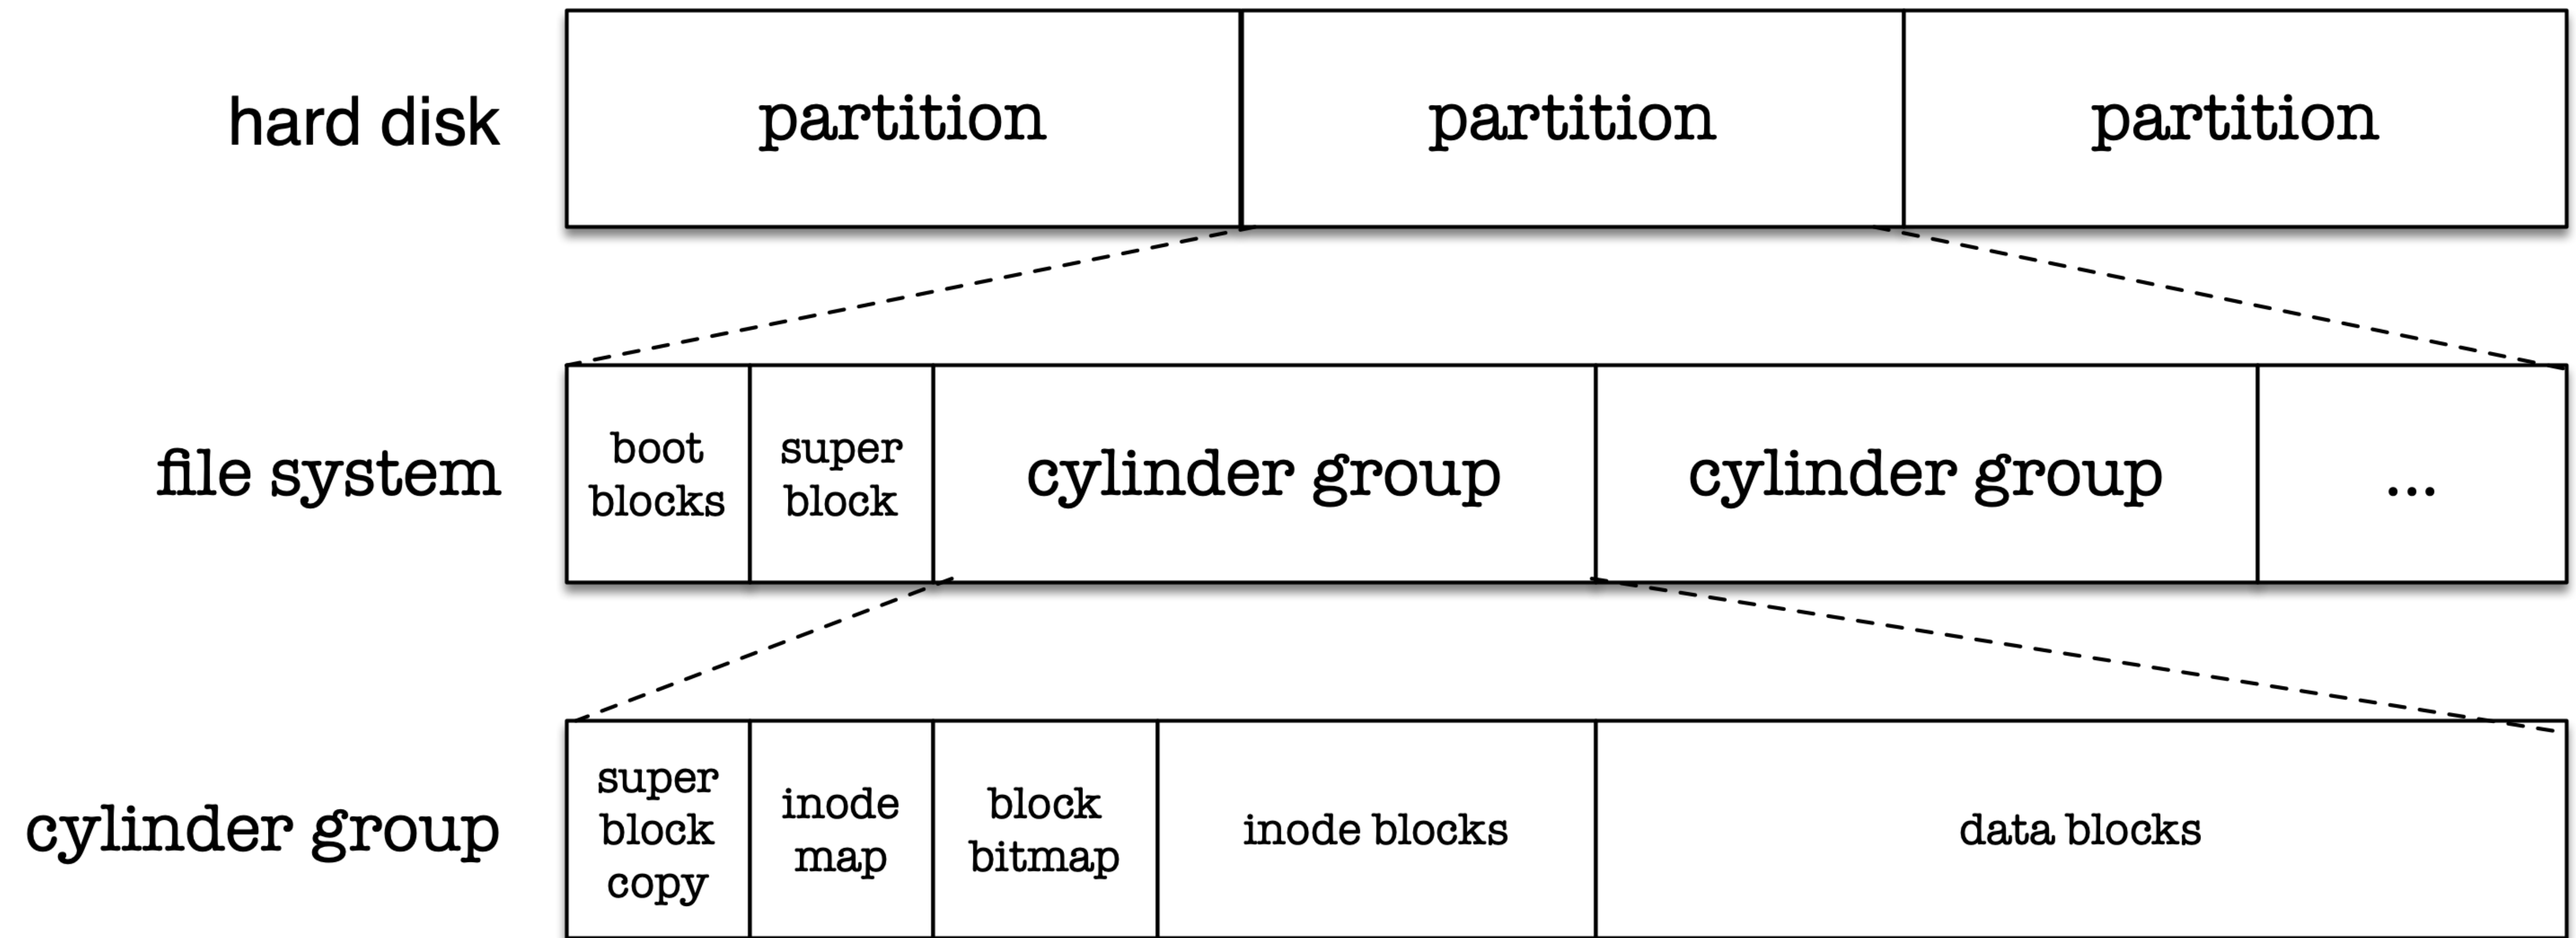
super  
block

cylinder group

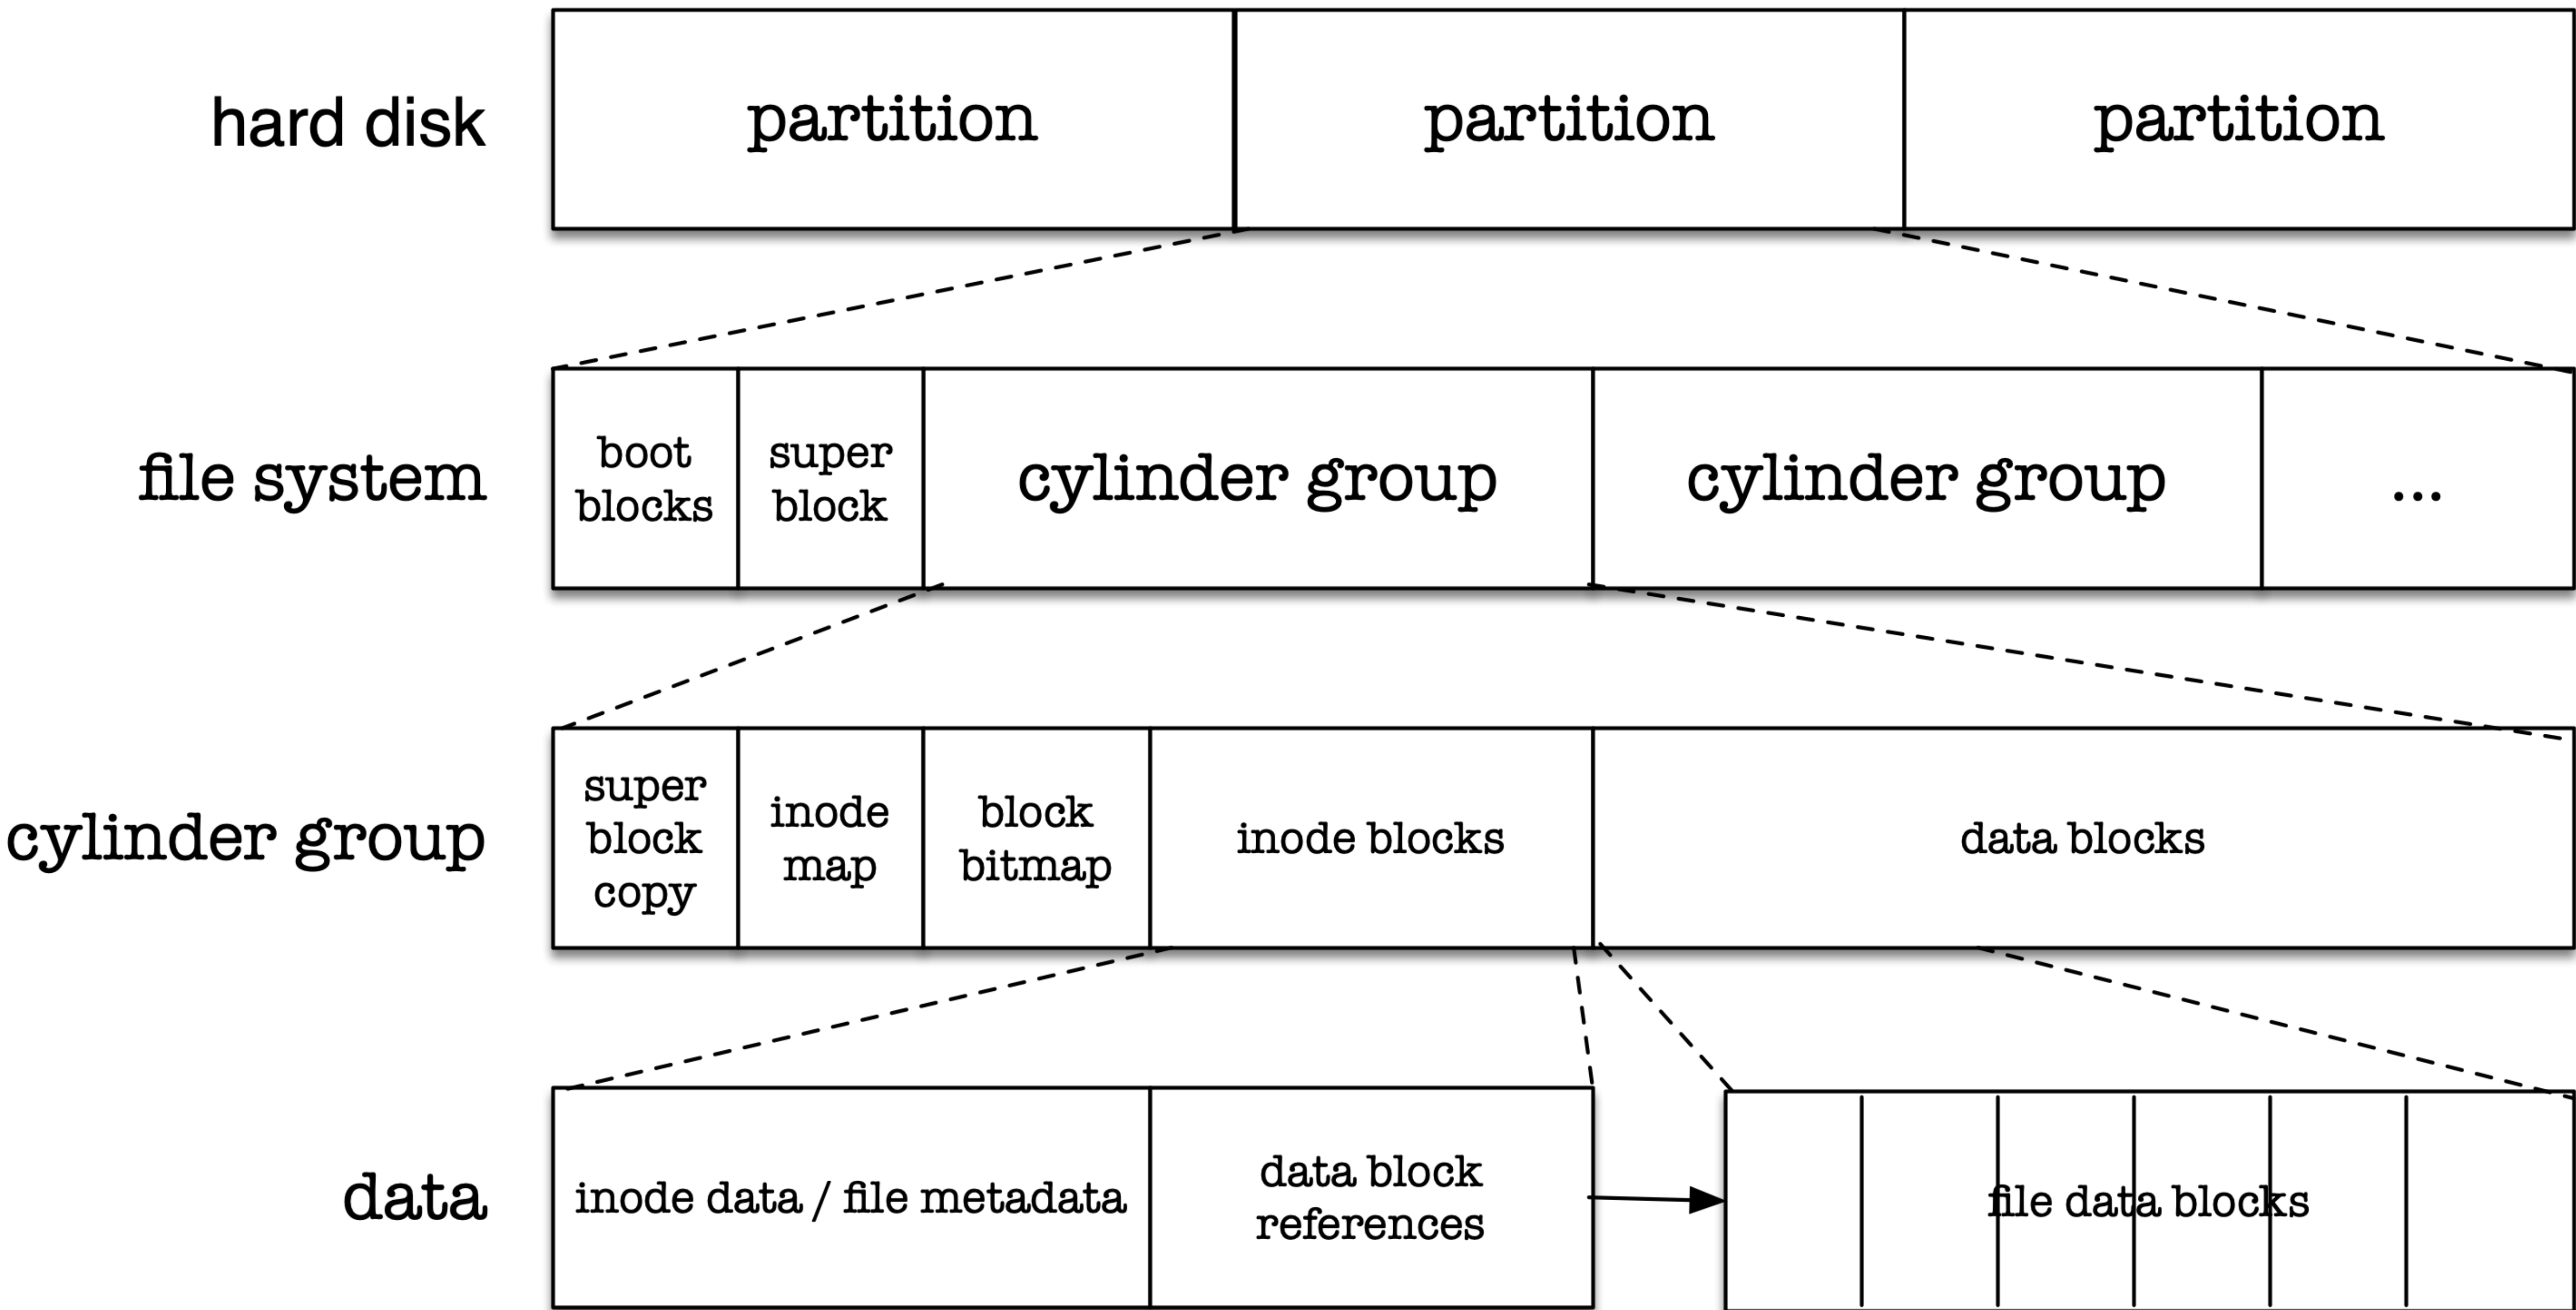
cylinder group

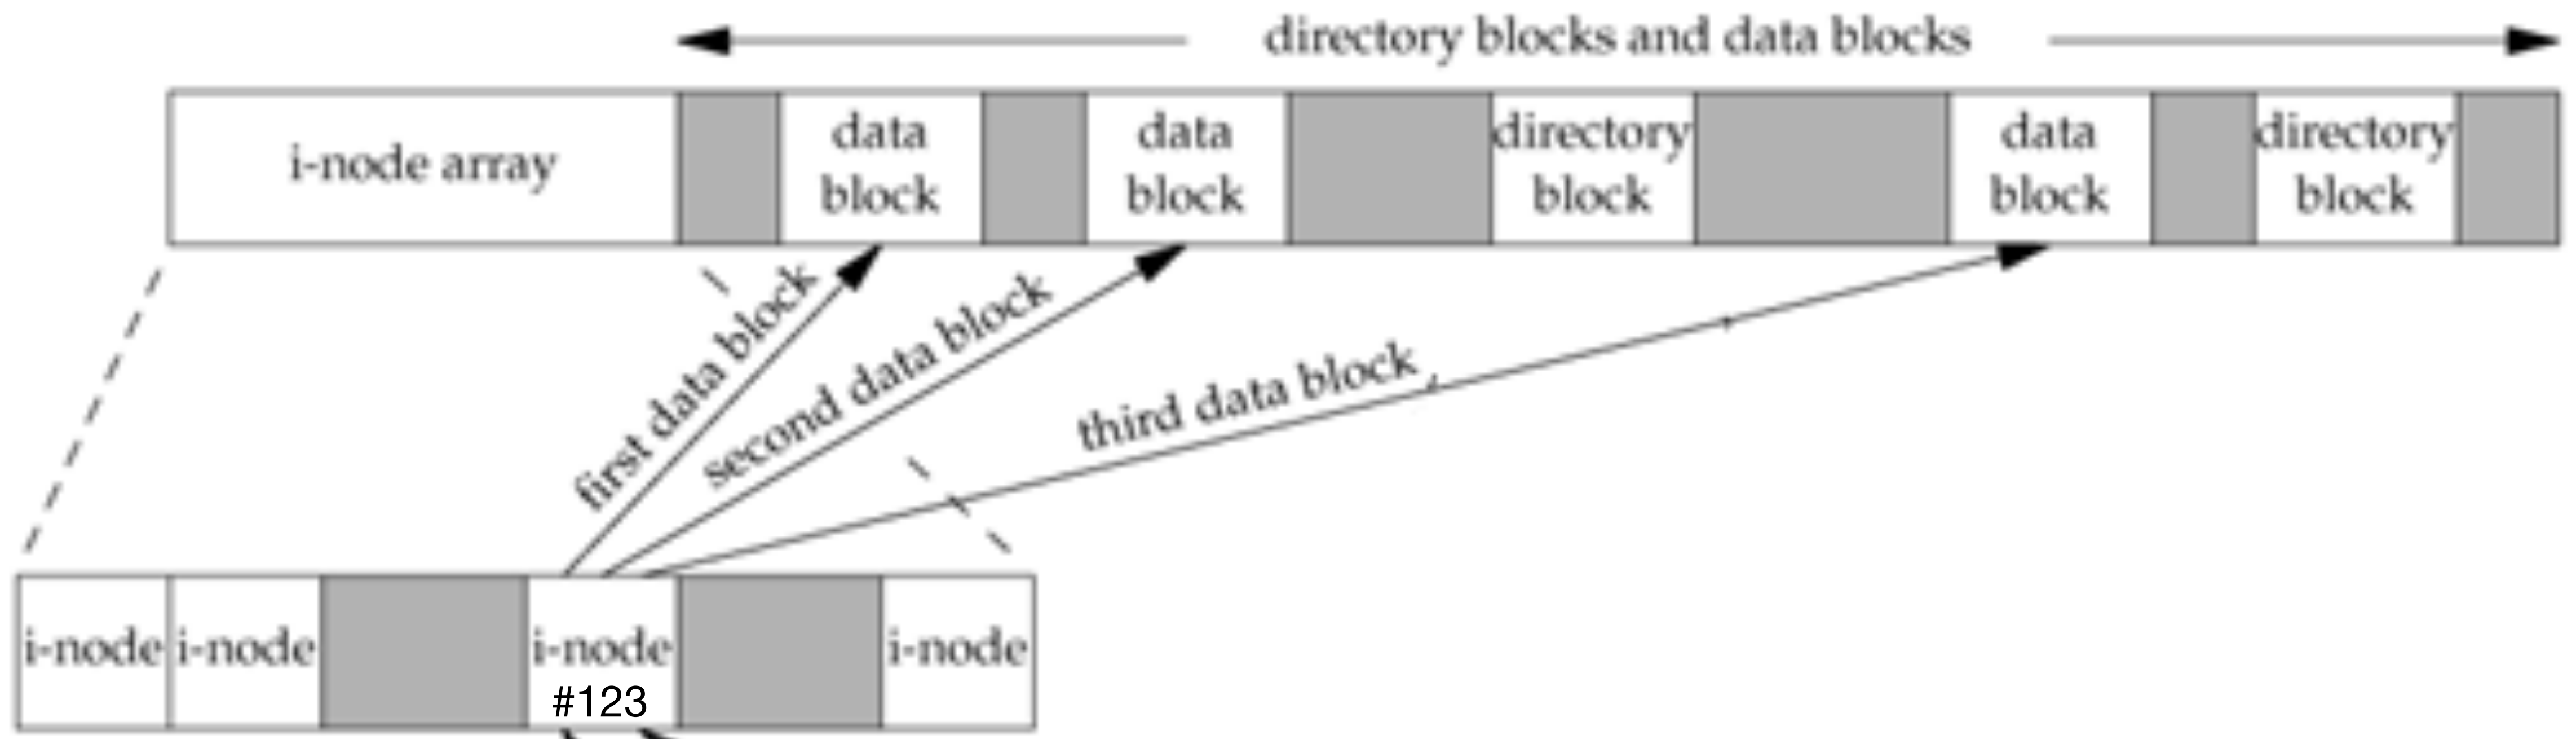
...

- on each logical *partition* you may create a *file system* containing the cylinder groups



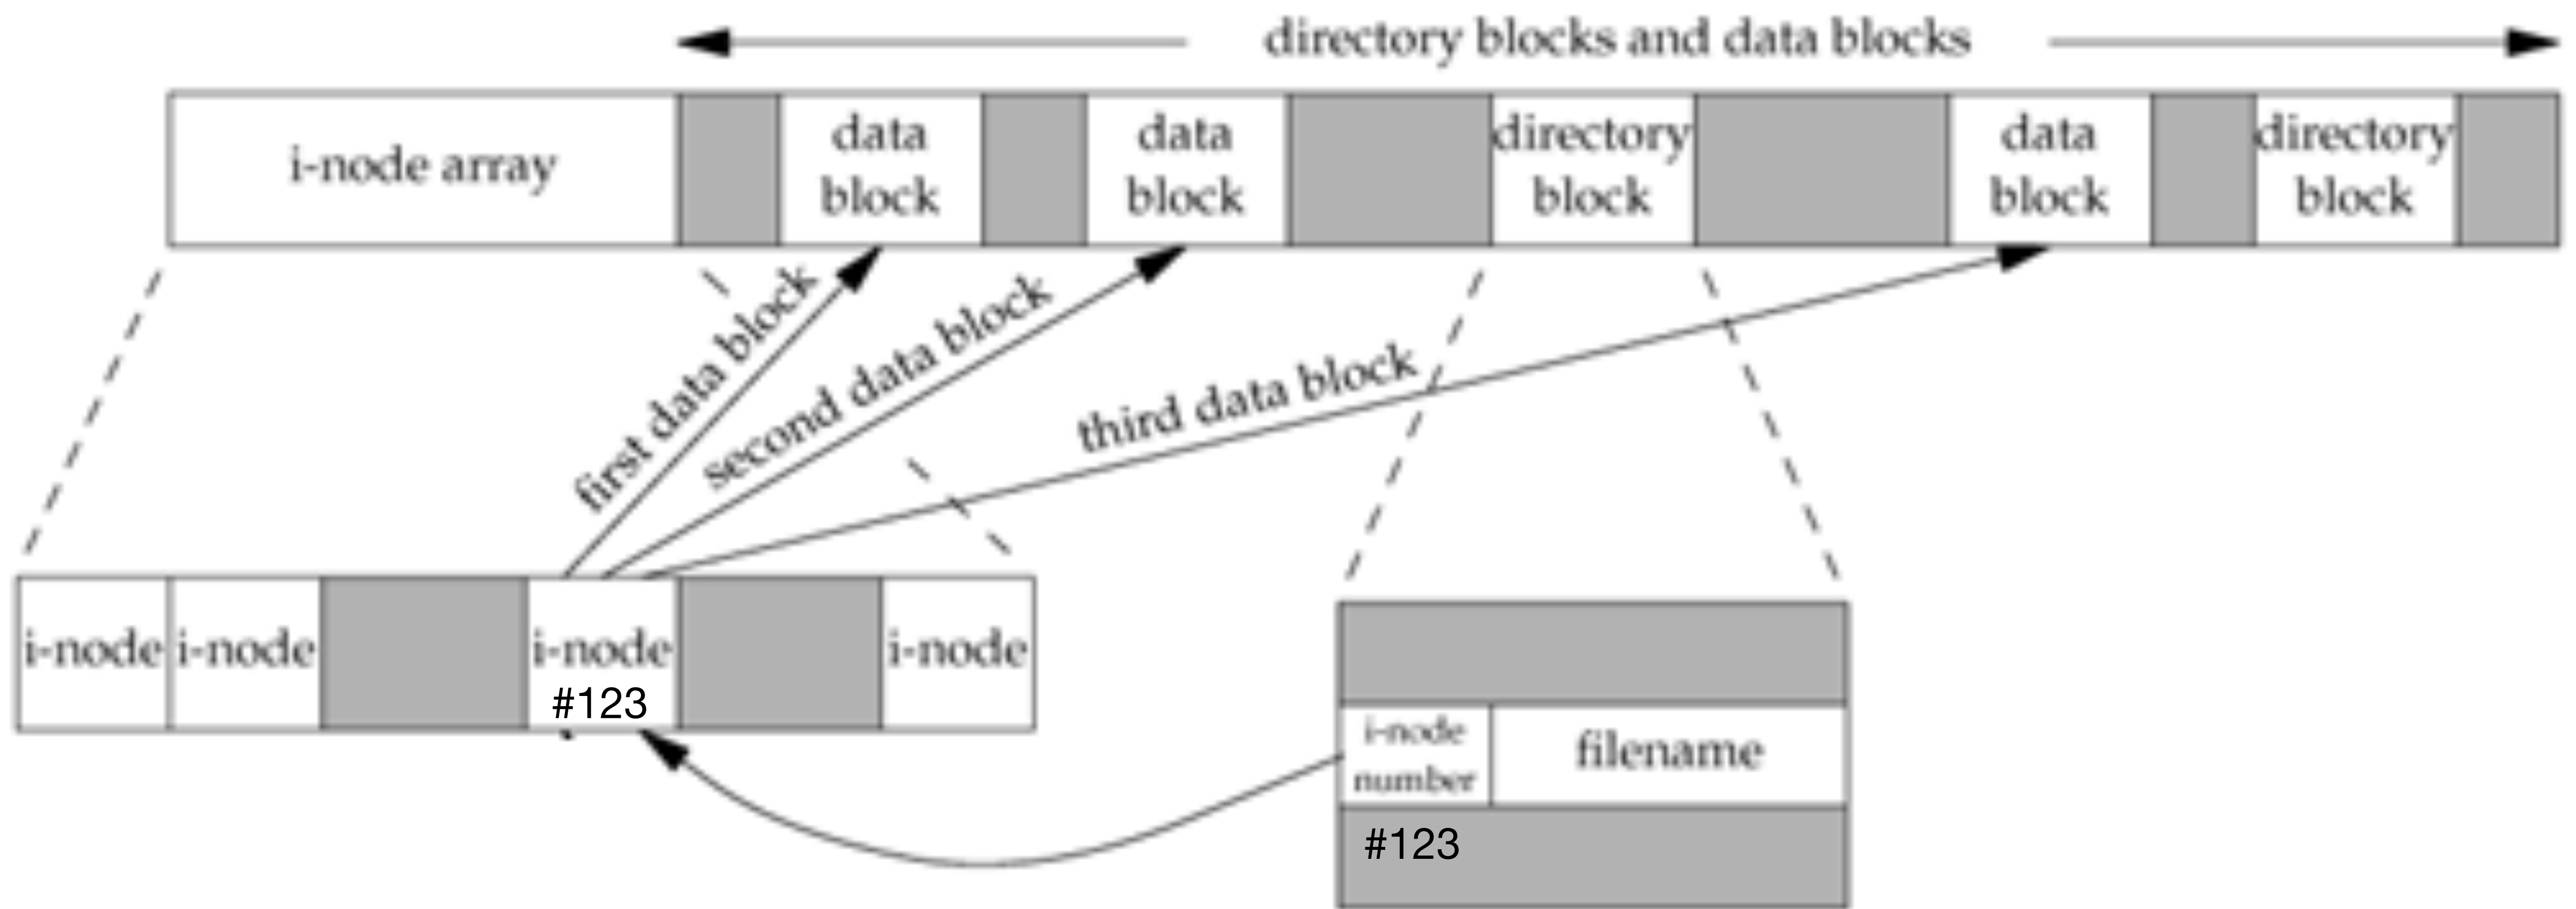
- each *cylinder group* contains a list of *inodes* (i-list) as well as the actual *directory*- and *data* blocks





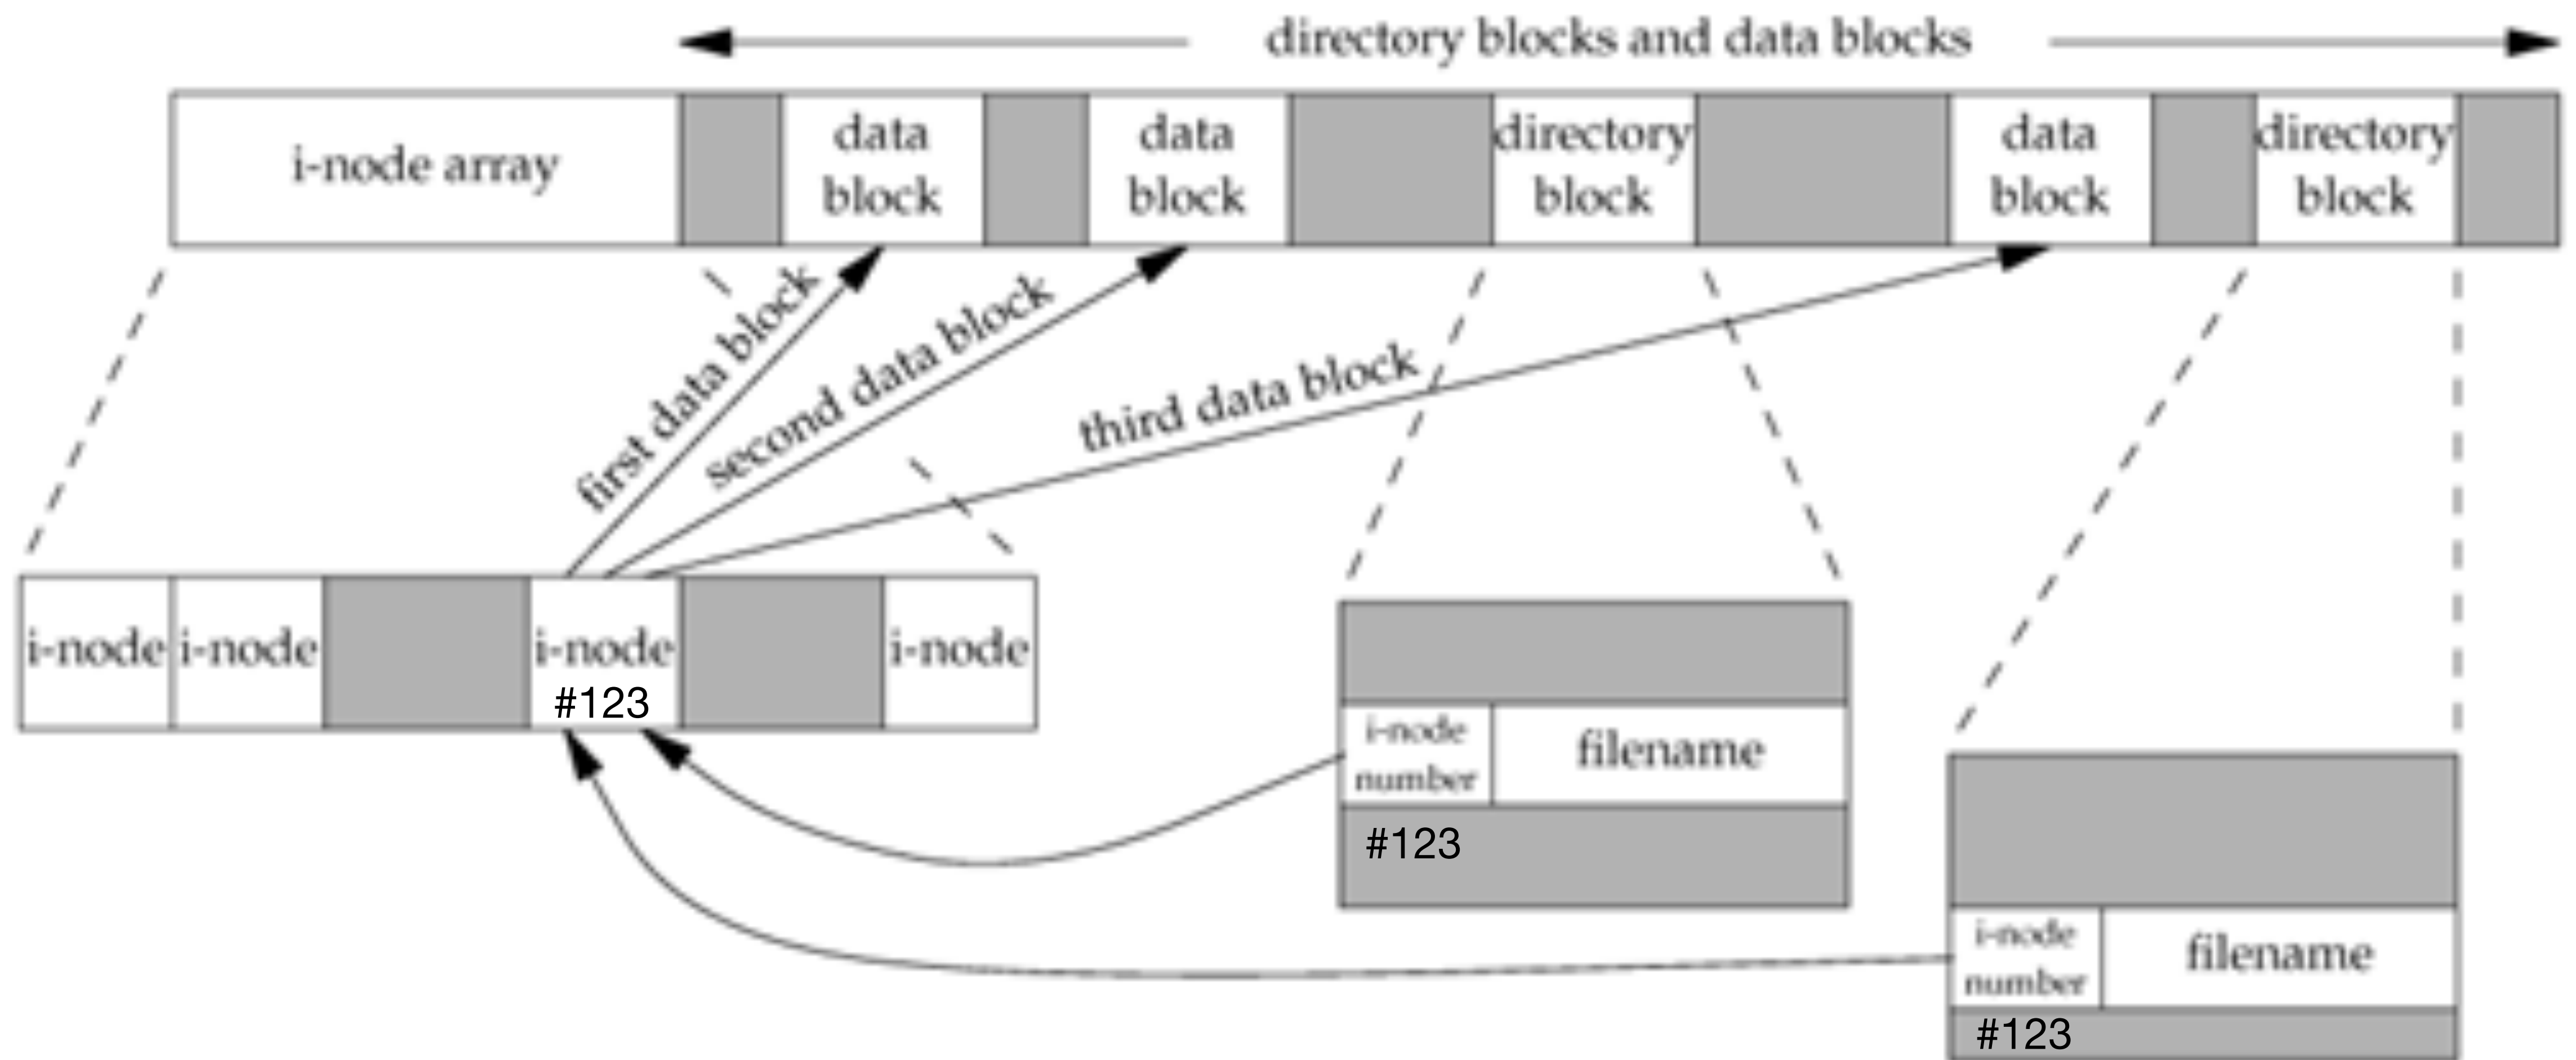
- data blocks containing the actual data (*i.e.*, contents of the file) are referenced from the inode



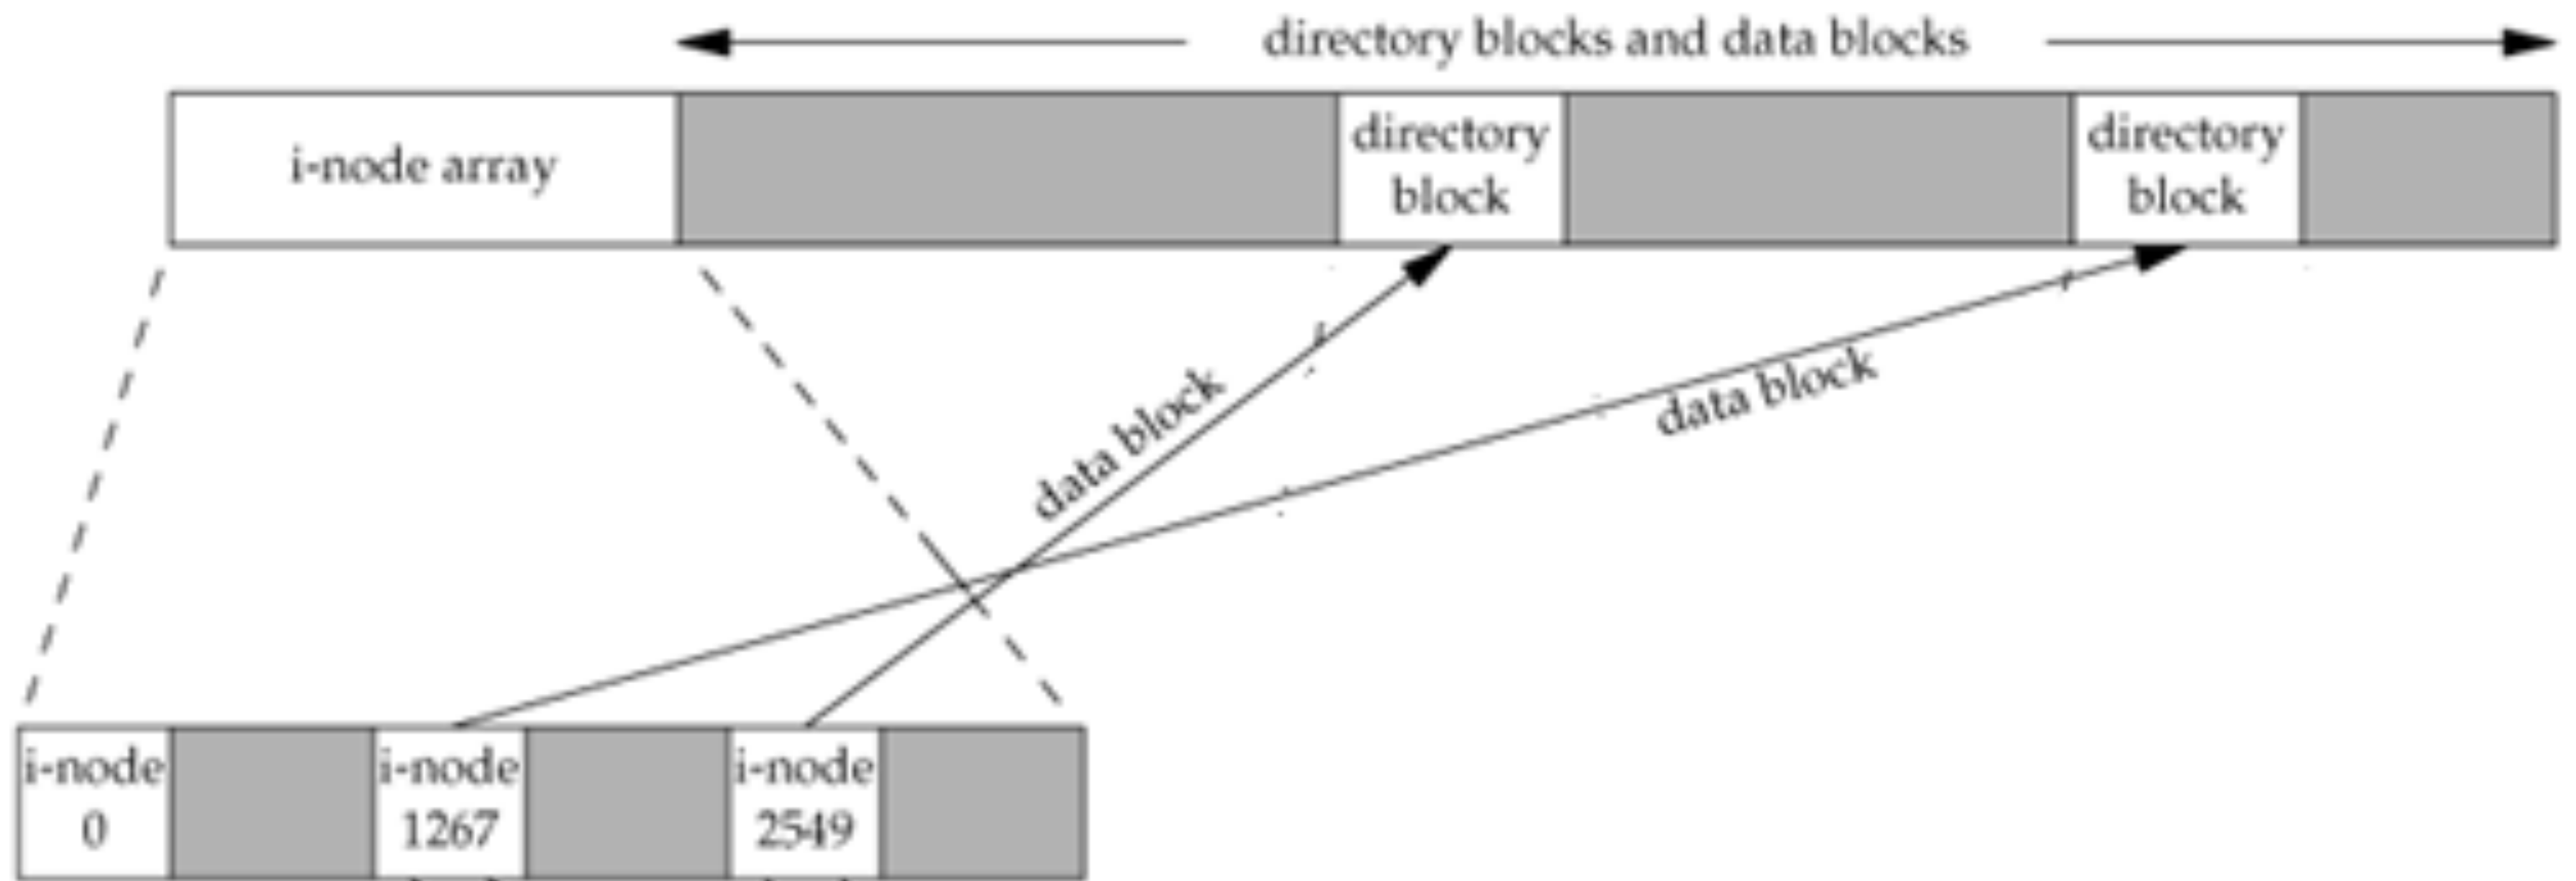


- a directory entry is really just a *hard link* mapping a “filename” to an inode

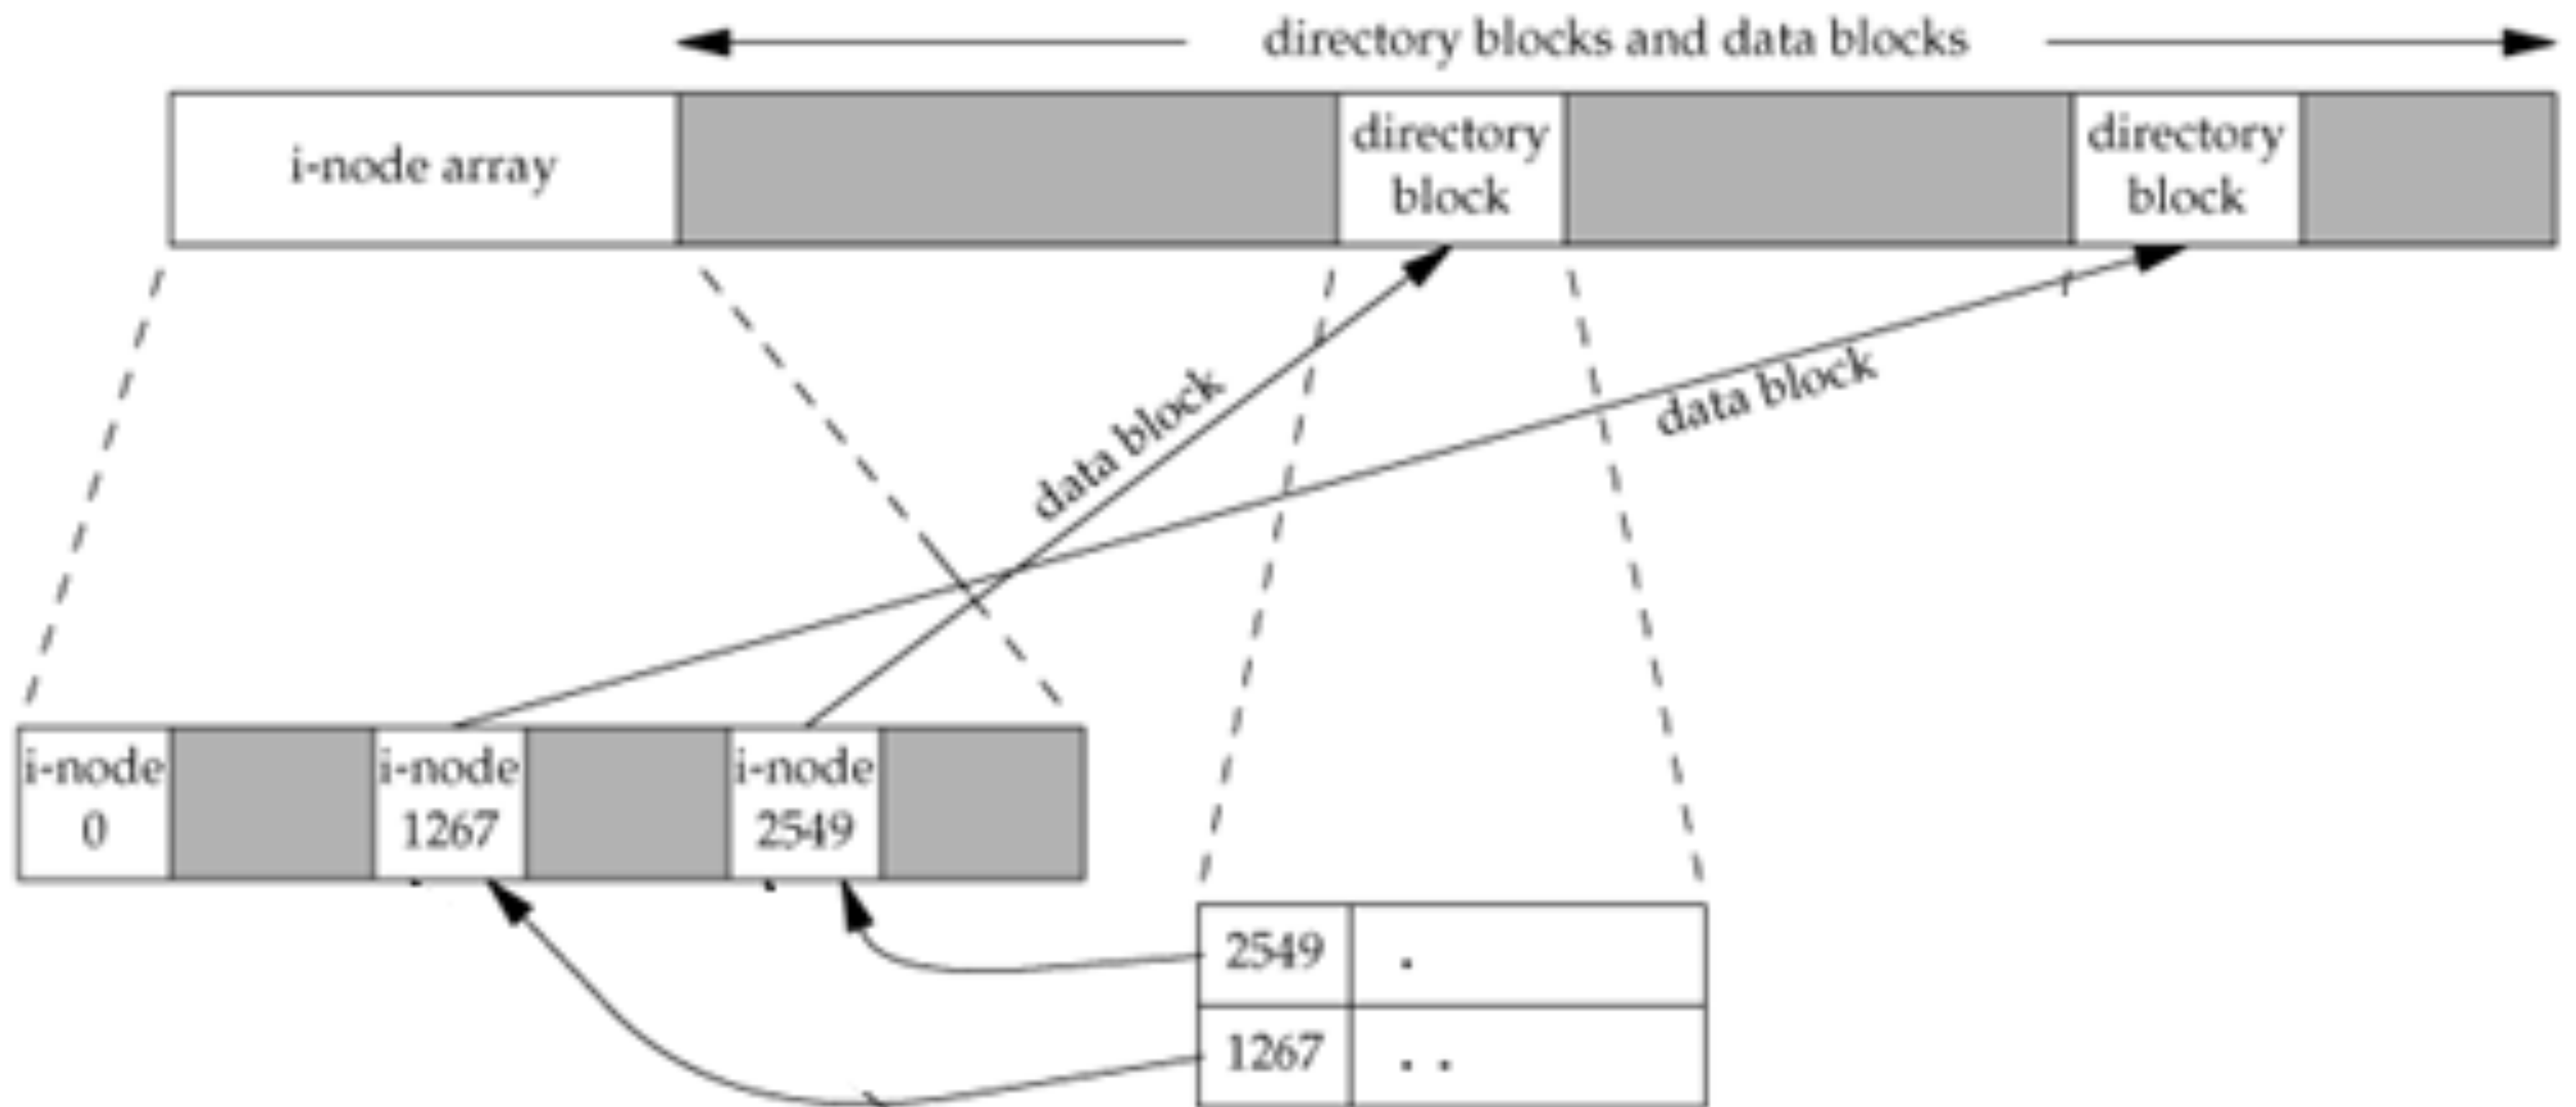




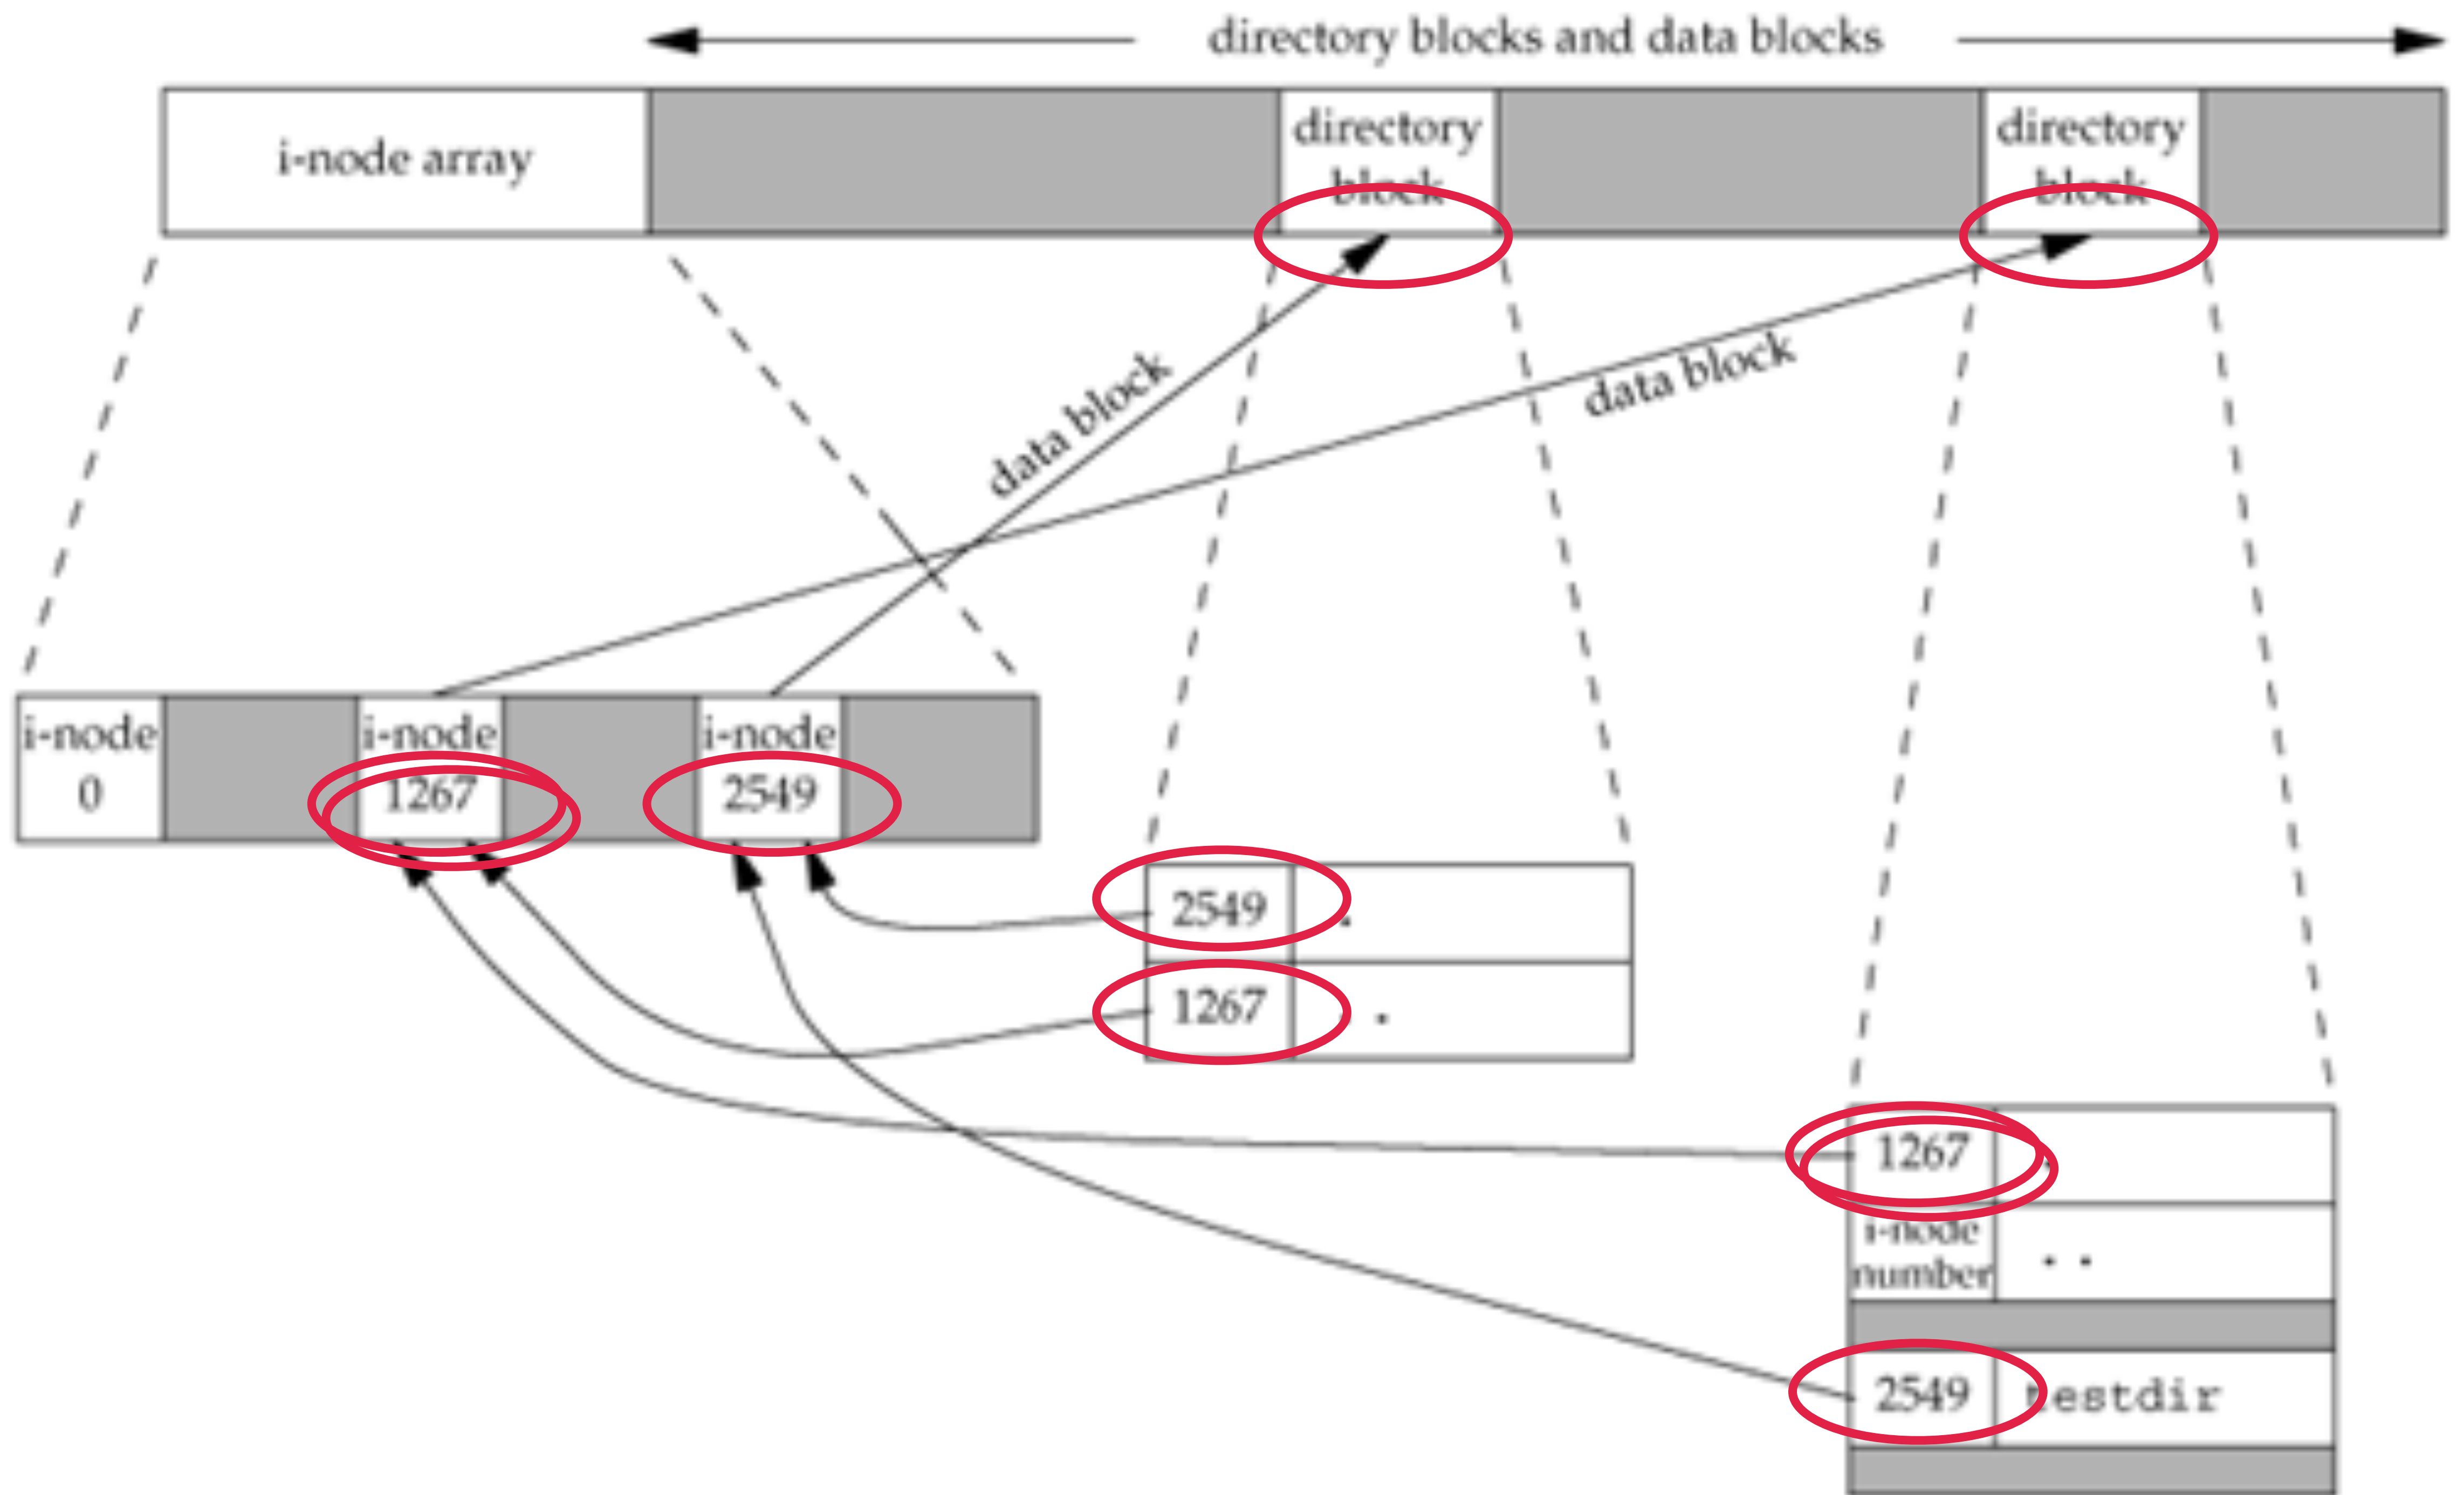
- a directory entry is really just a *hard link* mapping a “filename” to an inode
- you can have many such mappings to the same inode



- directories are special "files" containing a list of hard links

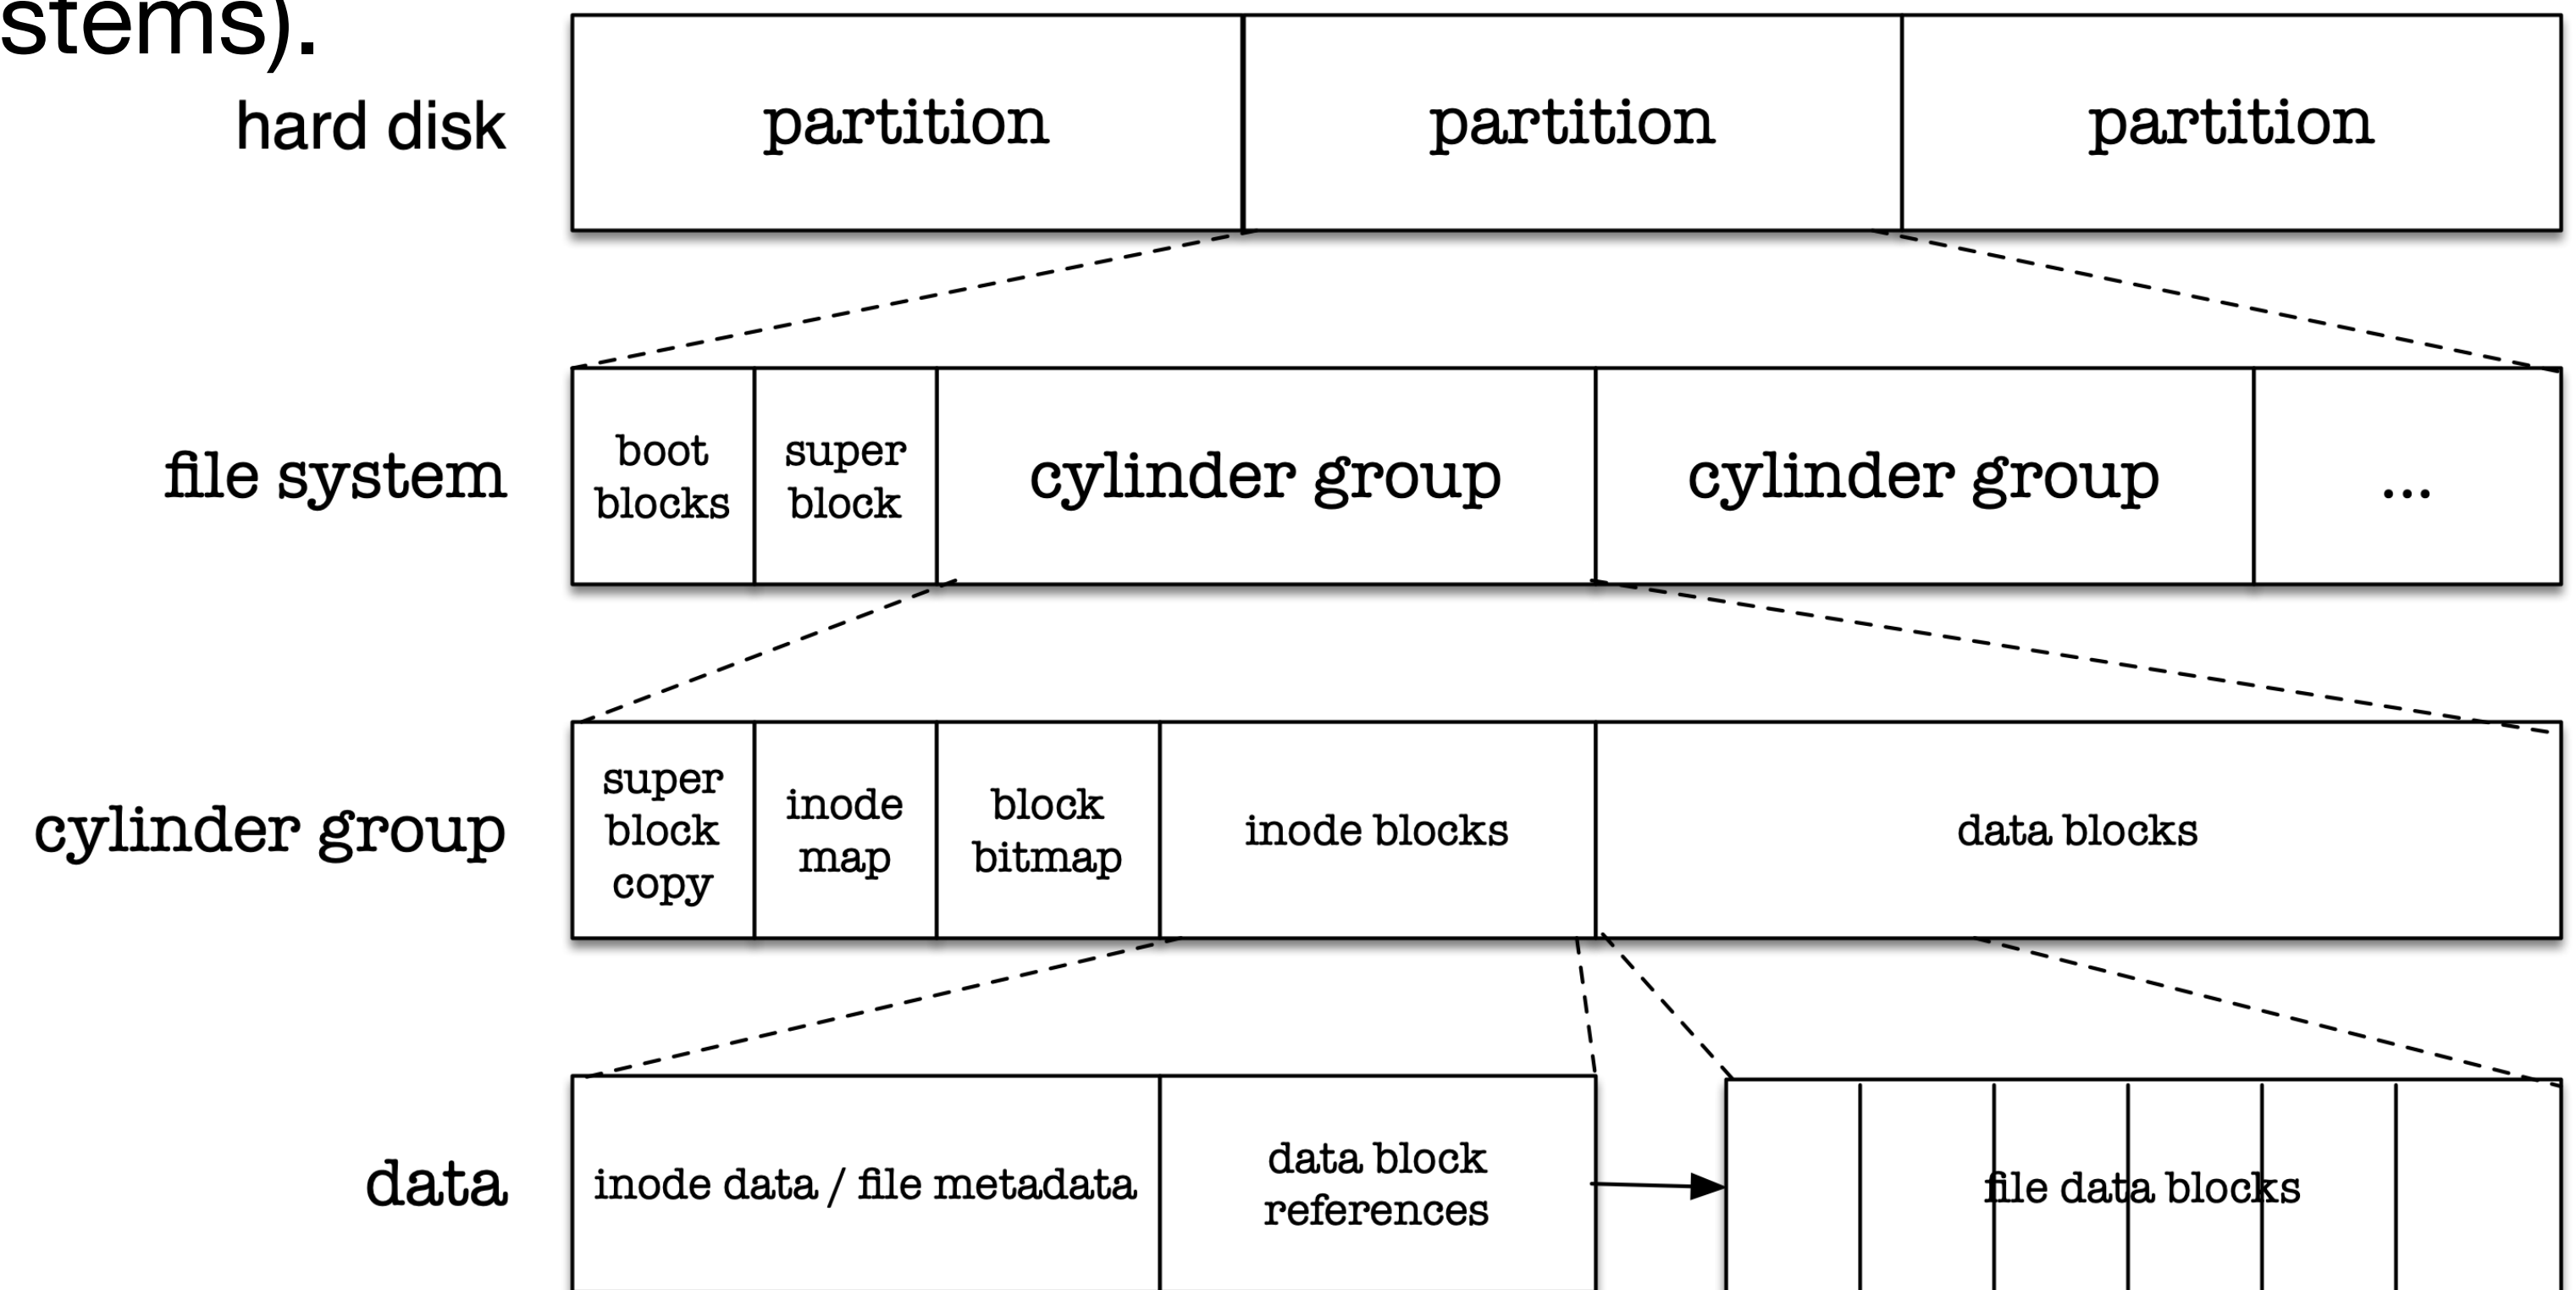


- each directory contains at least two entries:
  - "." -- this directory
  - ".." -- the parent directory



# Inodes

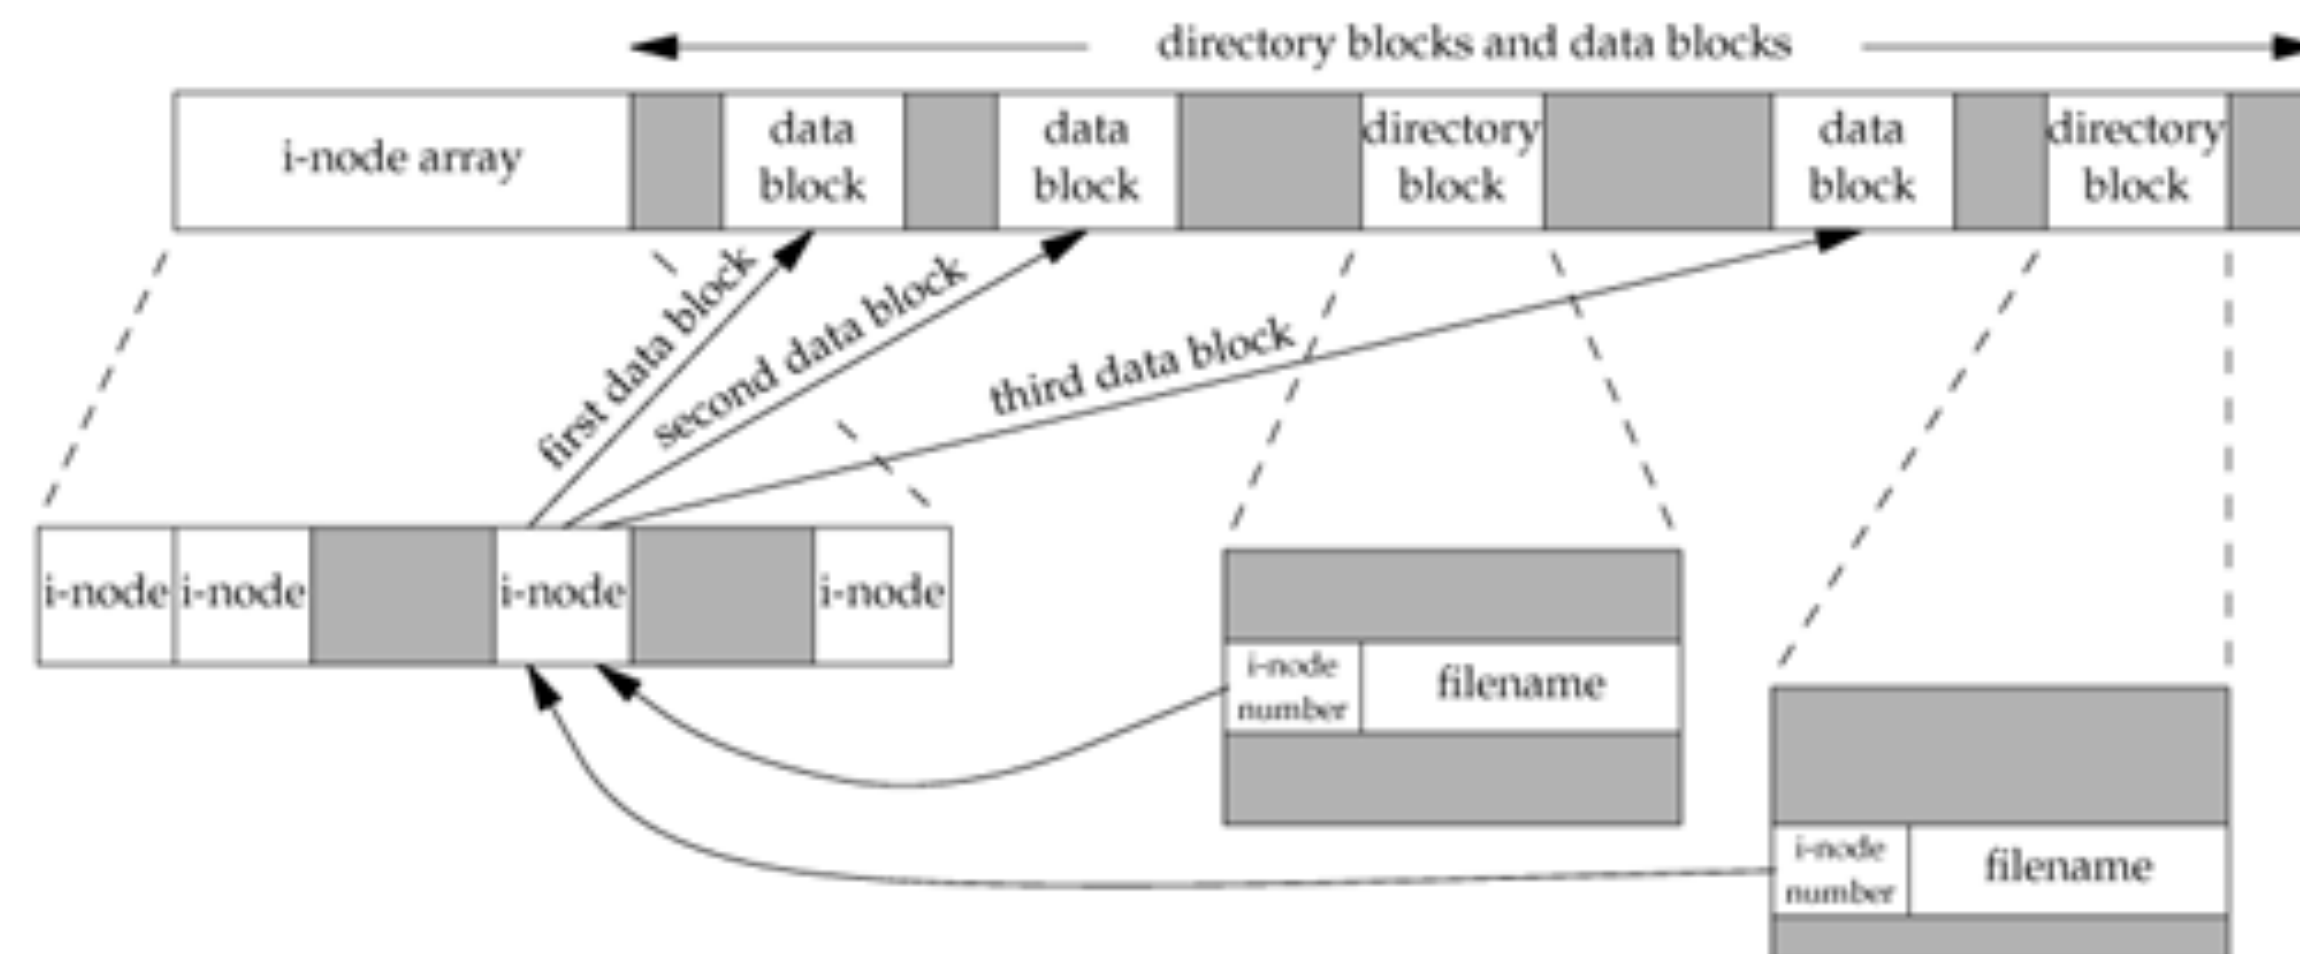
- The inode number in a directory entry must point to an inode on the same file system (no hardlinks across filesystems).





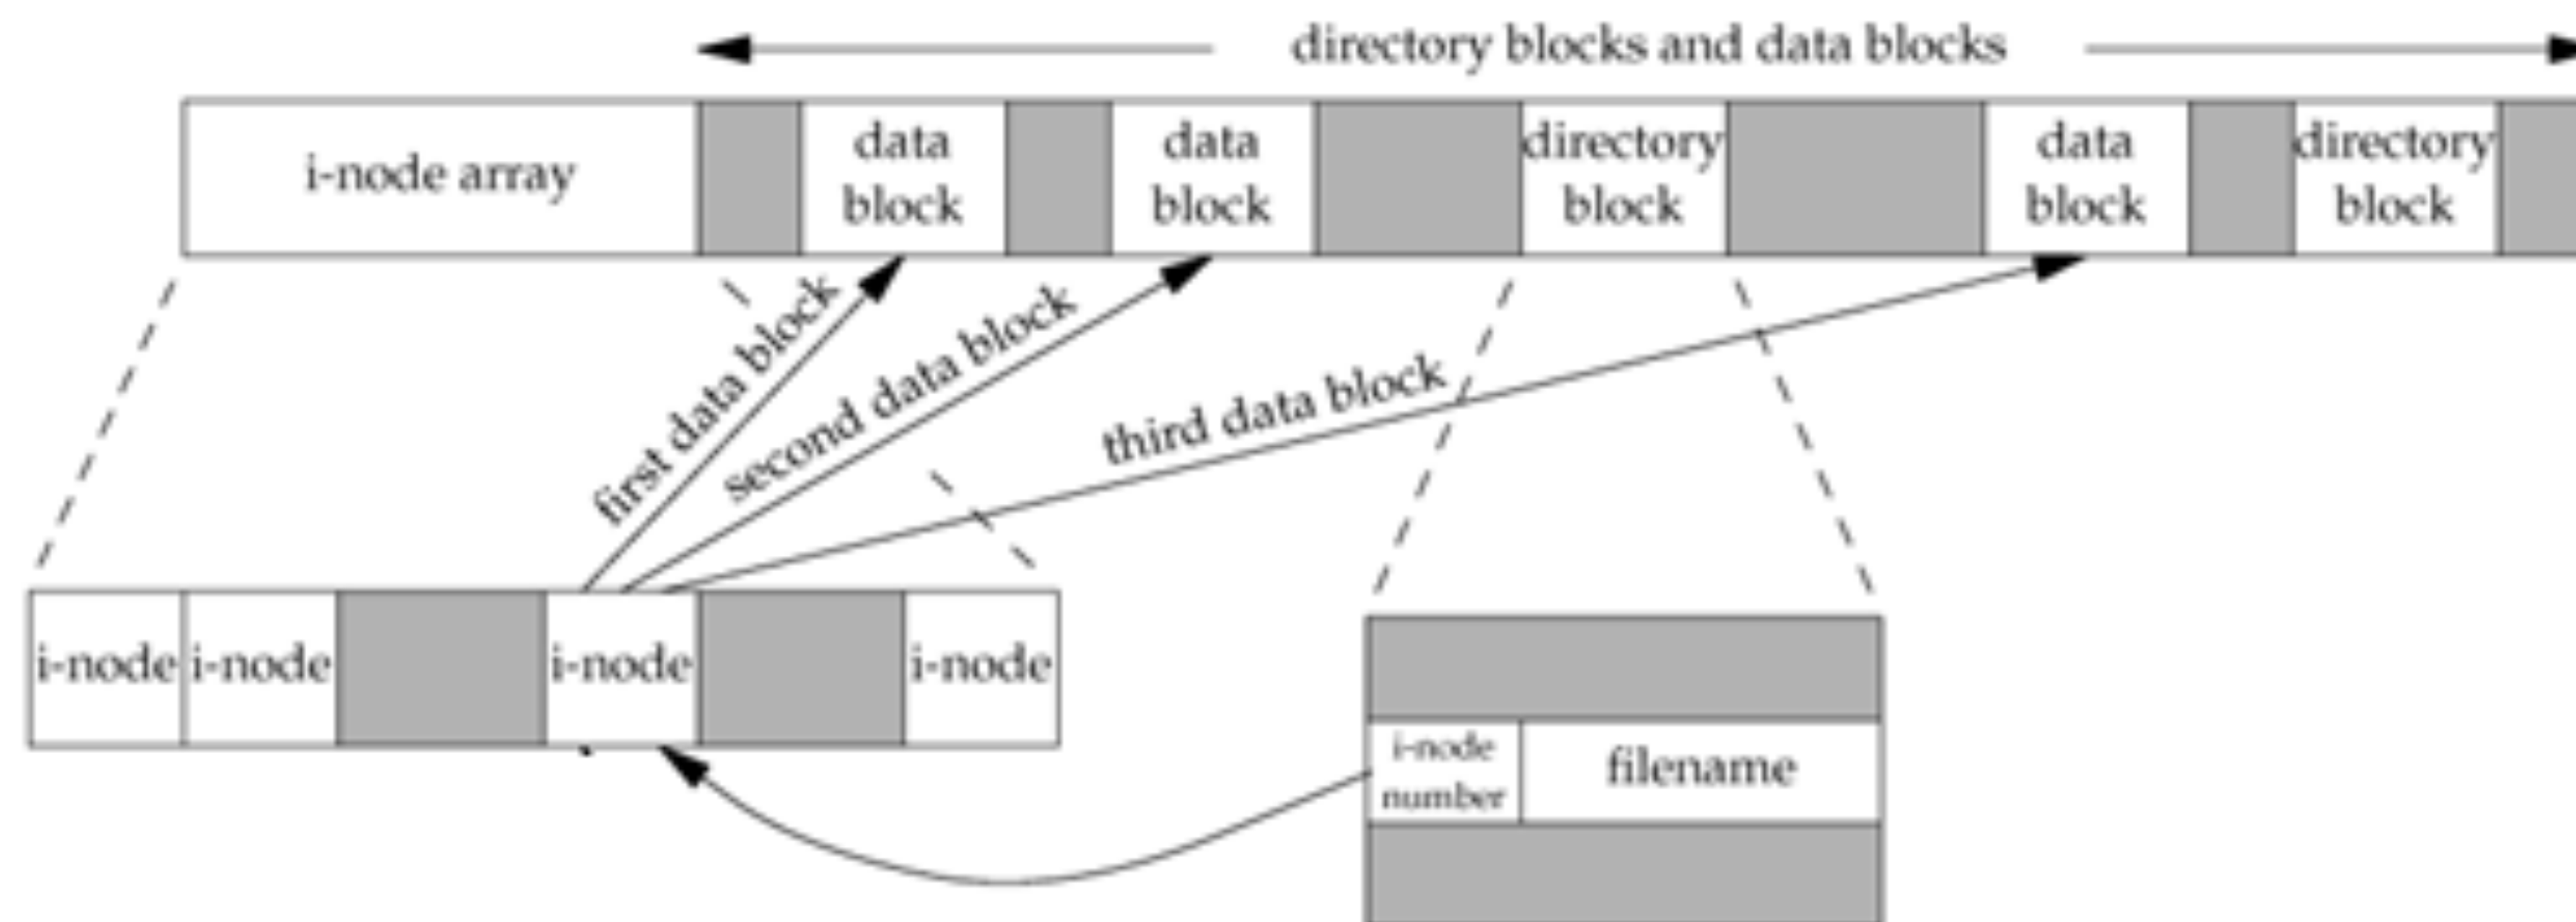
# Inodes

- The inode number in a directory entry must point to an inode on the same file system (no hardlinks across filesystems).
- The inode contains most of the information found in the struct stat.
- Every inode has a link count (st\_nlink): it shows how many “things” point to this inode. Only if this link count is 0 (and no process has the file open) are the data blocks freed.



# Inodes

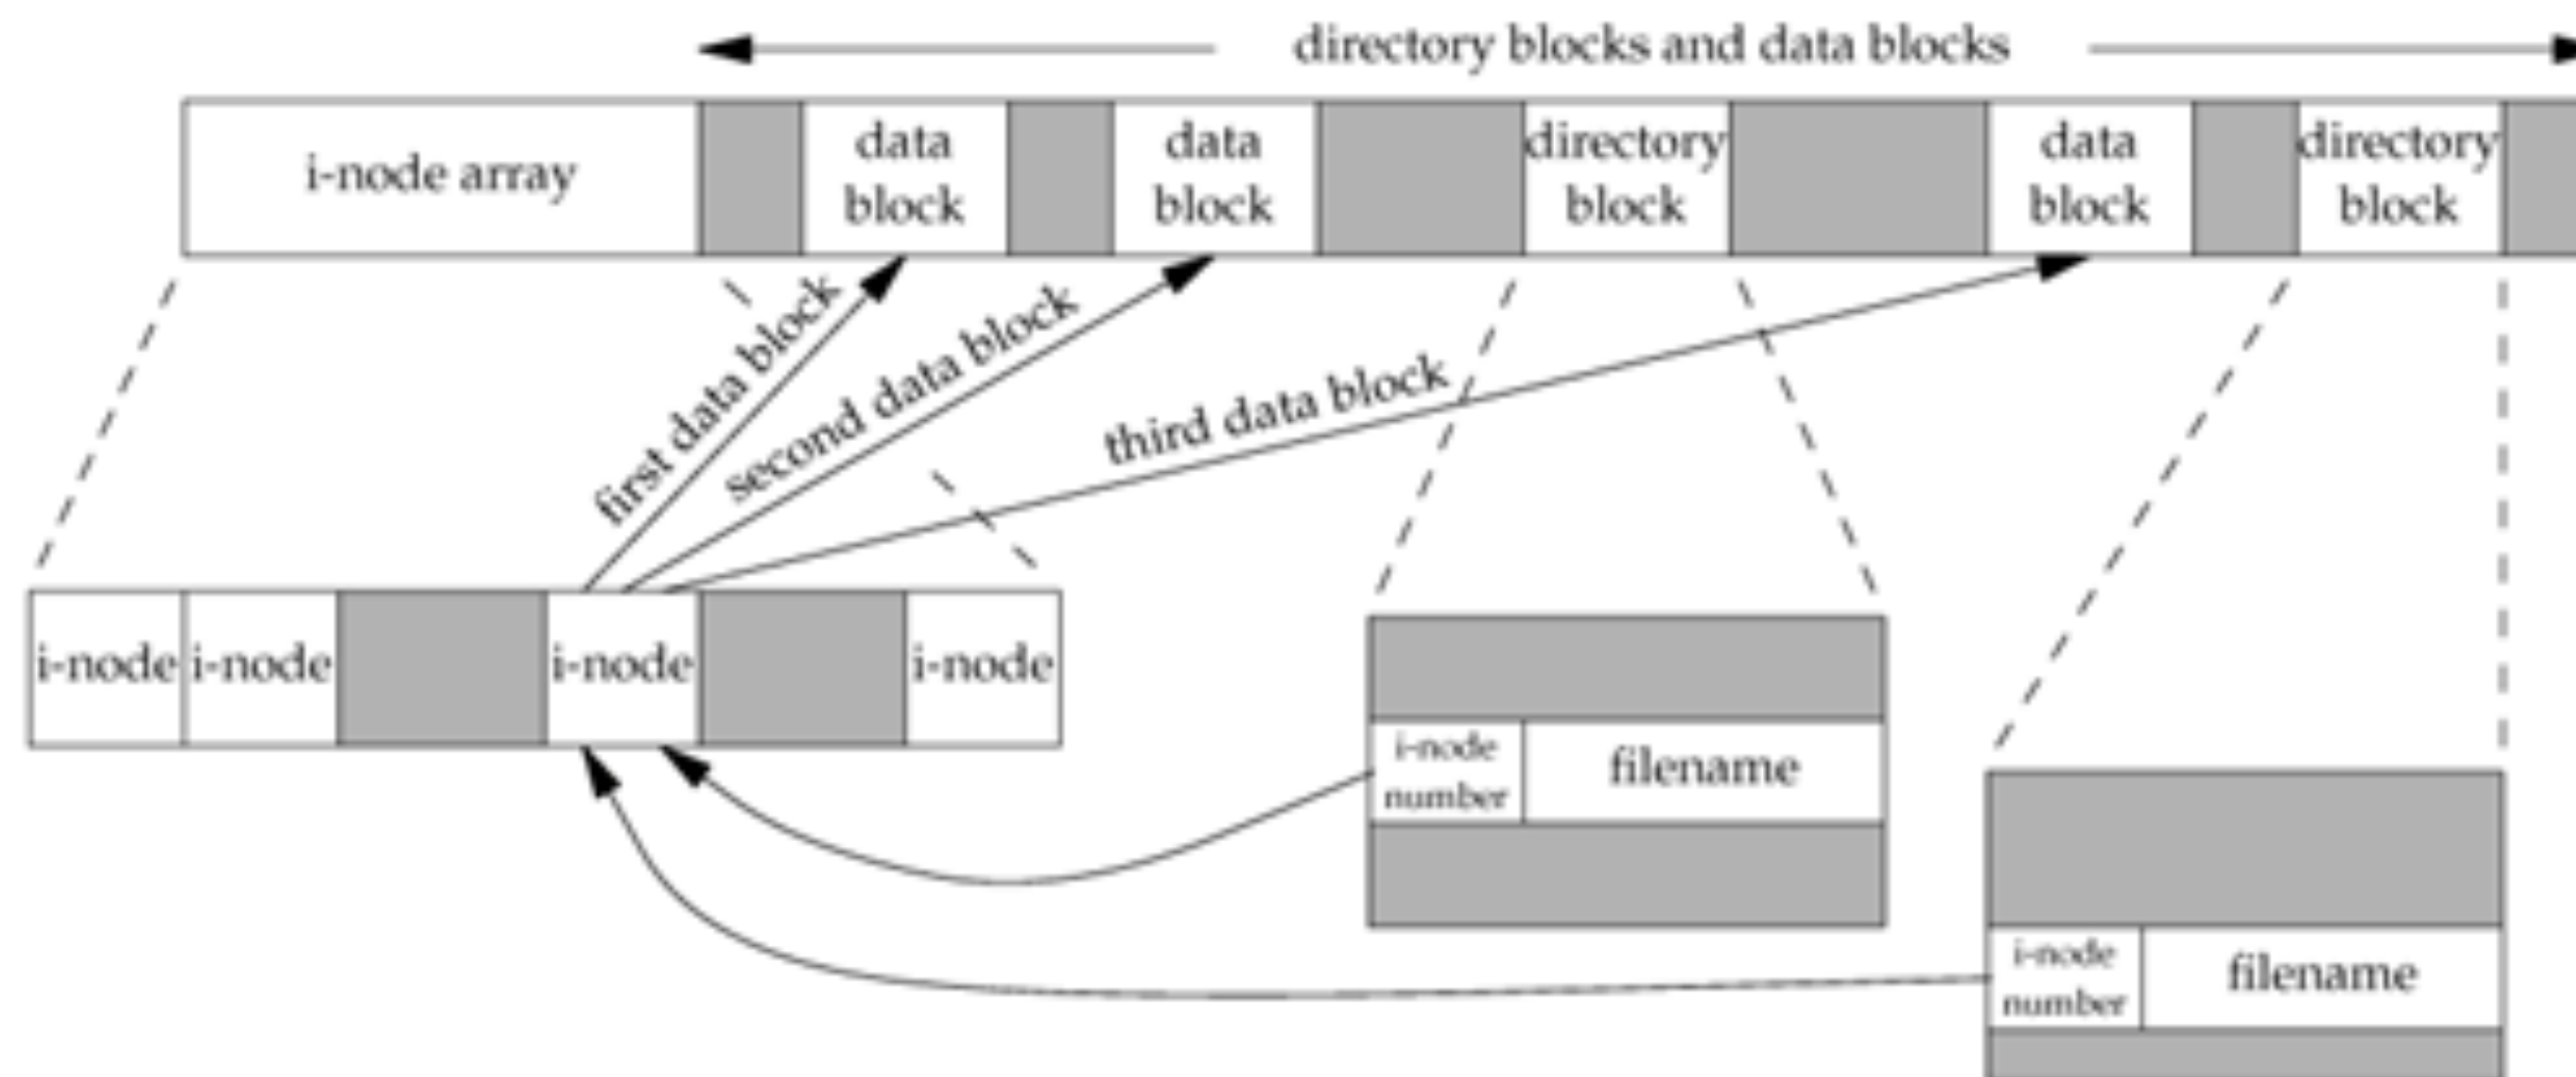
- To move a file within a single filesystem, we can just "move" the directory entry (actually done by creating a new entry, and deleting the old one).





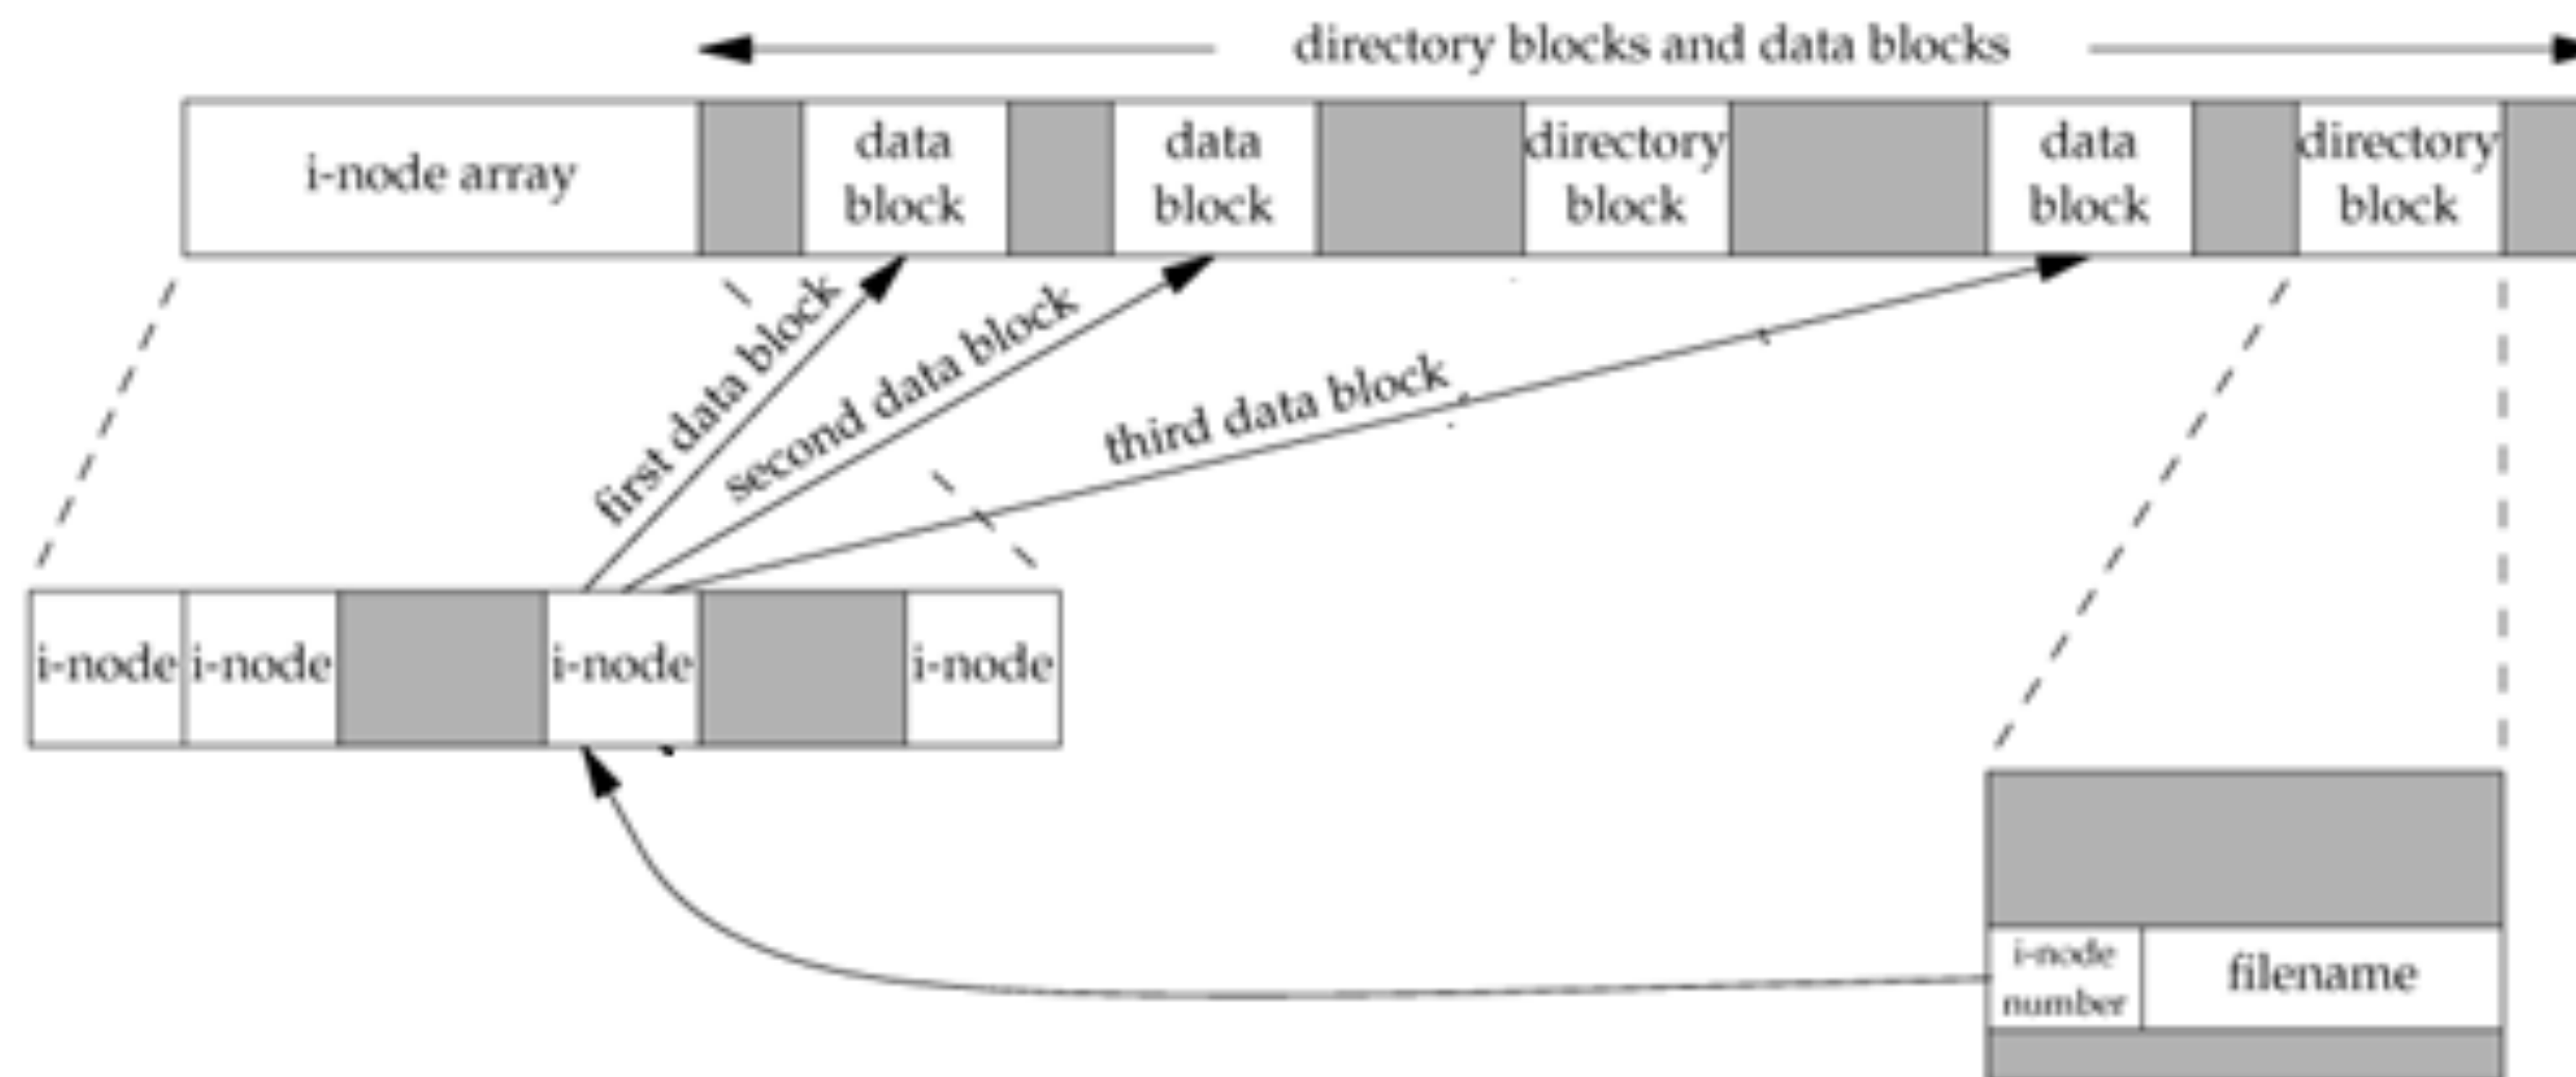
# Inodes

- To move a file within a single filesystem, we can just "move" the directory entry (actually done by creating a new entry, and deleting the old one).



# Inodes

- To move a file within a single filesystem, we can just "move" the directory entry (actually done by creating a new entry, and deleting the old one).



## The Unix Filesystem

---

Visualizing the Unix Filesystem helps us understand the concept of *hard links*, what directories "look like", and how operations on a directory are independent of the files and their data.

Coming up: creating, removing, and renaming links (*hard* and *symbolic*)

link(2)          unlink(2)          rename(2)          symlink(2) / readlink(2)