

Advanced Programming in the UNIX Environment

Week 06, Segment 6: Process Control

**Department of Computer Science
Stevens Institute of Technology**

Jan Schaumann

`jschauma@stevens.edu`

`https://stevens.netmeister.org/631/`

```
while (getinput(buf, sizeof(buf))) {
    buf[strlen(buf) - 1] = '\0';

    if((pid=fork()) == -1) {
        fprintf(stderr, "shell: can't fork: %s\n",
                strerror(errno));
        continue;
    } else if (pid == 0) { /* child */
        execlp(buf, buf, (char *)0);
        fprintf(stderr, "shell: couldn't exec %s: %s\n", buf,
                strerror(errno));
        exit(EX_UNAVAILABLE);
    }

    /* parent waits */
    if ((pid=waitpid(pid, &status, 0)) < 0) {
        fprintf(stderr, "shell: waitpid error: %s\n",
                strerror(errno));
    }
}

exit(EX_OK);
```

fork(2)

```
#include <unistd.h>
```

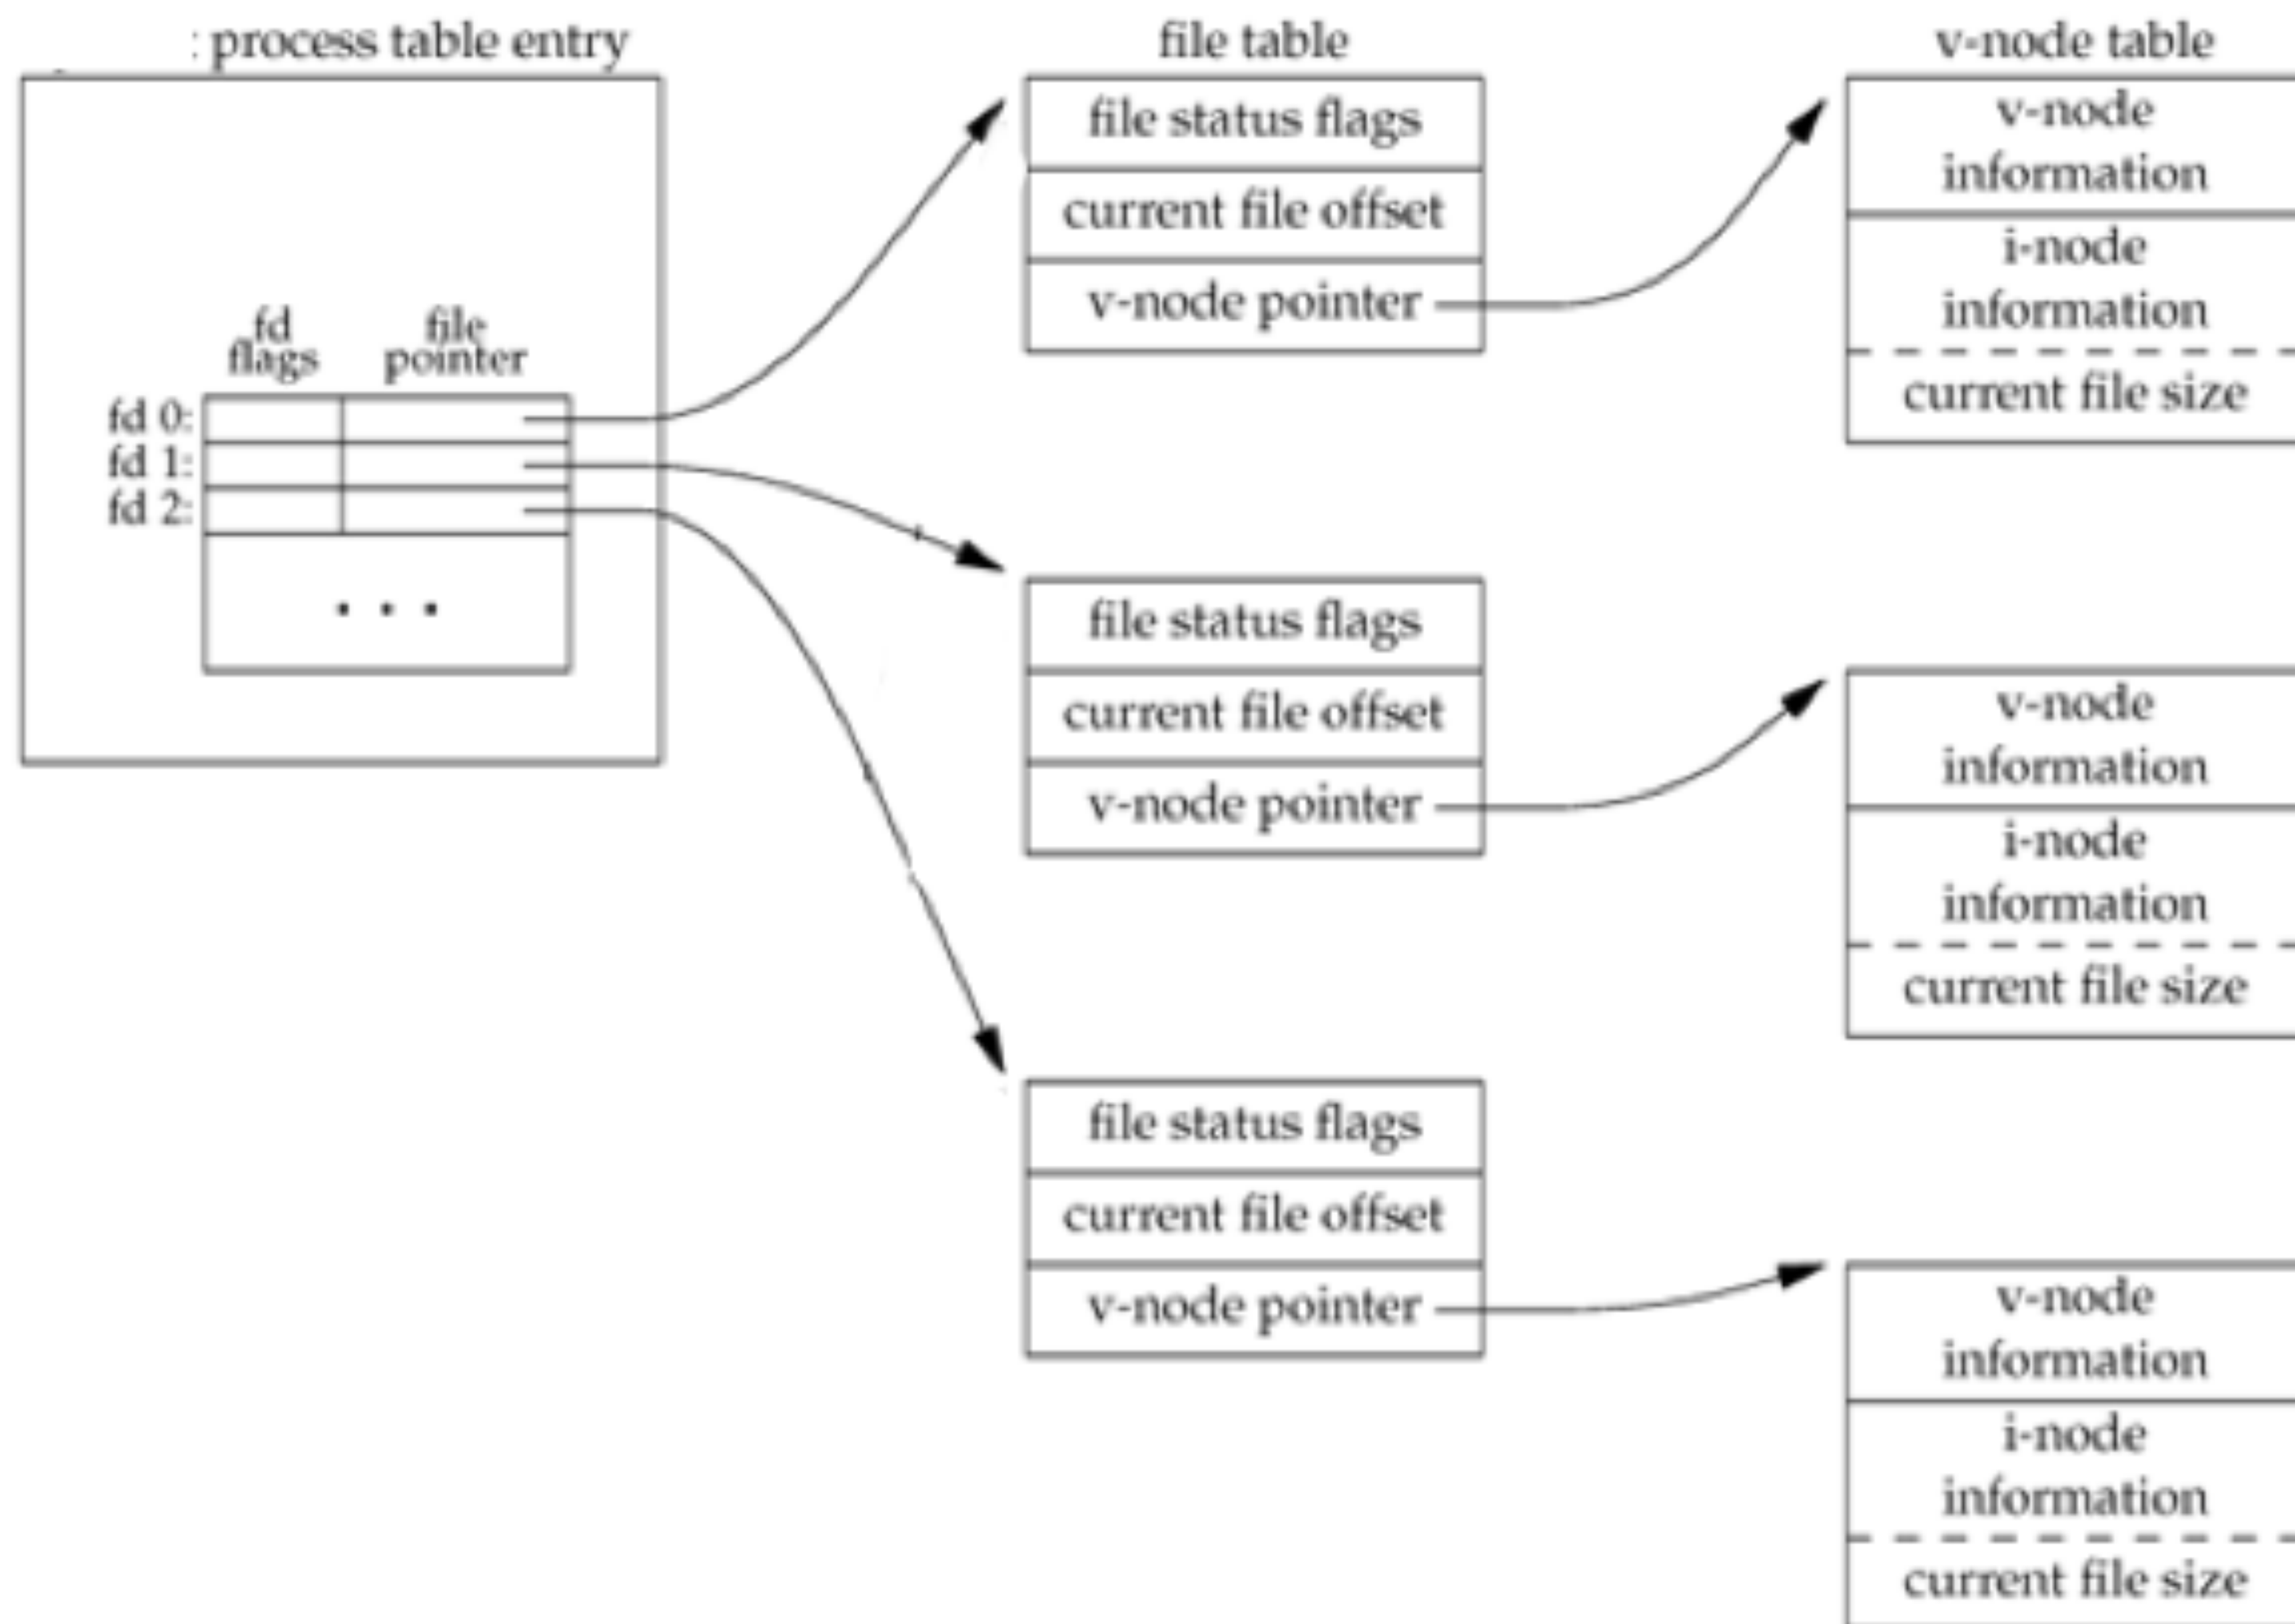
```
pid_t fork(void);
```

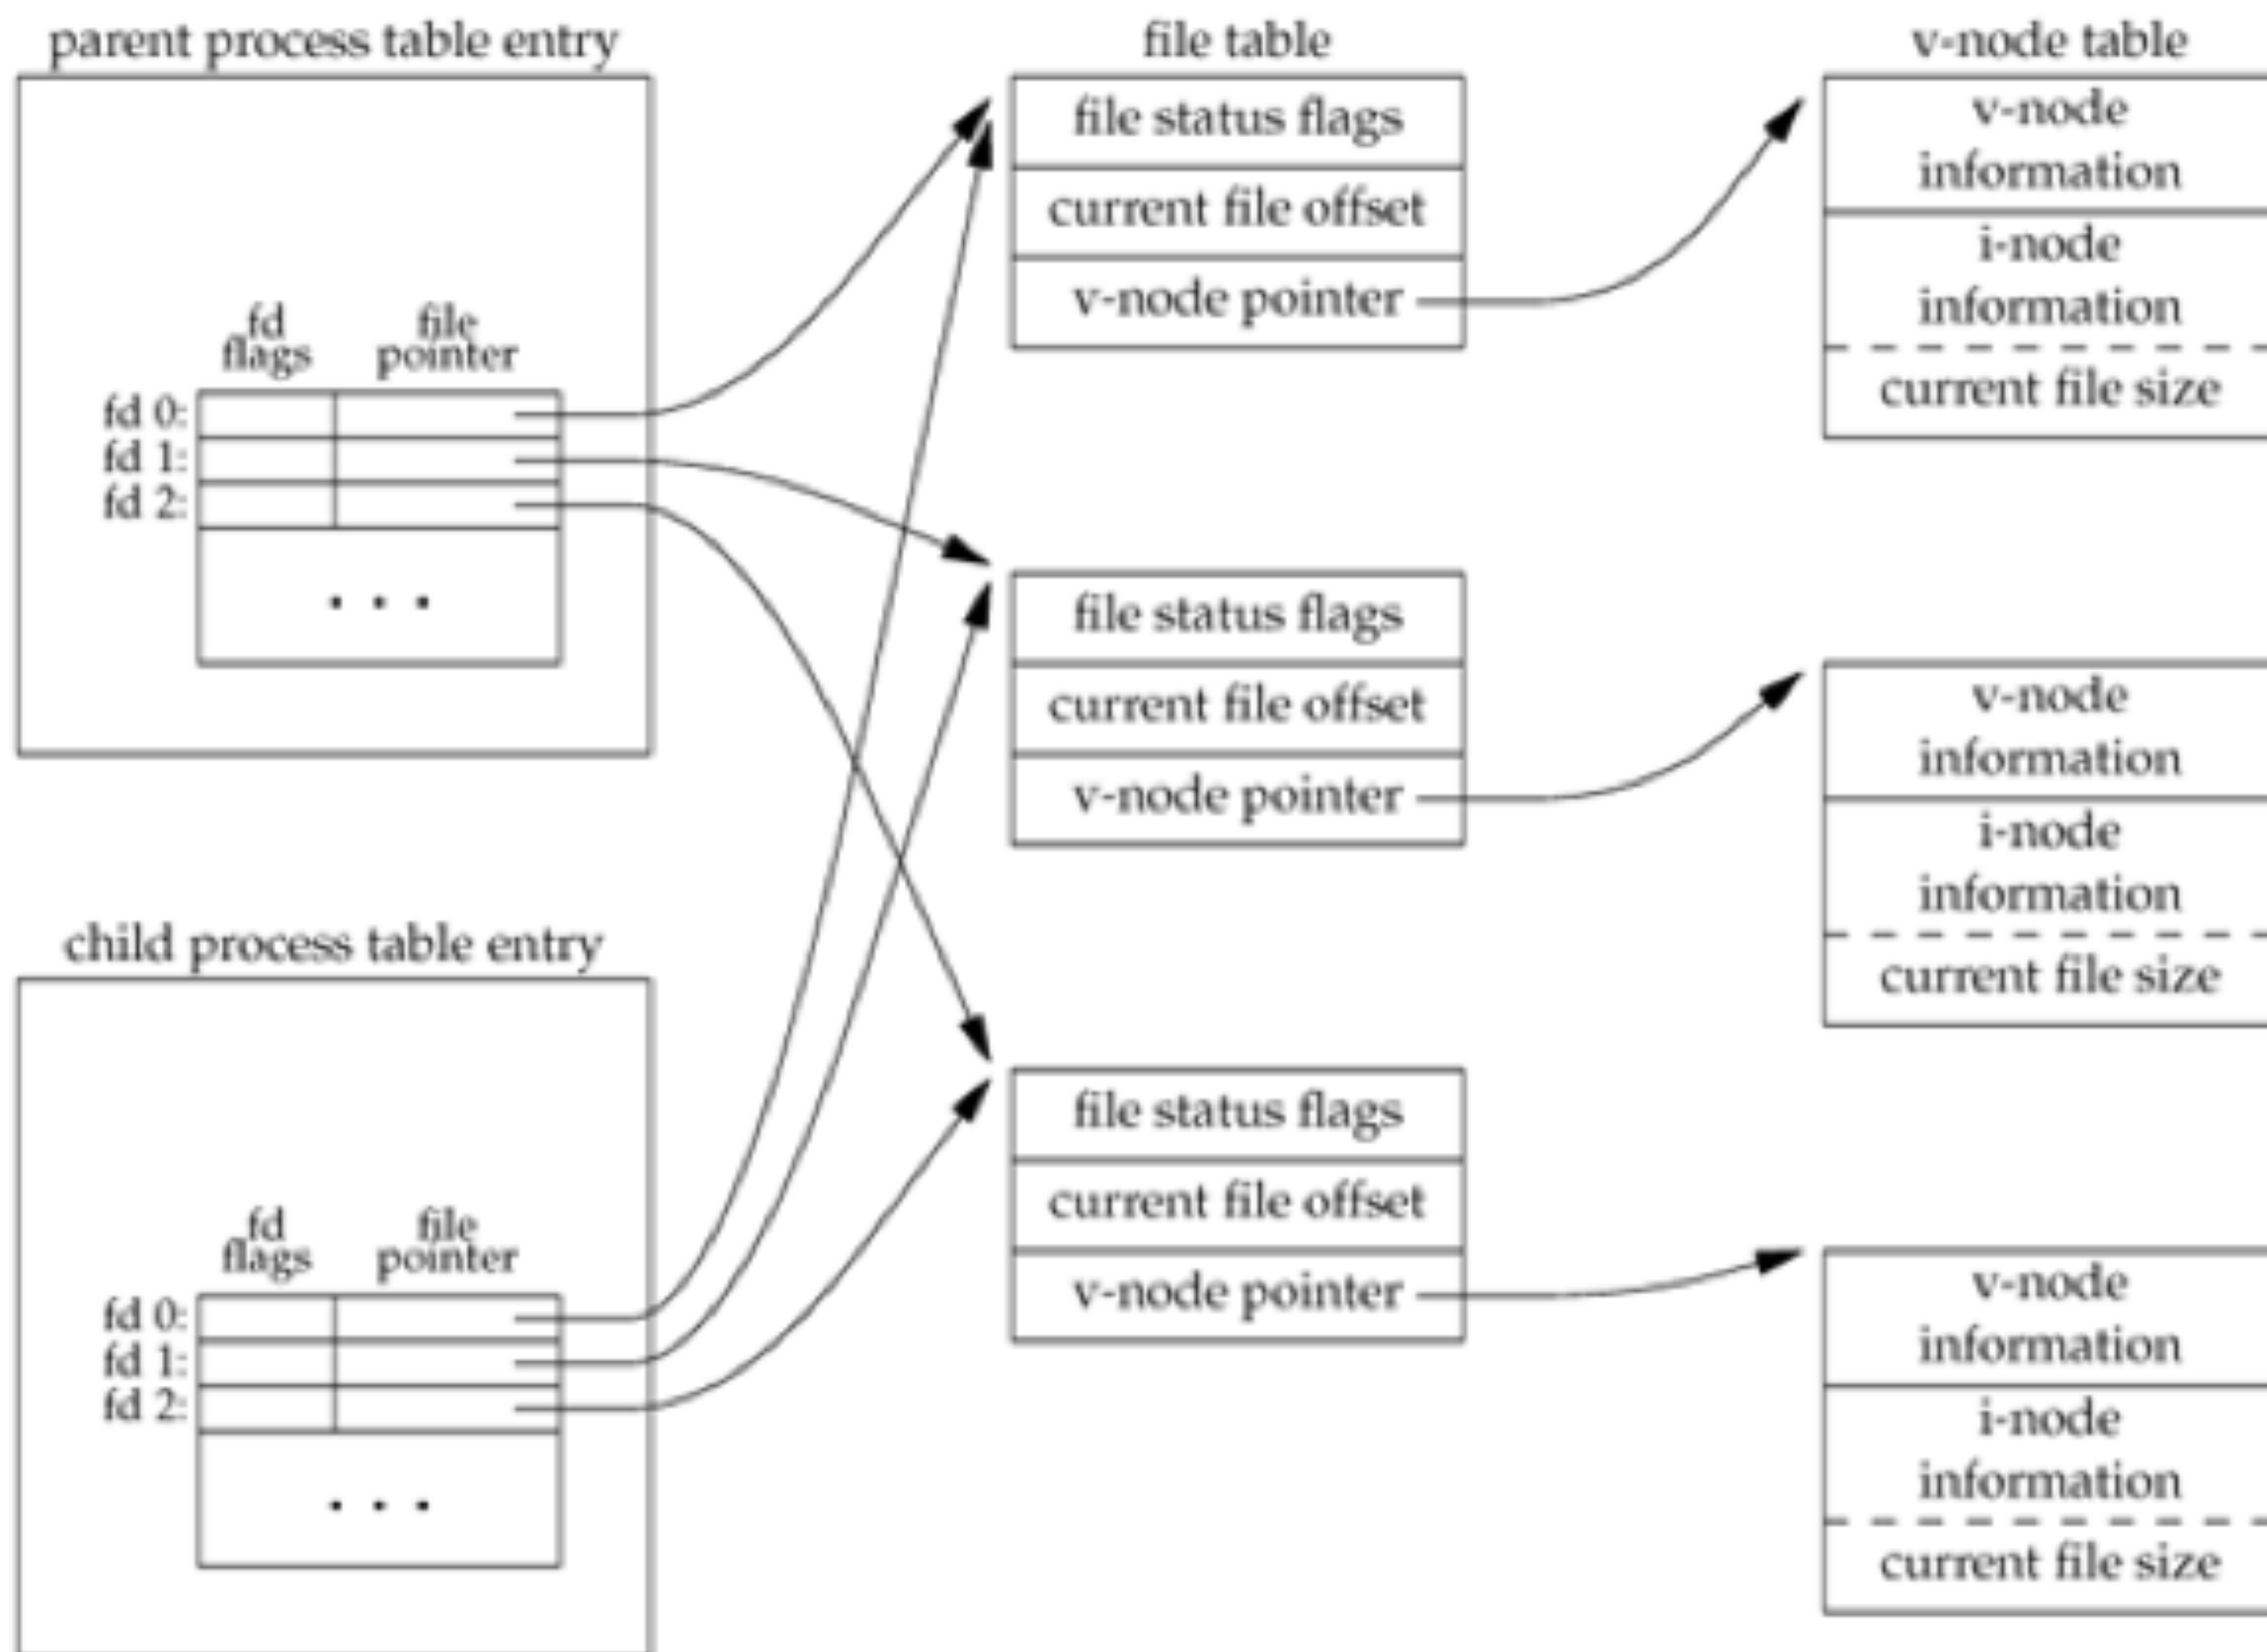
Returns: twice(!): 0 to the child, new pid to the parent; -1 on error

fork(2) causes creation of a new process. The new process (child) is an exact copy of the calling process (parent) except for the following:

- The child process has a unique process ID.
- The child process has a different parent process ID (*i.e.*, the processID of the parent process).
- The child process has its own copy of the parent's descriptors.
- The child process's resource utilizations are set to 0.

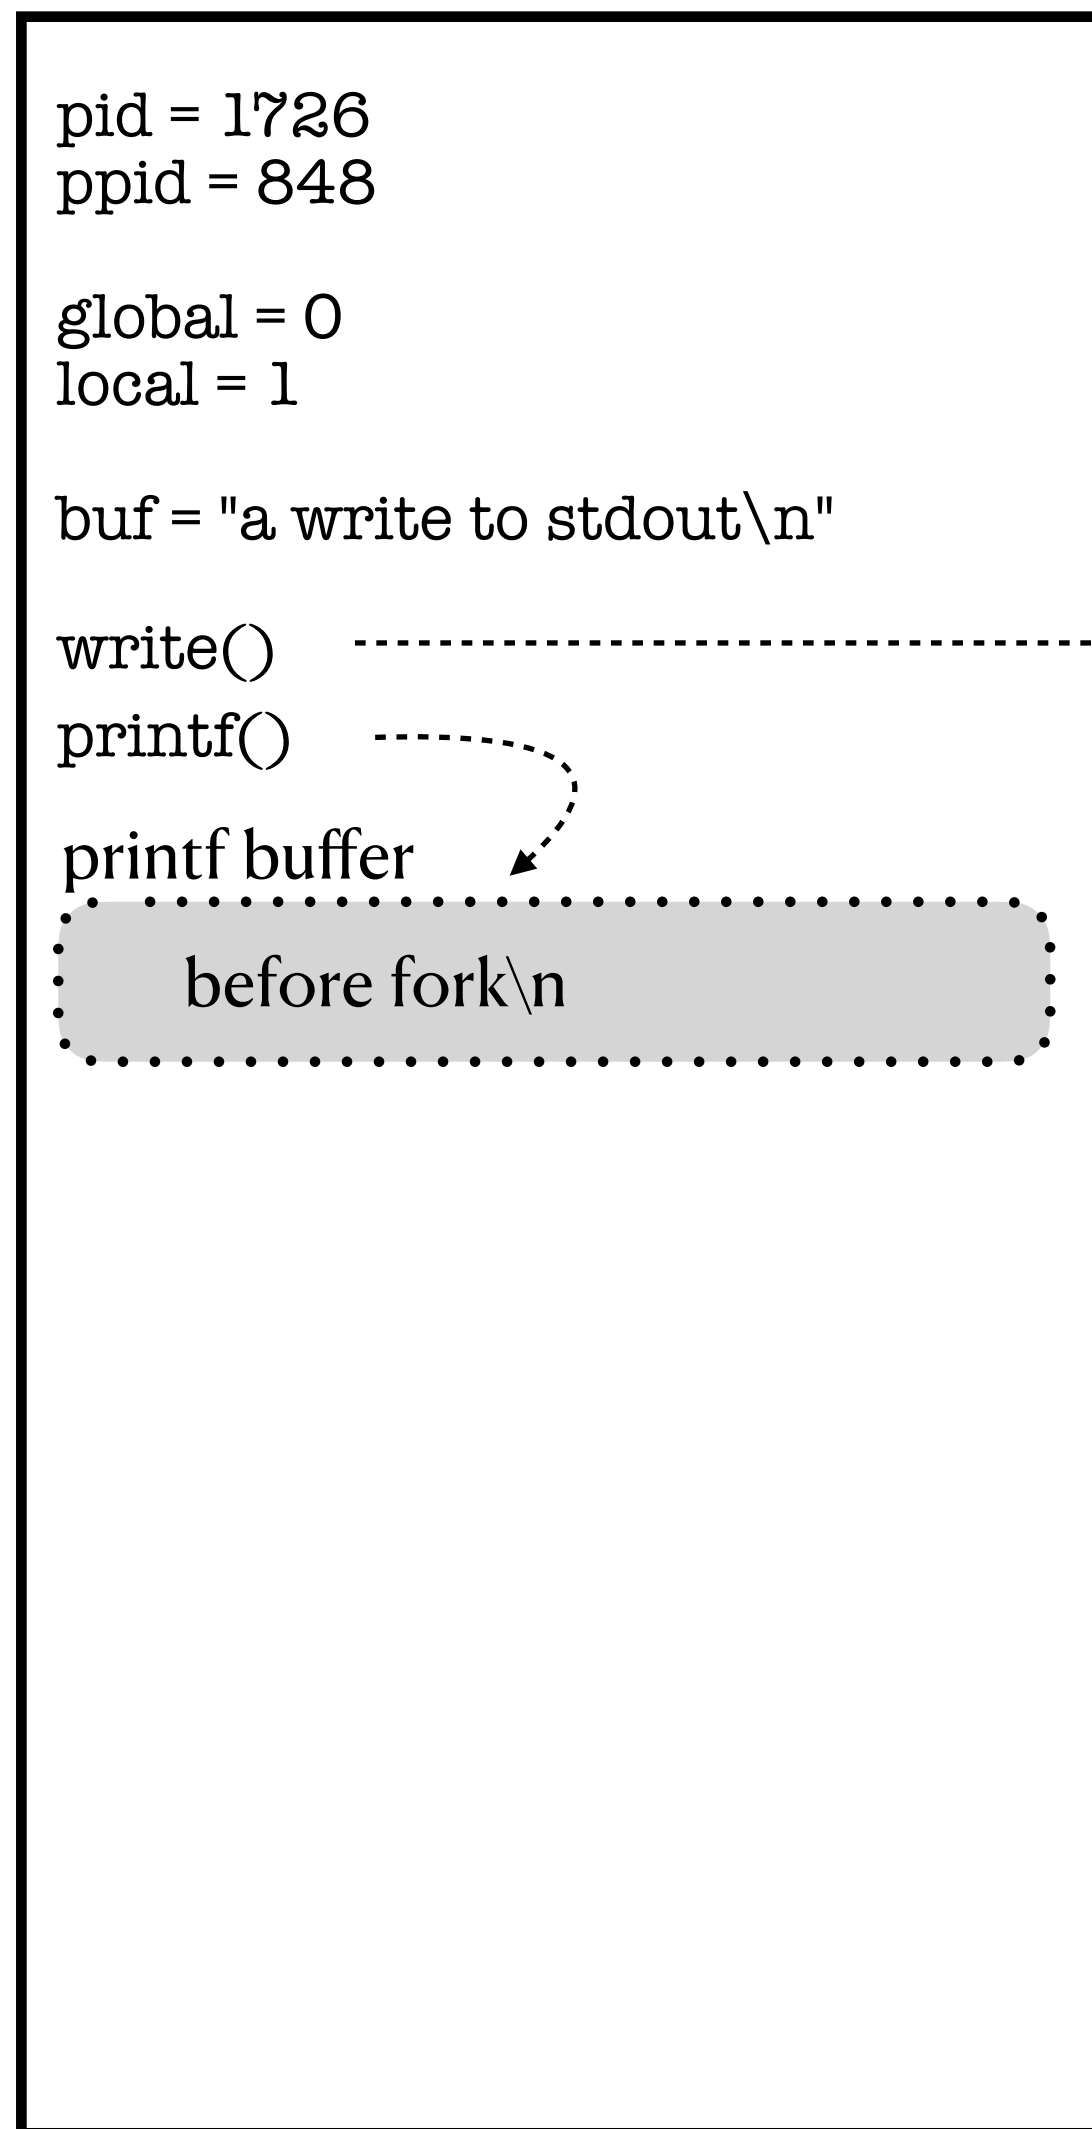
Note: no order of execution between child and parent is guaranteed!





```
[apue$ vim forkseek.c  
[apue$ cc -Wall -Werror -Wextra forkseek.c  
[apue$ ./a.out forkseek.c  
Starting pid is: 361  
361 offset is now: 0  
child 999 done seeking  
361 offset is now: 64  
999 offset is now: 96  
apue$ █
```

```
[apue$ vi forkflush.c
[apue$ cc -Wall -Werror -Wextra forkflush.c
[apue$ ./a.out
a write to stdout
before fork
pid = 2154, ppid = 1726, global = 1, local = 2
pid = 1726, ppid = 848, global = -1, local = 0
[apue$ echo $$
848
[apue$ ./a.out | cat
a write to stdout
before fork
pid = 1449, ppid = 1269, global = 1, local = 2
before fork
pid = 1269, ppid = 848, global = -1, local = 0
apue$ █
```



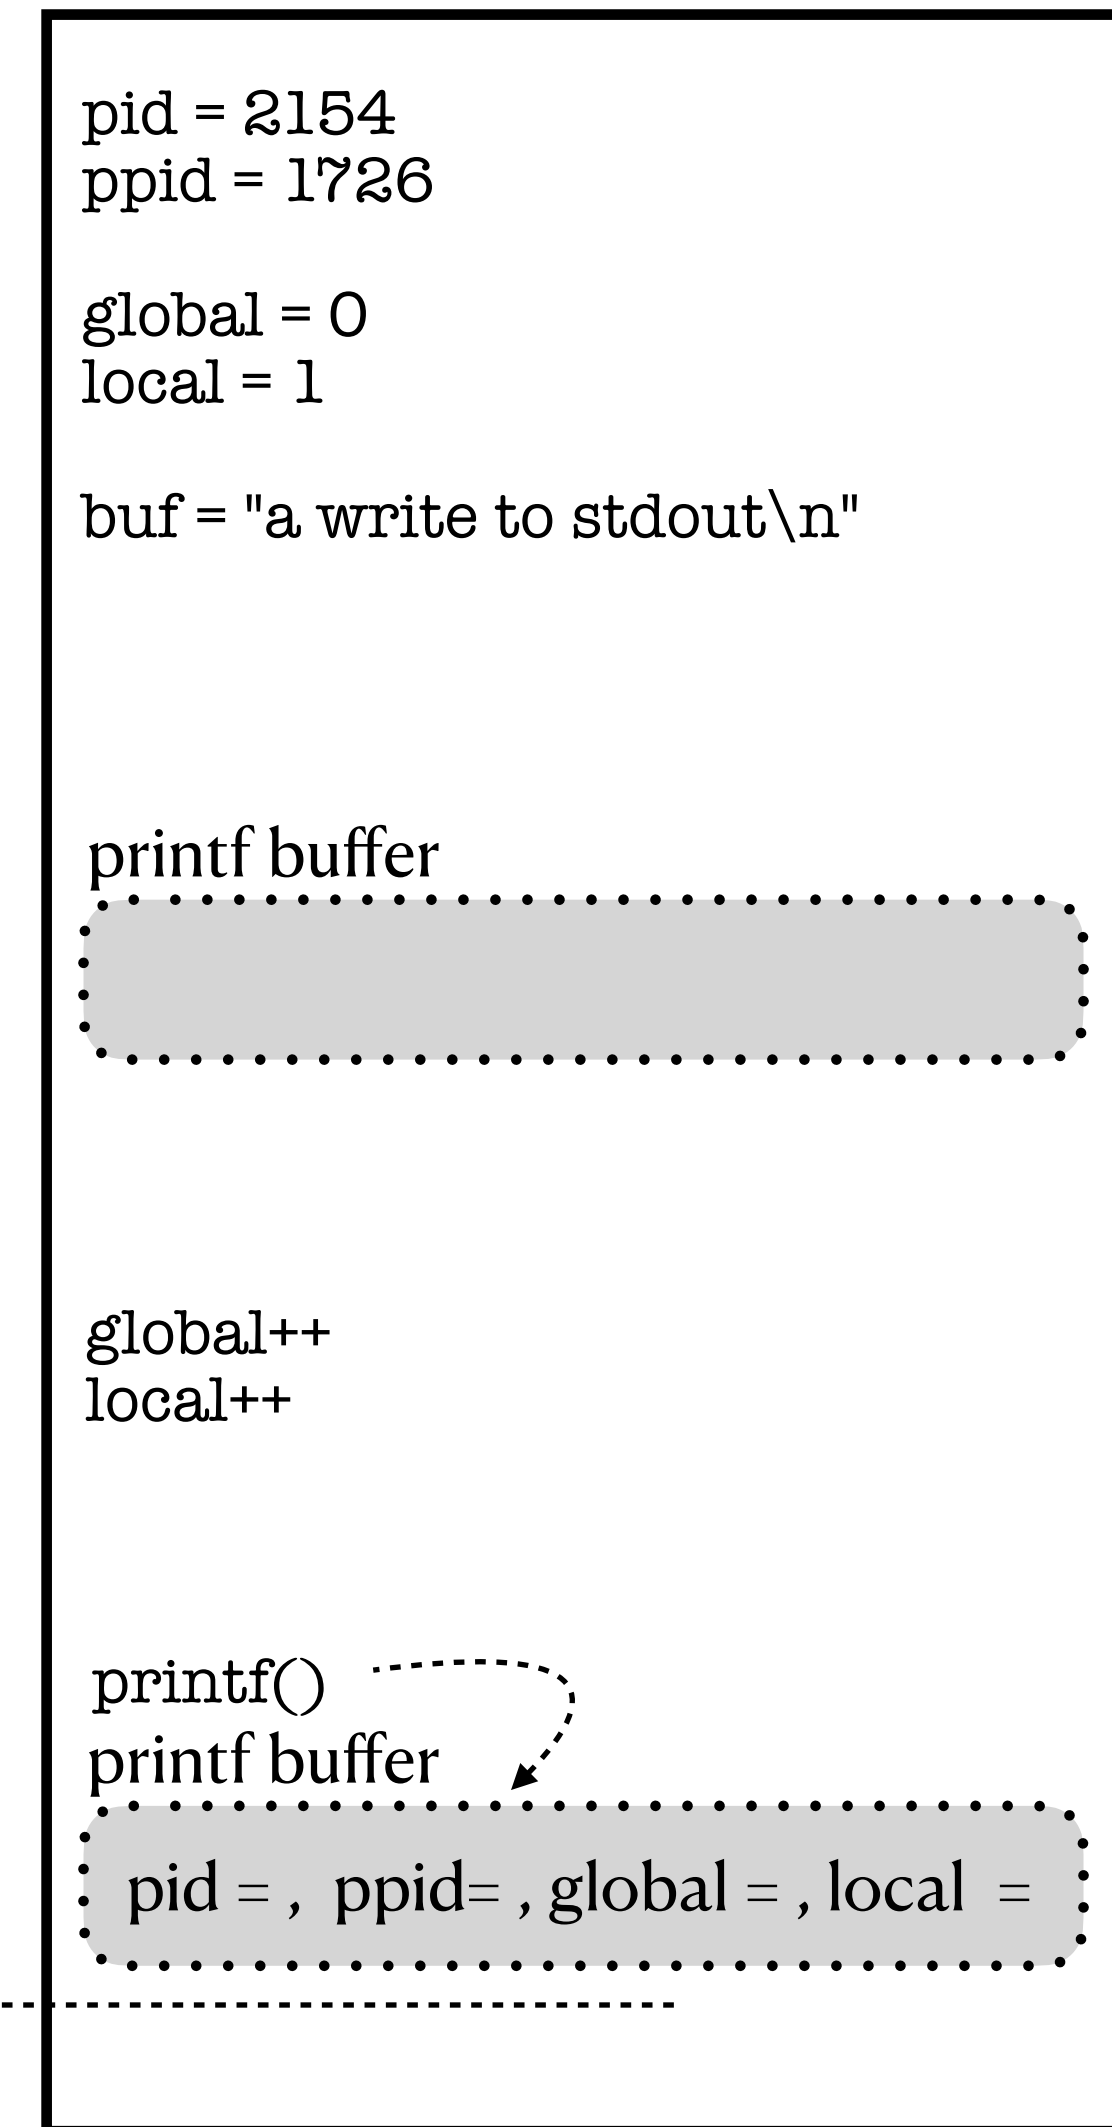

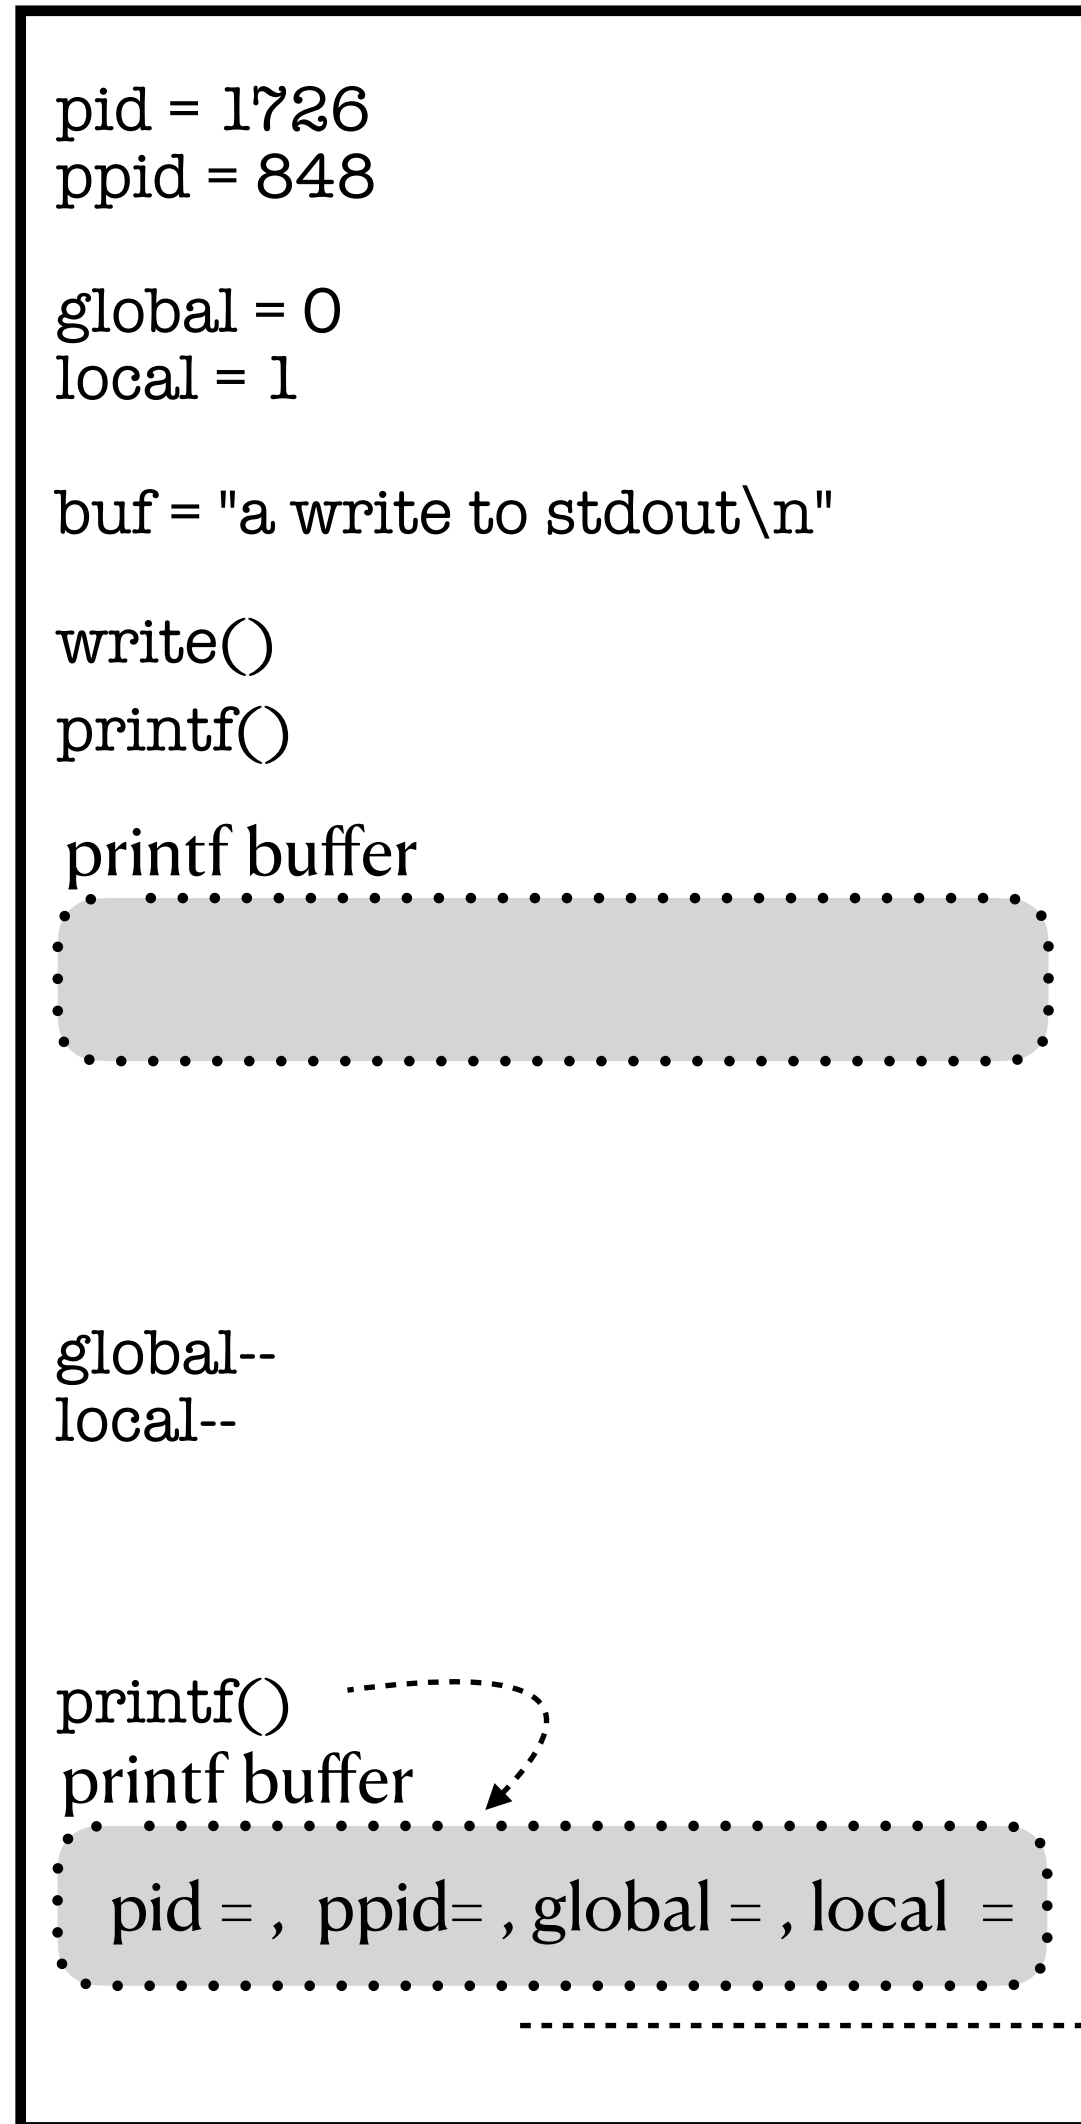

```
pid = 1726
ppid = 848


global = 0
local = 1

buf = "a write to stdout\n"

write()
printf()
printf buffer
before fork\n
fork()
```







exec(3)

```
#include <unistd.h>

int execl(const char *path, const char *arg, ...);
int execlp(const char *path, const char *arg, ...);
int execlpe(const char *path, const char *arg, ..., char *const envp[]);
int execl_e(const char *path, const char *arg, ..., char *const envp[]);
int execv(const char *path, char *const argv[]);
int execve(const char *path, char *const argv[], char *const envp[]);
int execvp(const char *path, char *const argv[]);
int execvpe(const char *path, char *const argv[], char *const envp[]);
```

Returns: doesn't; -1 on error

The `exec()` family of functions are used to completely replace a running process with a new process image. They all are front-ends for the `execve(2)` system call.

The exec(3) functions

- if it has a *v* in its name, argv's are a vector: `const * char argv[]`
- if it has an */* in its name, argv's are a list: `const char *arg0, .../* (char *) 0 */`
- if it has an *e* in its name, it takes a `char * const envp[]` array of environment variables
- if it has a *p* in its name, it uses the `PATH` environment variable to search for the file
- open file descriptors are inherited, unless the close-on-exec file flag was set
- ignored signals in the calling process are ignored after exec, but caught signals are reset to default
- real UID/GID is inherited; effective UID/GID is inherited unless the executable was `setuid/setgid`

wait(2) and waitpid(2)

```
#include <sys/wait.h>
```

```
pid_t wait(int *status);
```

```
pid_t waitpid(pid_t wpid, int *status, int options);
```

```
# include <sys/resource.h>
```

```
pid_t wait3(int *status, int options, struct rusage *rusage);
```

```
pid_t wait4(pid_t wpid, int *status, int options, struct rusage *rusage);
```

Returns: child PID on success; -1 on error

`wait()` suspends execution of the calling process until status information is available for a terminated child process.

`waitpid()` / `wait4()` allow waiting for a specific process or process group; `wait3()` / `wait4()` allow inspection of resource usage.

`wait(2)` and `waitpid(2)`

Once we get a termination status back in `status`, we'd like to be able to determine how a child died. We do this with the following macros:

- `WIFEXITED(status)` – true if the child terminated normally; use `WEXITSTATUS(status)` to get the exit status
- `WIFSIGNALED(status)` – true if child terminated abnormally (by receiving a signal it didn't catch); use
 - `WTERMSIG(status)` to retrieve the signal number
 - `WCOREDUMP(status)` to see if the child left a core image
- `WIFSTOPPED(status)` – true if the child is currently stopped; use `WSTOPSIG(status)` to determine the signal that caused this

Additionally, `wait(2)` will block until a child terminates; pass `WNOHANG` to `waitpid(2)` / `wait(4)` to return immediately.

What happens if we don't wait(2)?



4175 pts/1 Z+ 0:00.00 (a.out)

====

1102 pts/1 Z+ 0:00.00 (a.out)

2668 pts/1 Z+ 0:00.00 (a.out)

2974 pts/1 Z+ 0:00.00 (a.out)

2996 pts/1 S+ 0:00.00 ./a.out

4175 pts/1 Z+ 0:00.00 (a.out)

====

1102 pts/1 Z+ 0:00.00 (a.out)

2668 pts/1 Z+ 0:00.00 (a.out)

2974 pts/1 Z+ 0:00.00 (a.out)

2996 pts/1 S+ 0:00.00 ./a.out

3723 pts/1 Z+ 0:00.00 (a.out)

4175 pts/1 Z+ 0:00.00 (a.out)

I'm going to sleep – try to kill my zombie children, if you like.

[1] Terminated ./a.out

apue\$ ps

PID	TTY	STAT	TIME	COMMAND
701	pts/0	Is	0:00.09	-sh
2262	pts/0	S+	0:00.01	screen -r -d
2134	pts/1	0+	0:00.00	ps
4217	pts/1	Ss	0:00.01	/bin/sh
2388	pts/2	Is+	0:00.01	/bin/sh
2361	pts/3	Ss+	0:00.01	/bin/sh

Process Control

All processes not explicitly instantiated by the kernel were created via `fork(2)`.

`fork(2)` creates a copy of the current process, including file descriptors and output buffers.

To replace the current process with a new process image, use the `exec(3)` family of function.

After creating a new process via `fork(2)`, the parent process can `wait(2)` for the child process to reap its exit status and resource utilization; failure to do so will create a zombie process until the parent is terminated, at which point `init` will reap it.