

Advanced Programming in the UNIX Environment

Week 11, Segment 2: Of Linkers and Loaders

**Department of Computer Science
Stevens Institute of Technology**

Jan Schaumann

`jschauma@stevens.edu`

`https://stevens.netmeister.org/631/`

\$ readelf -l a.out

Elf file type is EXEC (Executable file)

Entry point 0x400570

There are 7 program headers, starting at offset 64

Program Headers:

Type	Offset	VirtAddr	PhysAddr
	FileSiz	MemSiz	Flags Align
PHDR	0x00000000000000000040	0x000000000000400040	0x000000000000400040
	0x0000000000000000188	0x00000000000000188	R 0x8
INTERP	0x00000000000000001c8	0x0000000000004001c8	0x0000000000004001c8
	0x0000000000000000017	0x000000000000000017	R 0x1

[Requesting program interpreter: /usr/libexec/ld.elf_so]

Linking and Loading

Let's revisit our lecture on Unix tools and the compile chain:

Linking and Loading

Let's revisit our lecture on Unix tools and the compile chain:

The compiler chain usually performs preprocessing (e.g., via `cpp(1)`)

```
$ cpp crypt.c crypt.i
```

Linking and Loading

Let's revisit our lecture on Unix tools and the compile chain:

The compiler chain usually performs preprocessing (e.g., via `cpp(1)`), compilation (`cc(1)`)

```
$ cpp crypt.c crypt.i
```

```
$ cc -S crypt.i
```

Linking and Loading

Let's revisit our lecture on Unix tools and the compile chain:

The compiler chain usually performs preprocessing (e.g., via `cpp(1)`), compilation (`cc(1)`), assembly (`as(1)`)

```
$ cpp crypt.c crypt.i
```

```
$ cc -S crypt.i
```

```
$ as -o crypt.o crypt.s
```

Linking and Loading


Let's revisit our lecture on Unix tools and the compile chain:

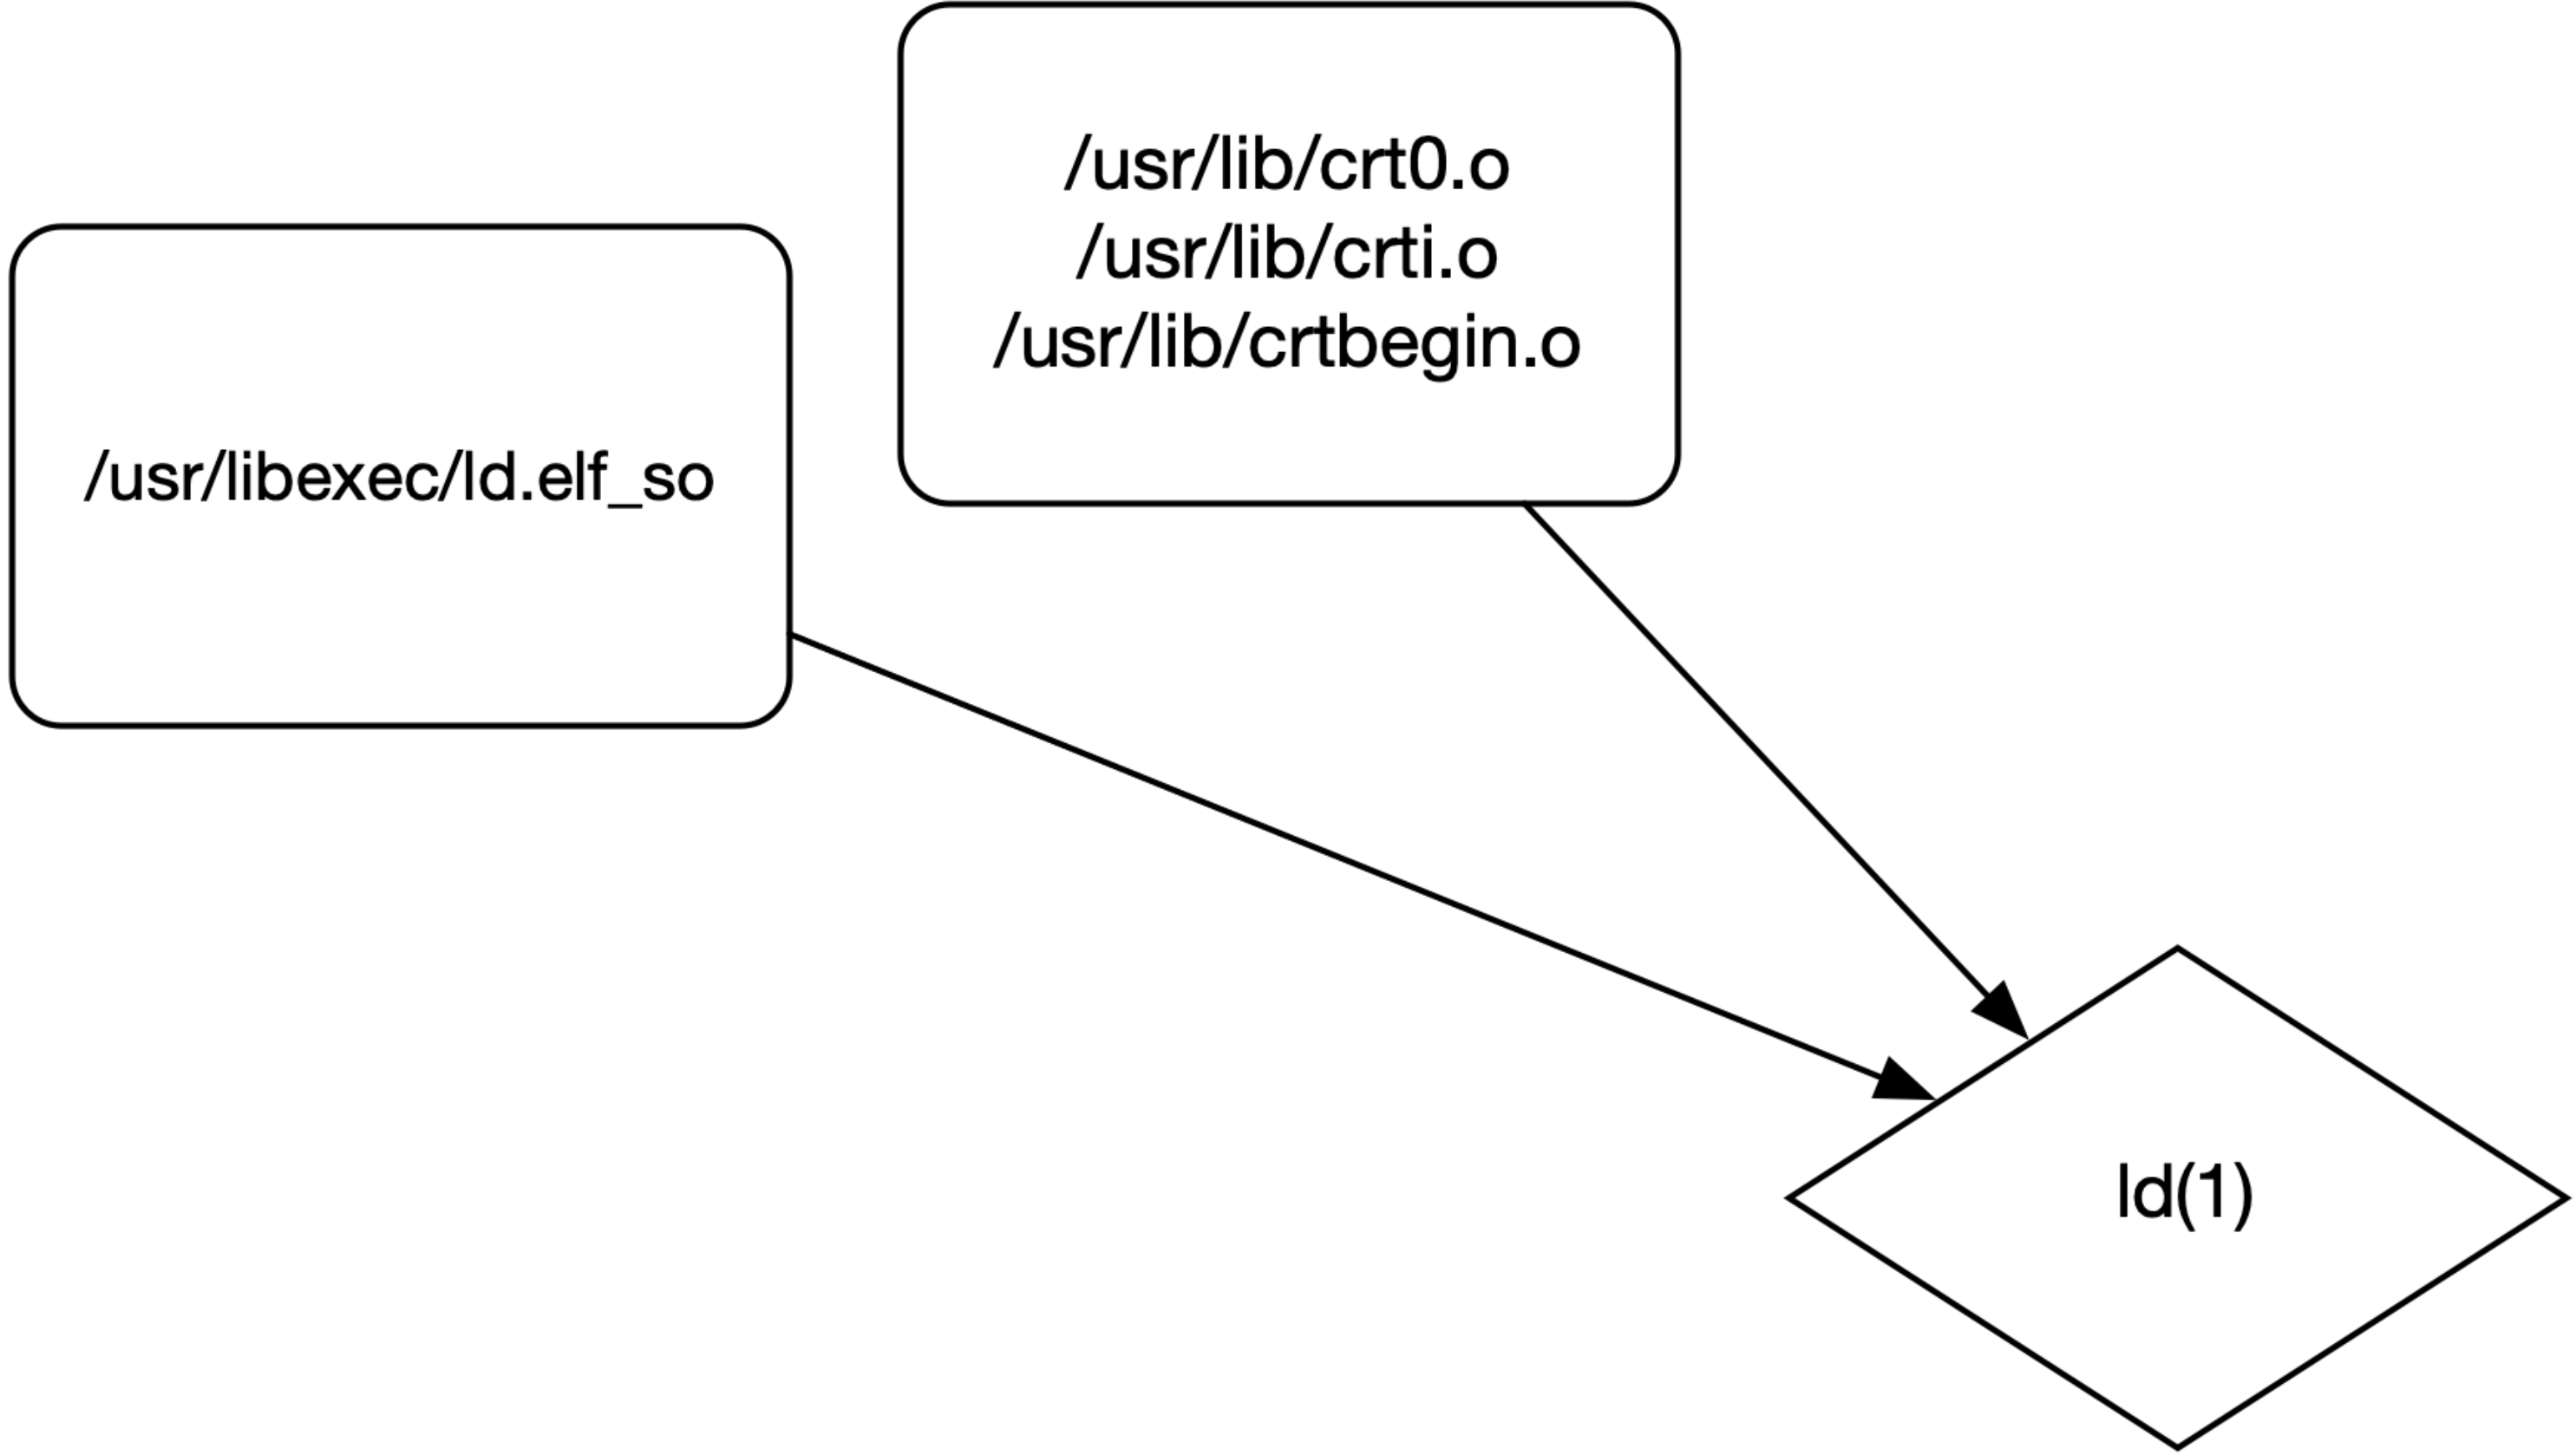
The compiler chain usually performs preprocessing (e.g., via `cpp(1)`), compilation (`cc(1)`), assembly (`as(1)`) and linking (`ld(1)`).

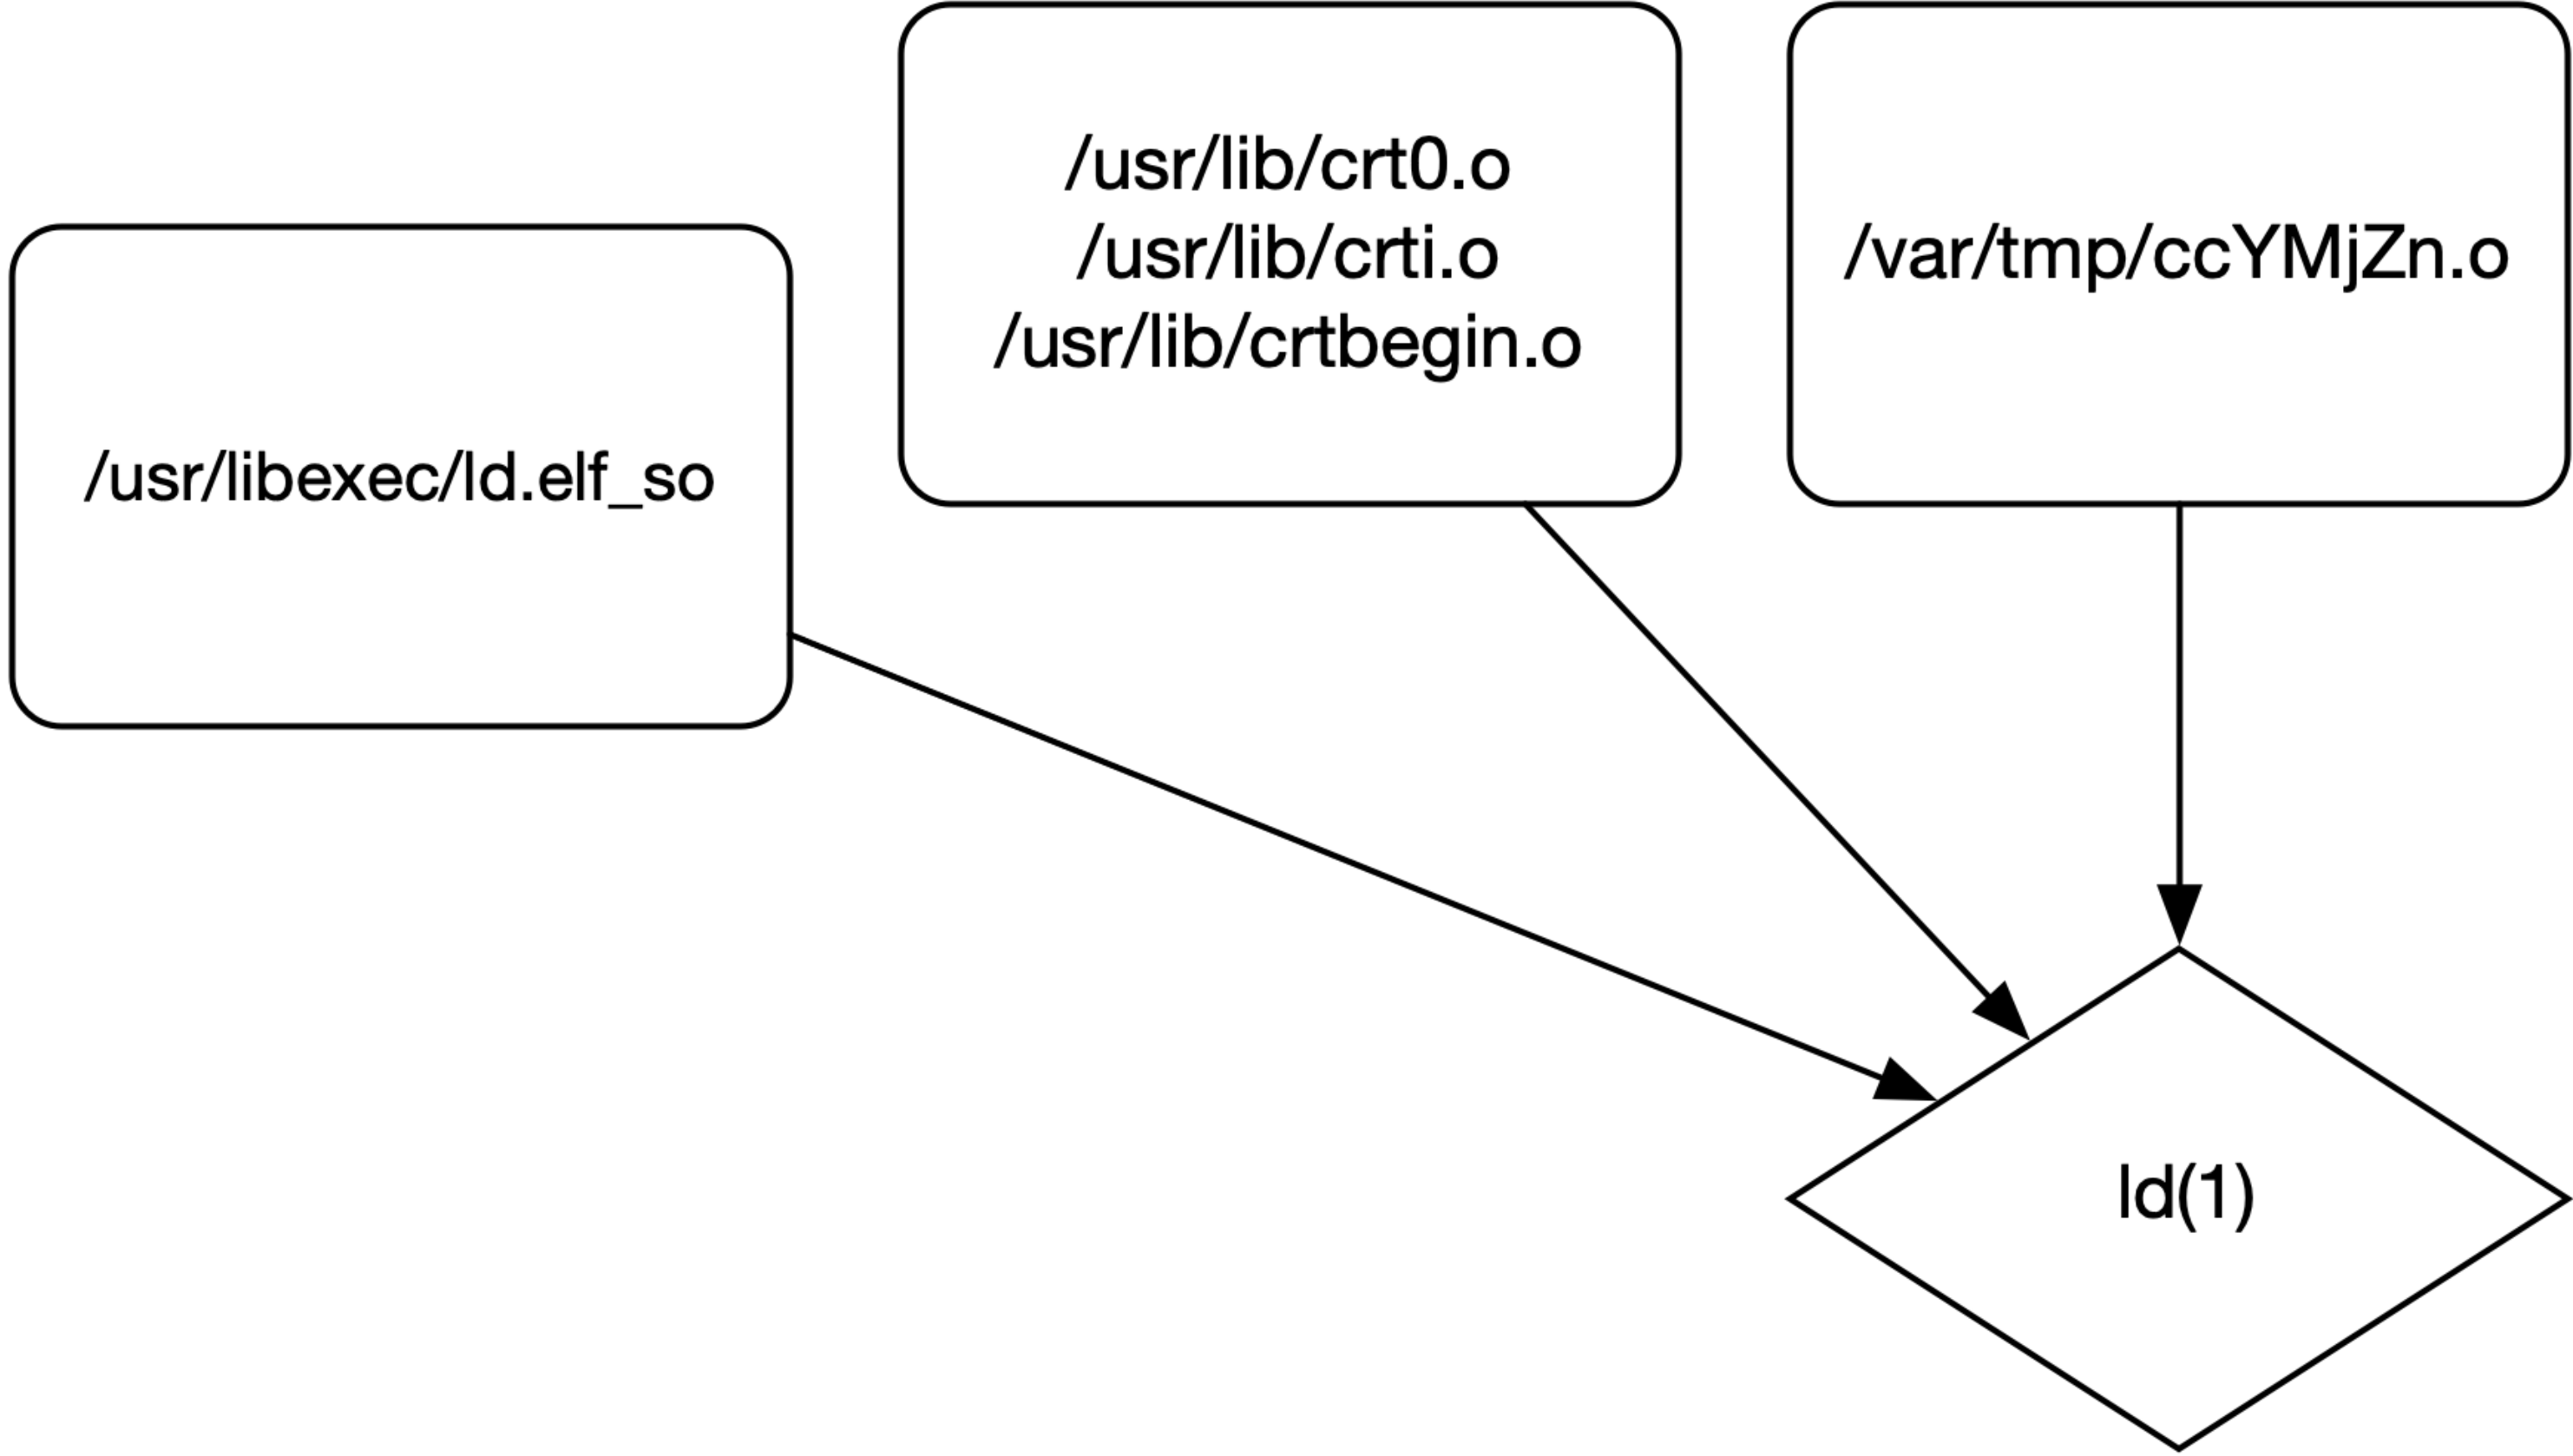
```
$ cpp crypt.c crypt.i
$ cc -S crypt.i
$ as -o crypt.o crypt.s
$ ld -dynamic-linker /usr/libexec/ld.elf_so \
    /usr/lib/crt0.o /usr/lib/crti.o /usr/lib/crtbegin.o \
    crypt.o -lcrypto -lc /usr/lib/crtend.o /usr/lib/crtn.o
```

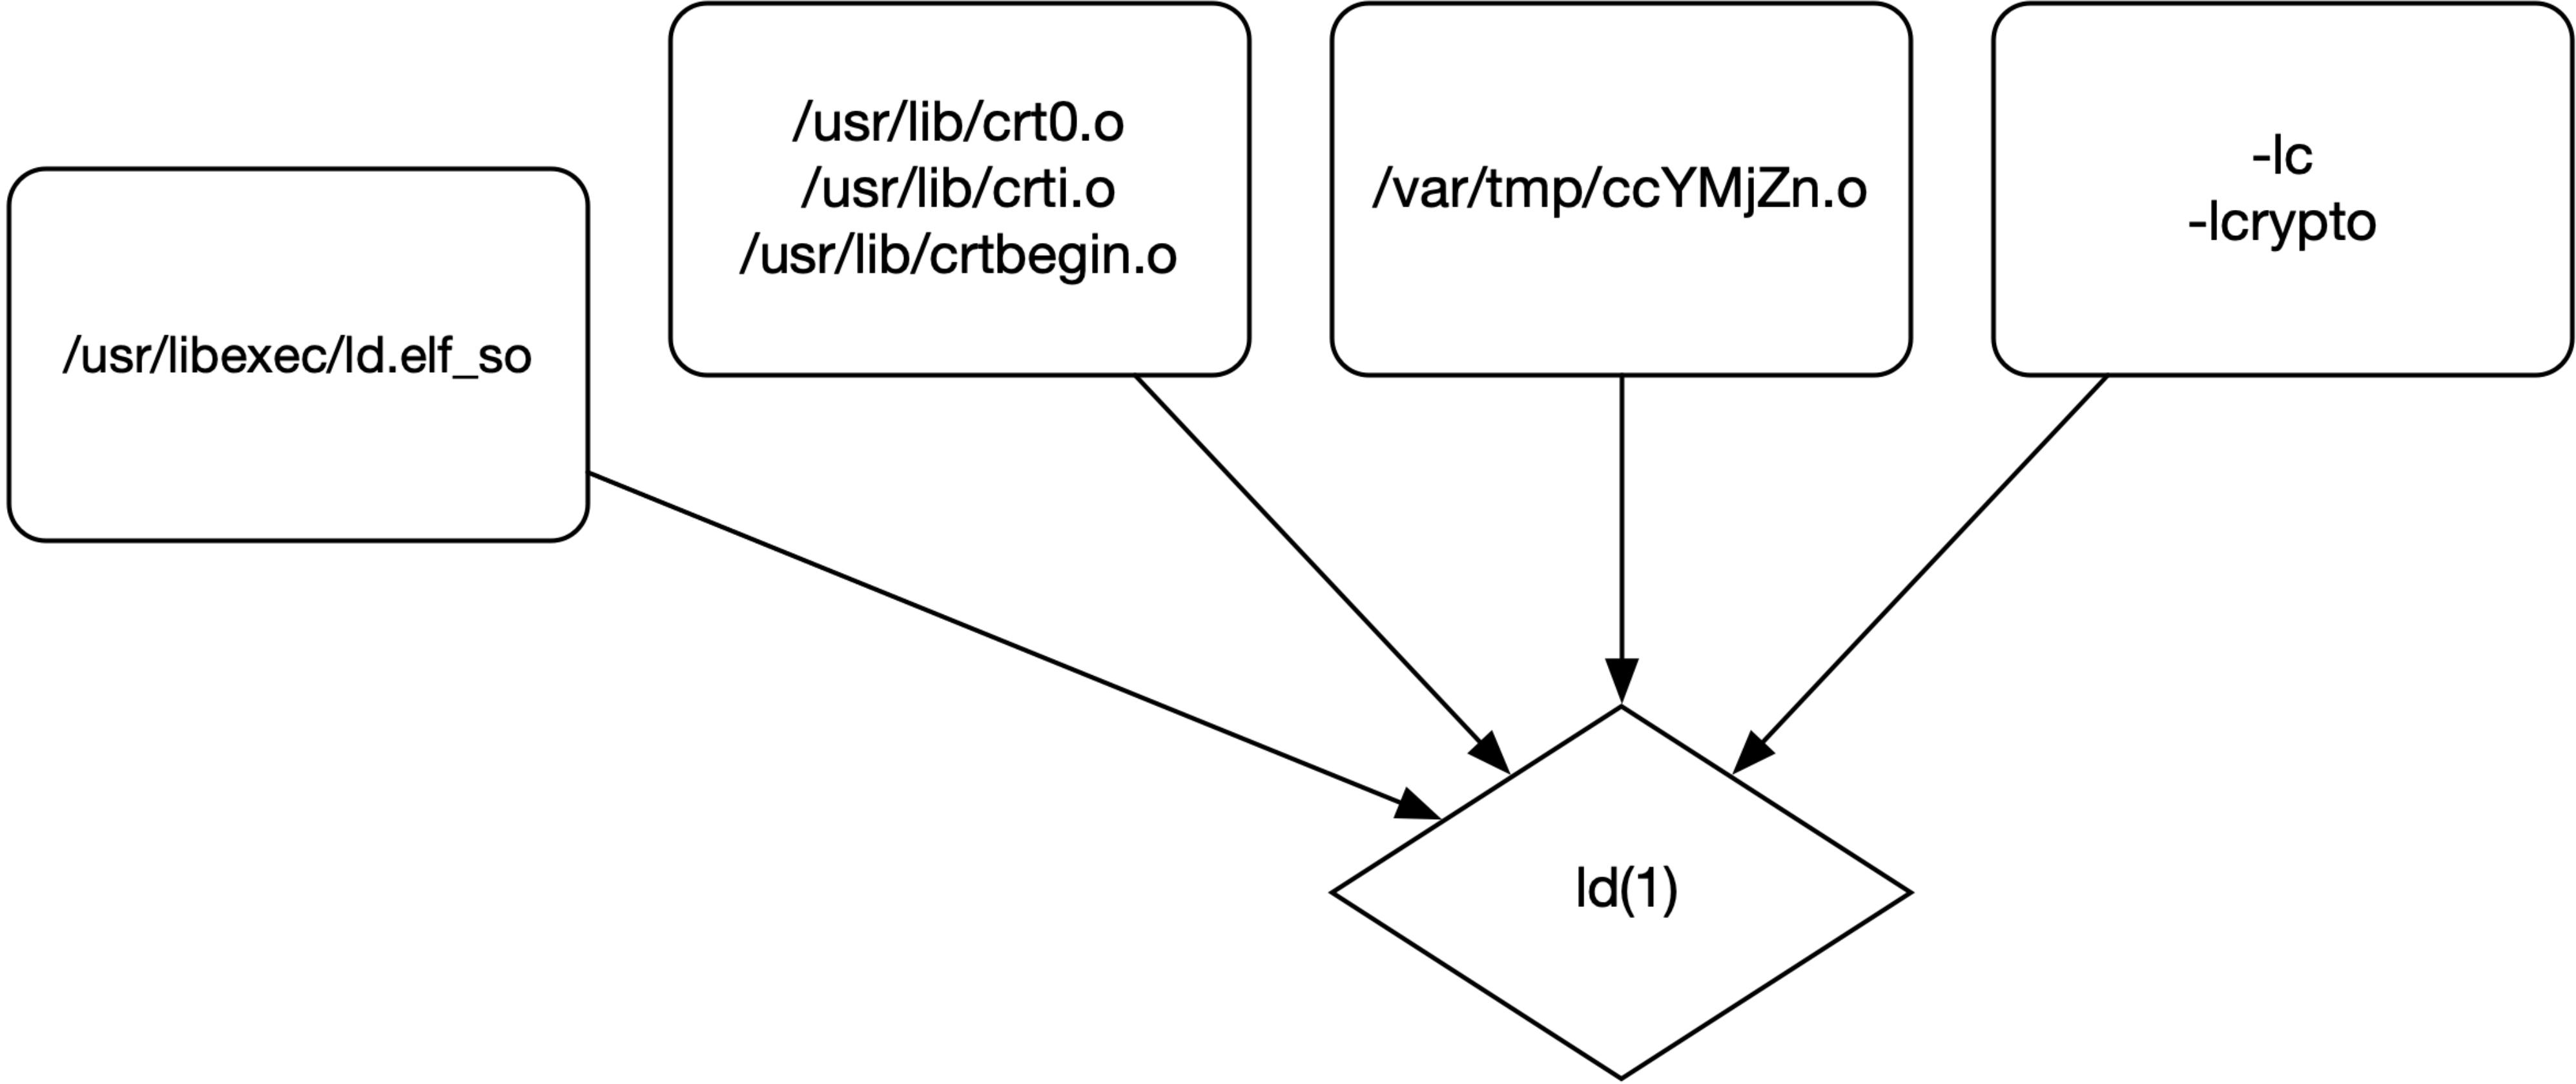
/usr/libexec/ld.elf_so

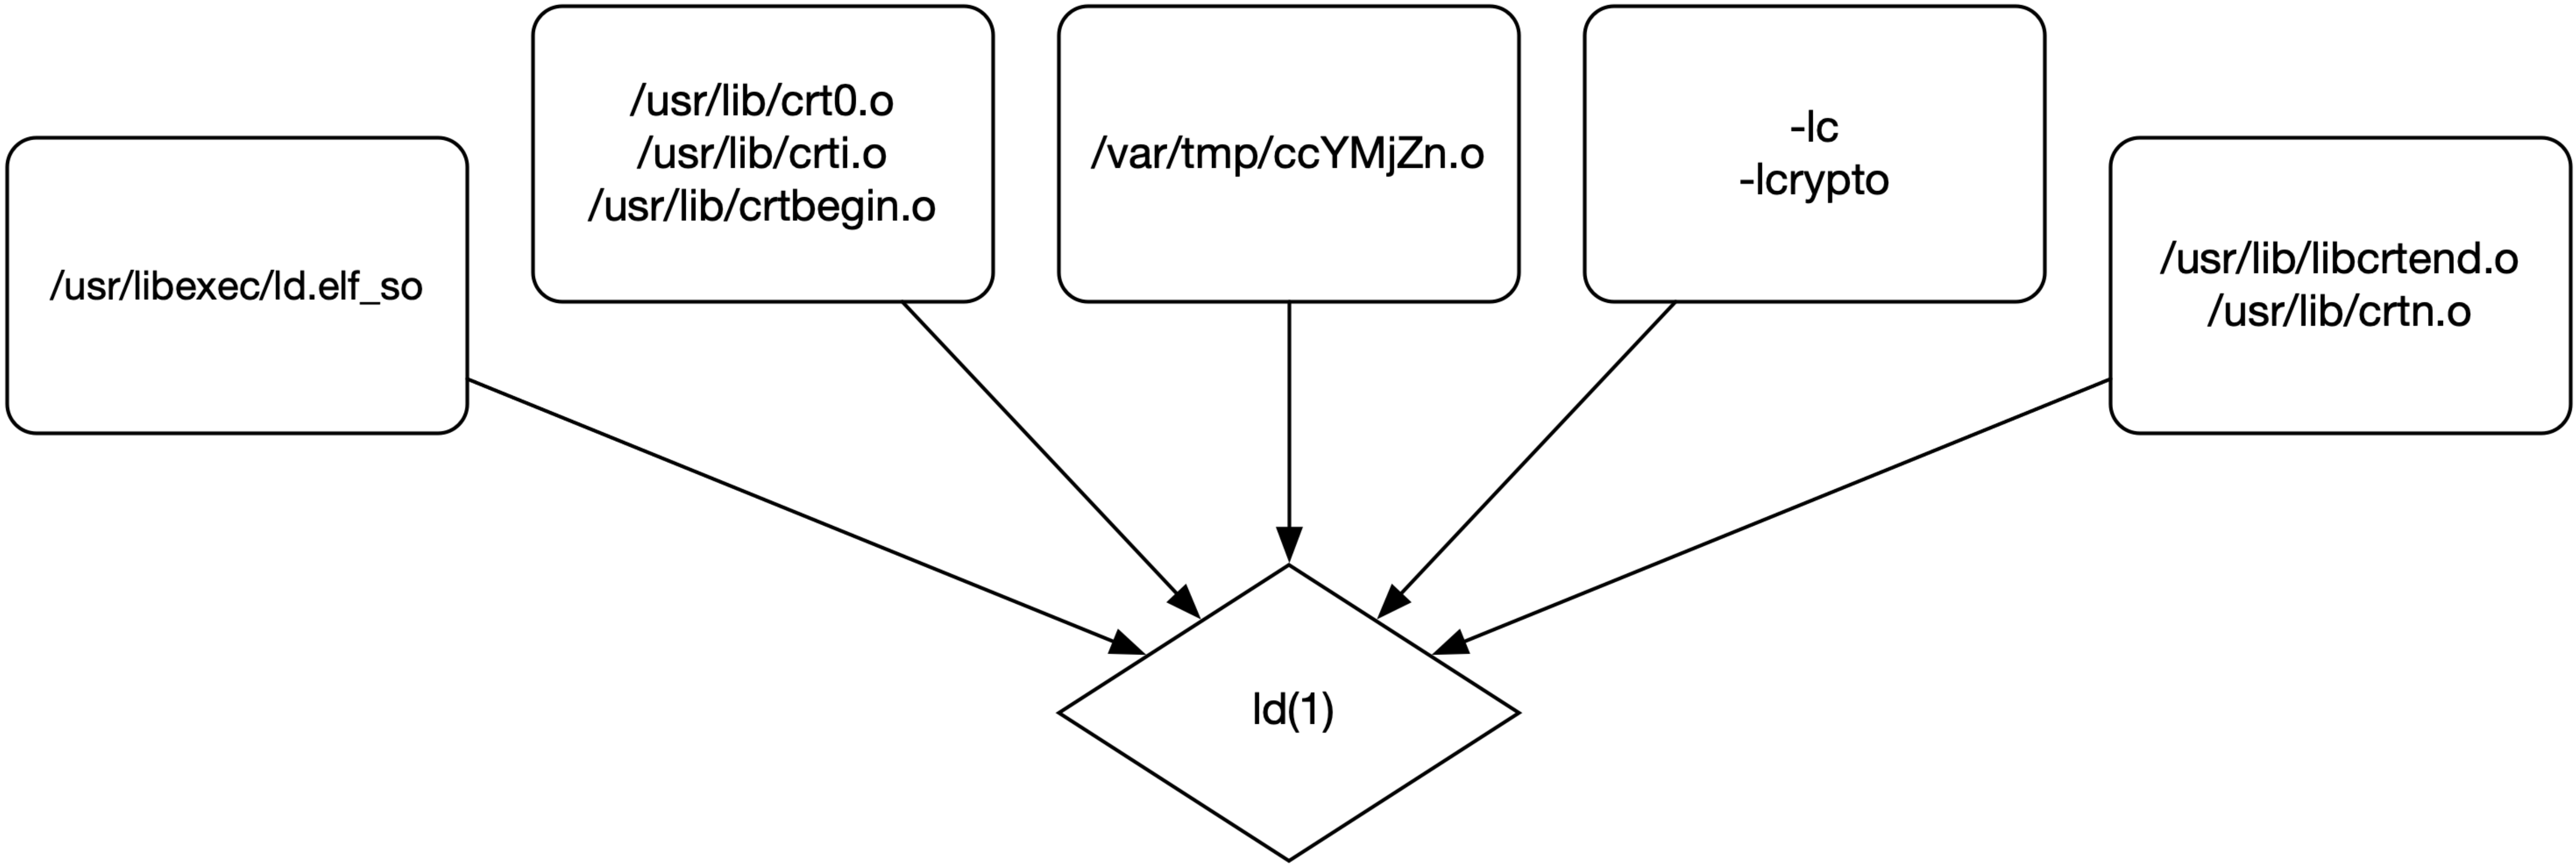
ld(1)

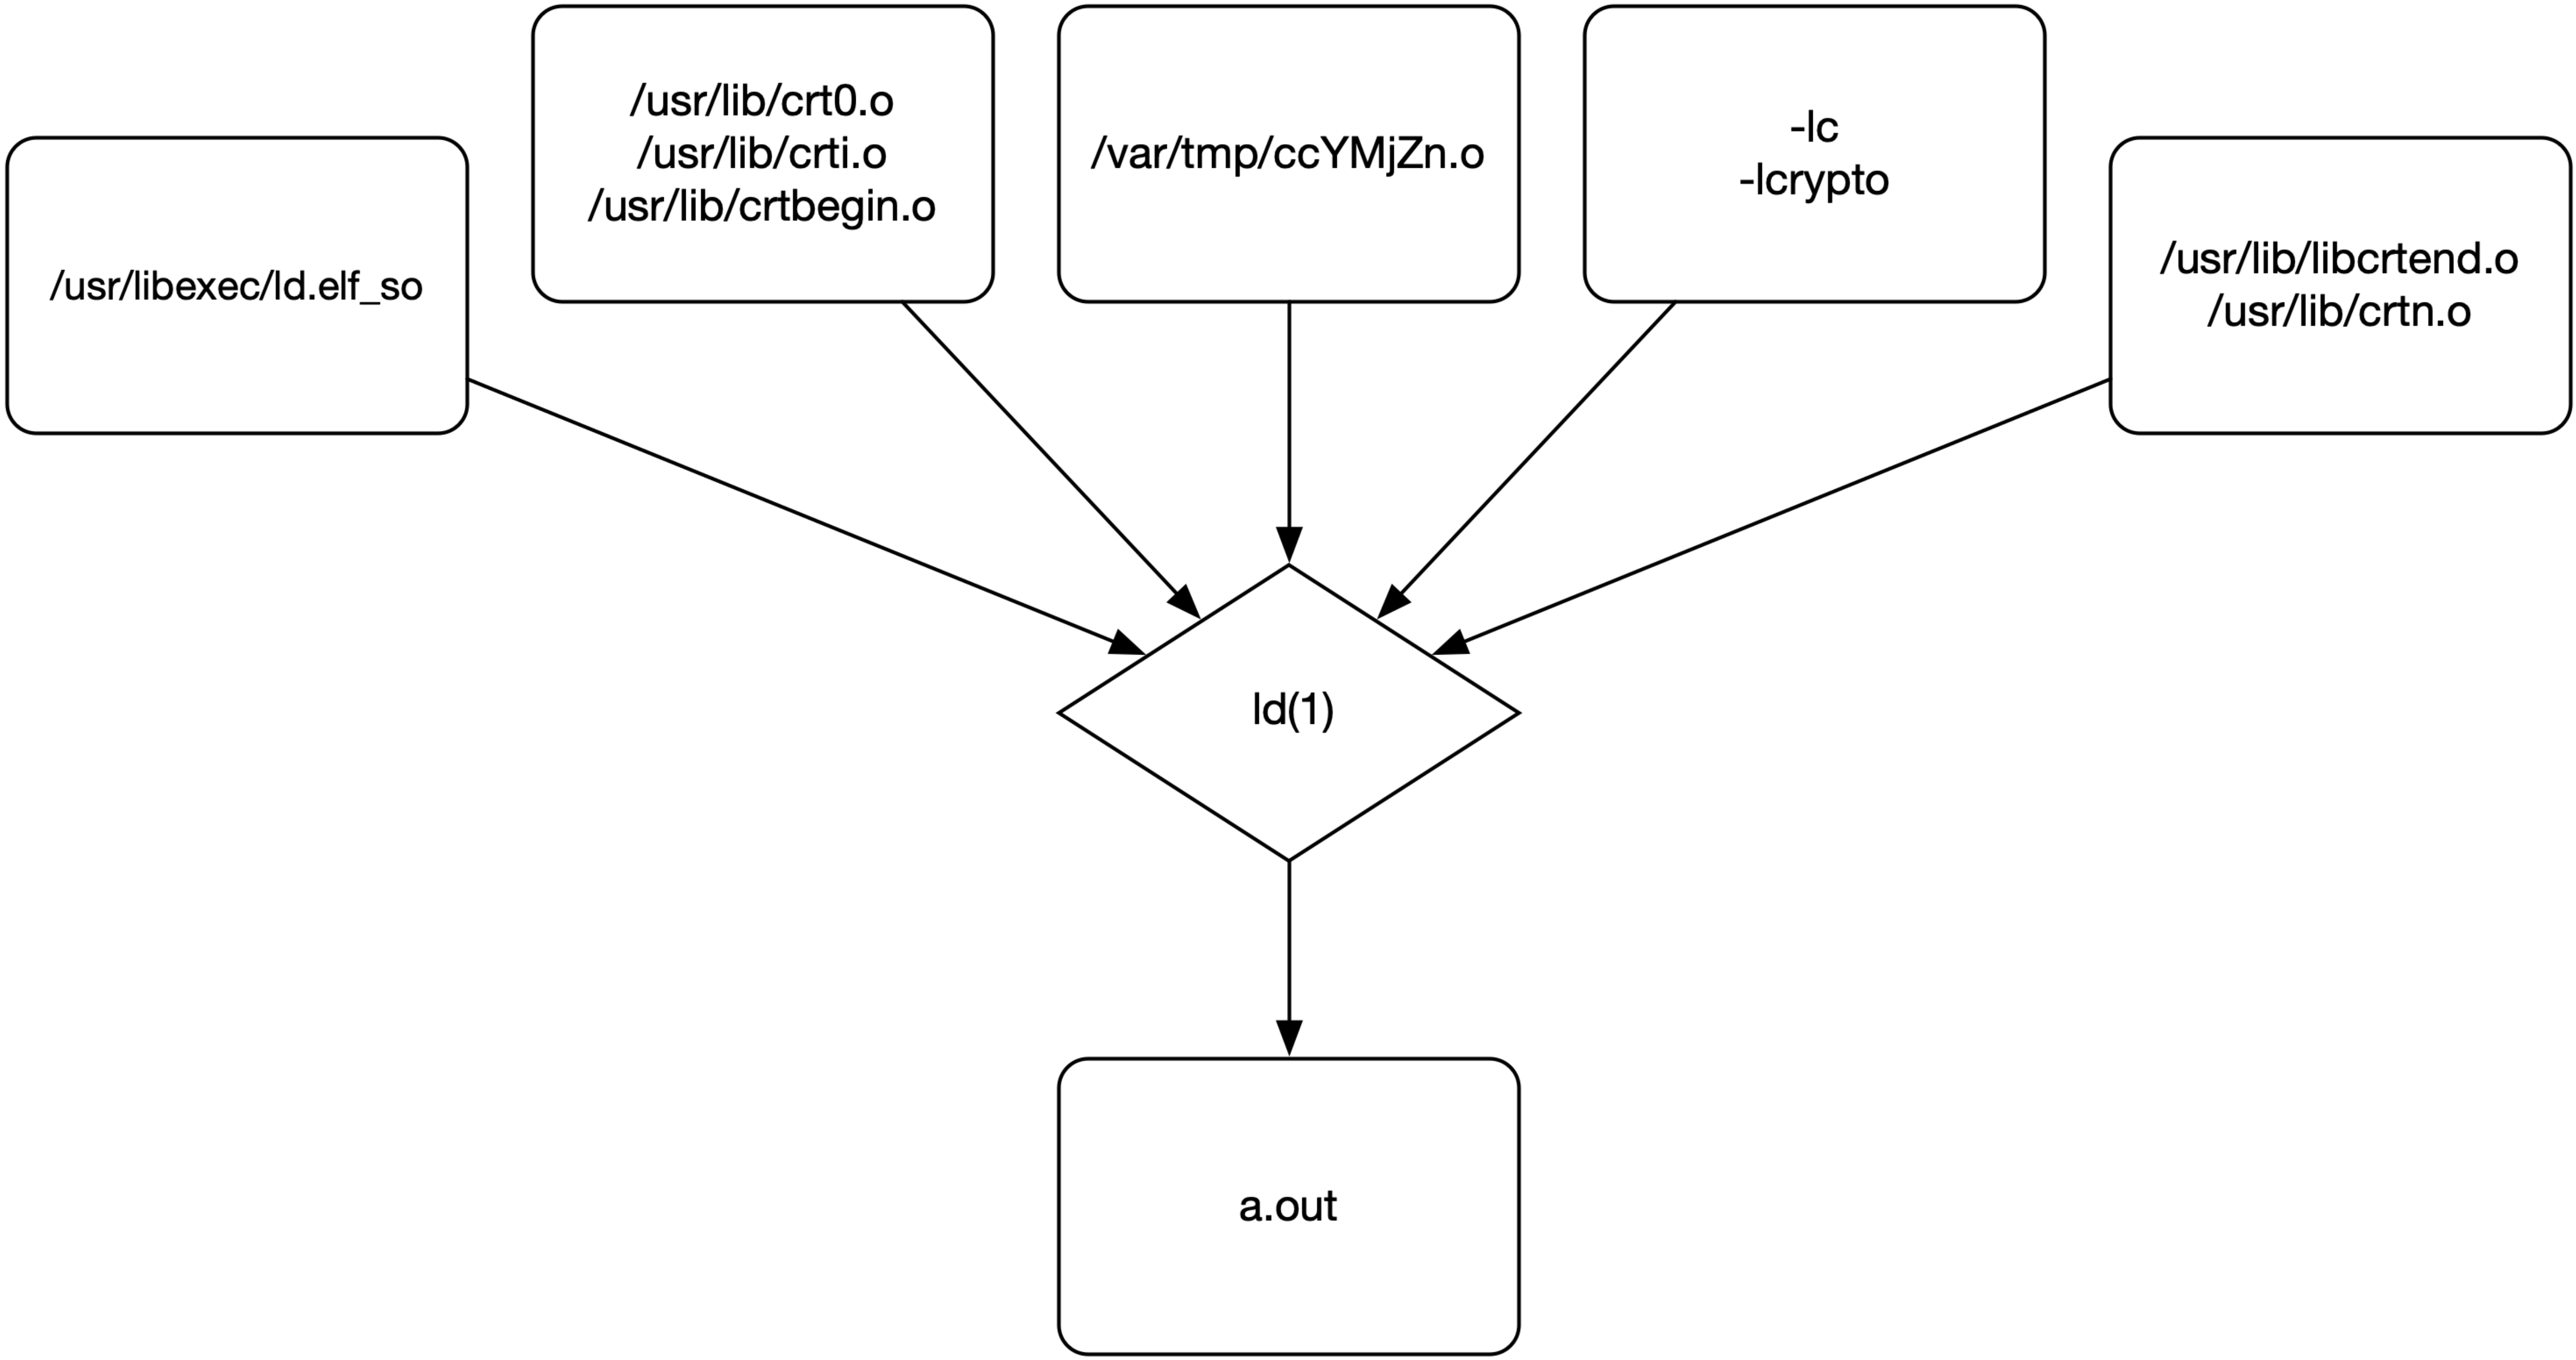












```

0x00000000000000190 0x00000000000000190 RW 0x8
NOTE 0x00000000000000158 0x000000000000400158 0x000000000000400158
0x000000000000002c 0x000000000000002c R 0x4

```

Section to Segment mapping:

Segment Sections...

00

[01 .note.netbsd.ident .note.netbsd.pax .hash .dynsym .dynstr .rela.dyn .r]

jschauma@apue\$ ktrace -i ./a.out foo

[2] Segmentation fault (core dumped) ktrace -i ./a.out foo

[jschauma@apue\$ kdump

1493 1 ktrace EMUL "netbsd"

1493 1 ktrace CALL execve(0x7f7fff1378ba,0x7f7fff137318,0x7f7fff137330

)

1493 1 ktrace NAMI "/home/jschauma/11/./a.out"

1493 1 a.out EMUL "netbsd"

1493 1 a.out RET execve JUSTRETURN

1493 1 a.out PSIG SIGSEGV SIG_DFL: code=SEGV_MAPERR, addr=0x0, trap=6

)


1493 1 a.out NAMI "a.out.core"

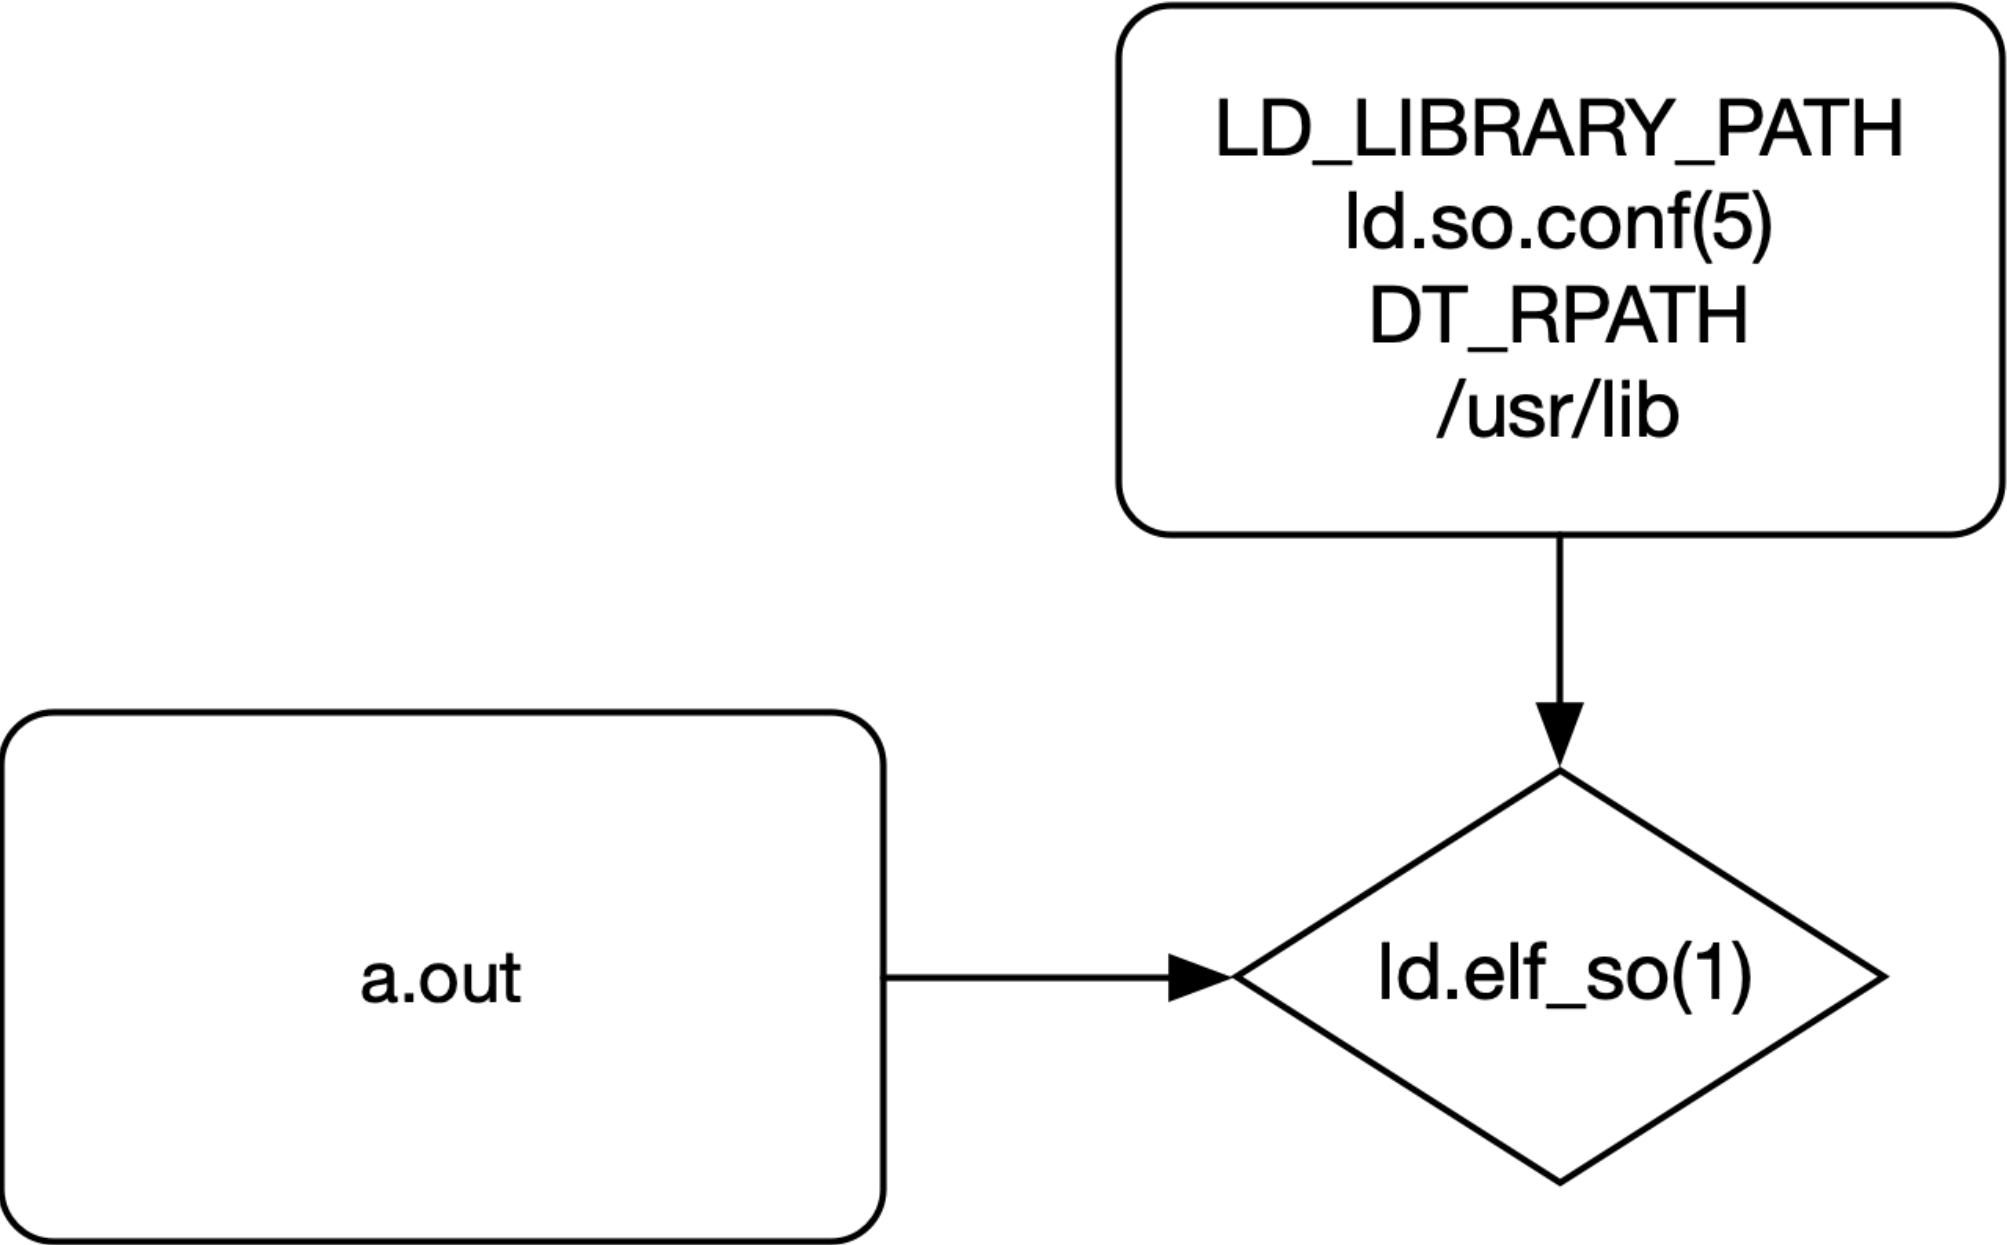
[jschauma@apue\$ ld -dynamic-linker /usr/libexec/ld.elf_so /usr/libexec/ld.elf_so

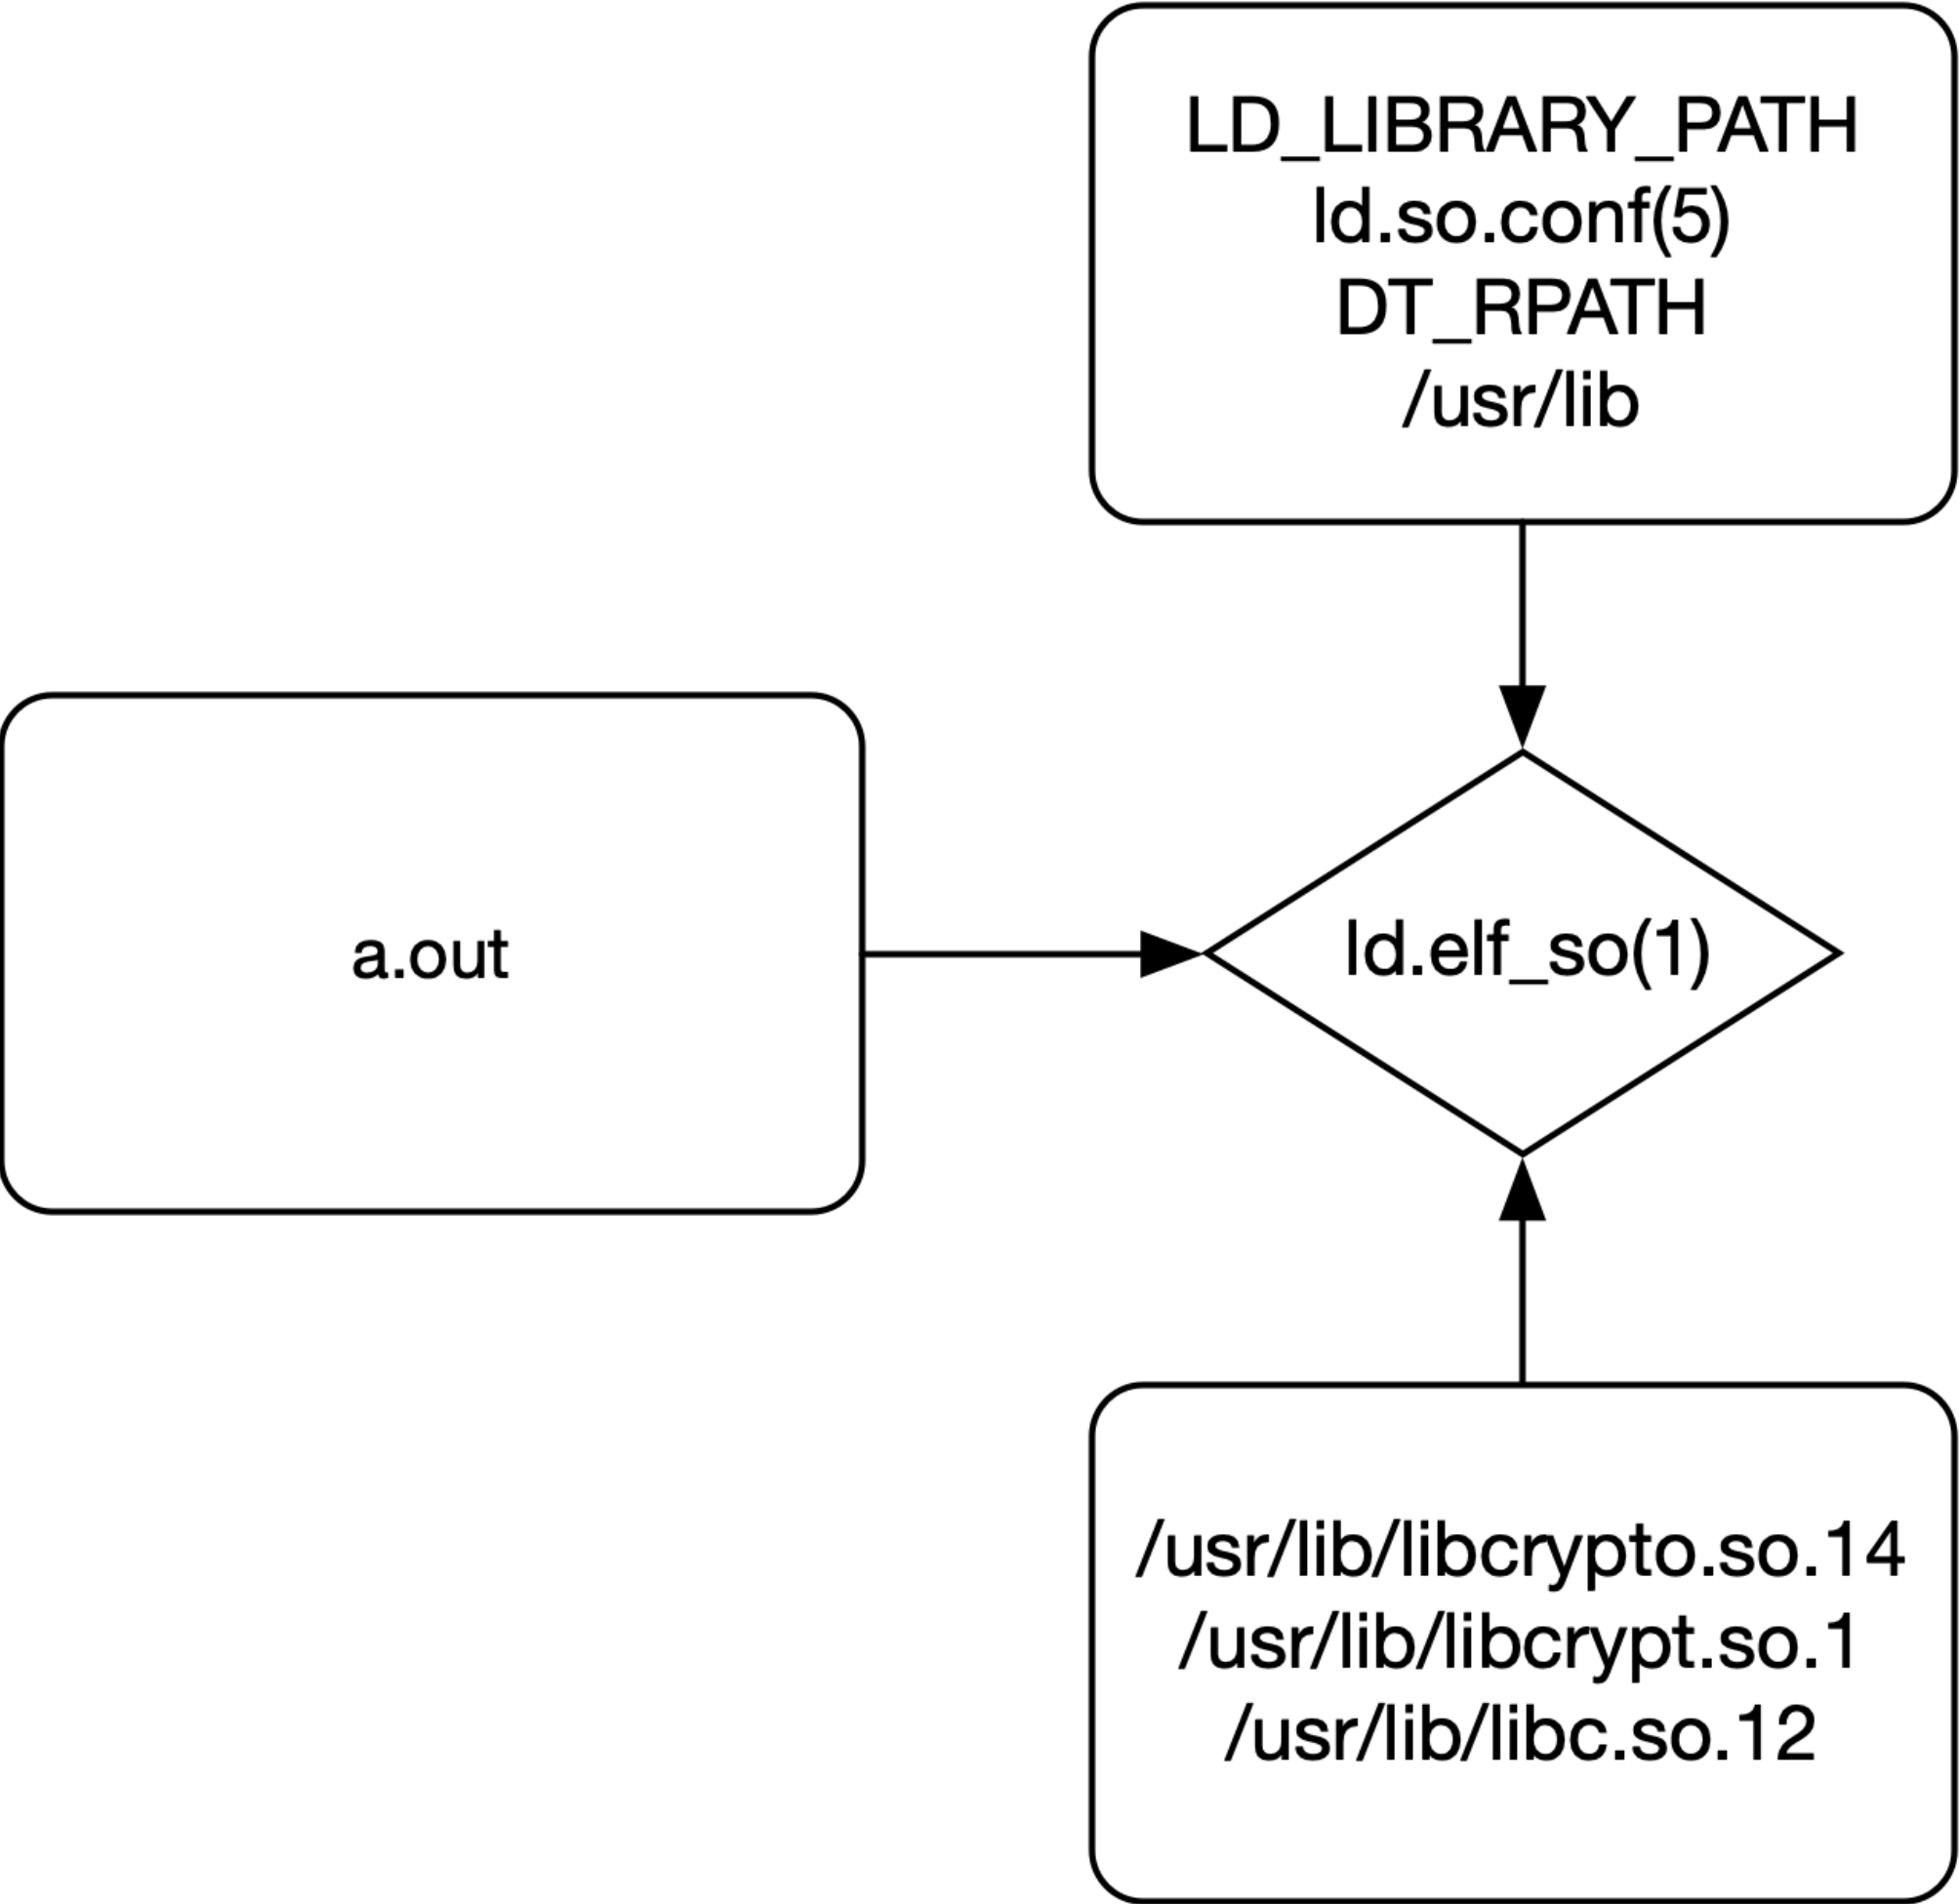
/usr/lib/crt0.o /usr/lib/crti.o crypt.o -lc -lcrypto /usr/lib/crtn.o

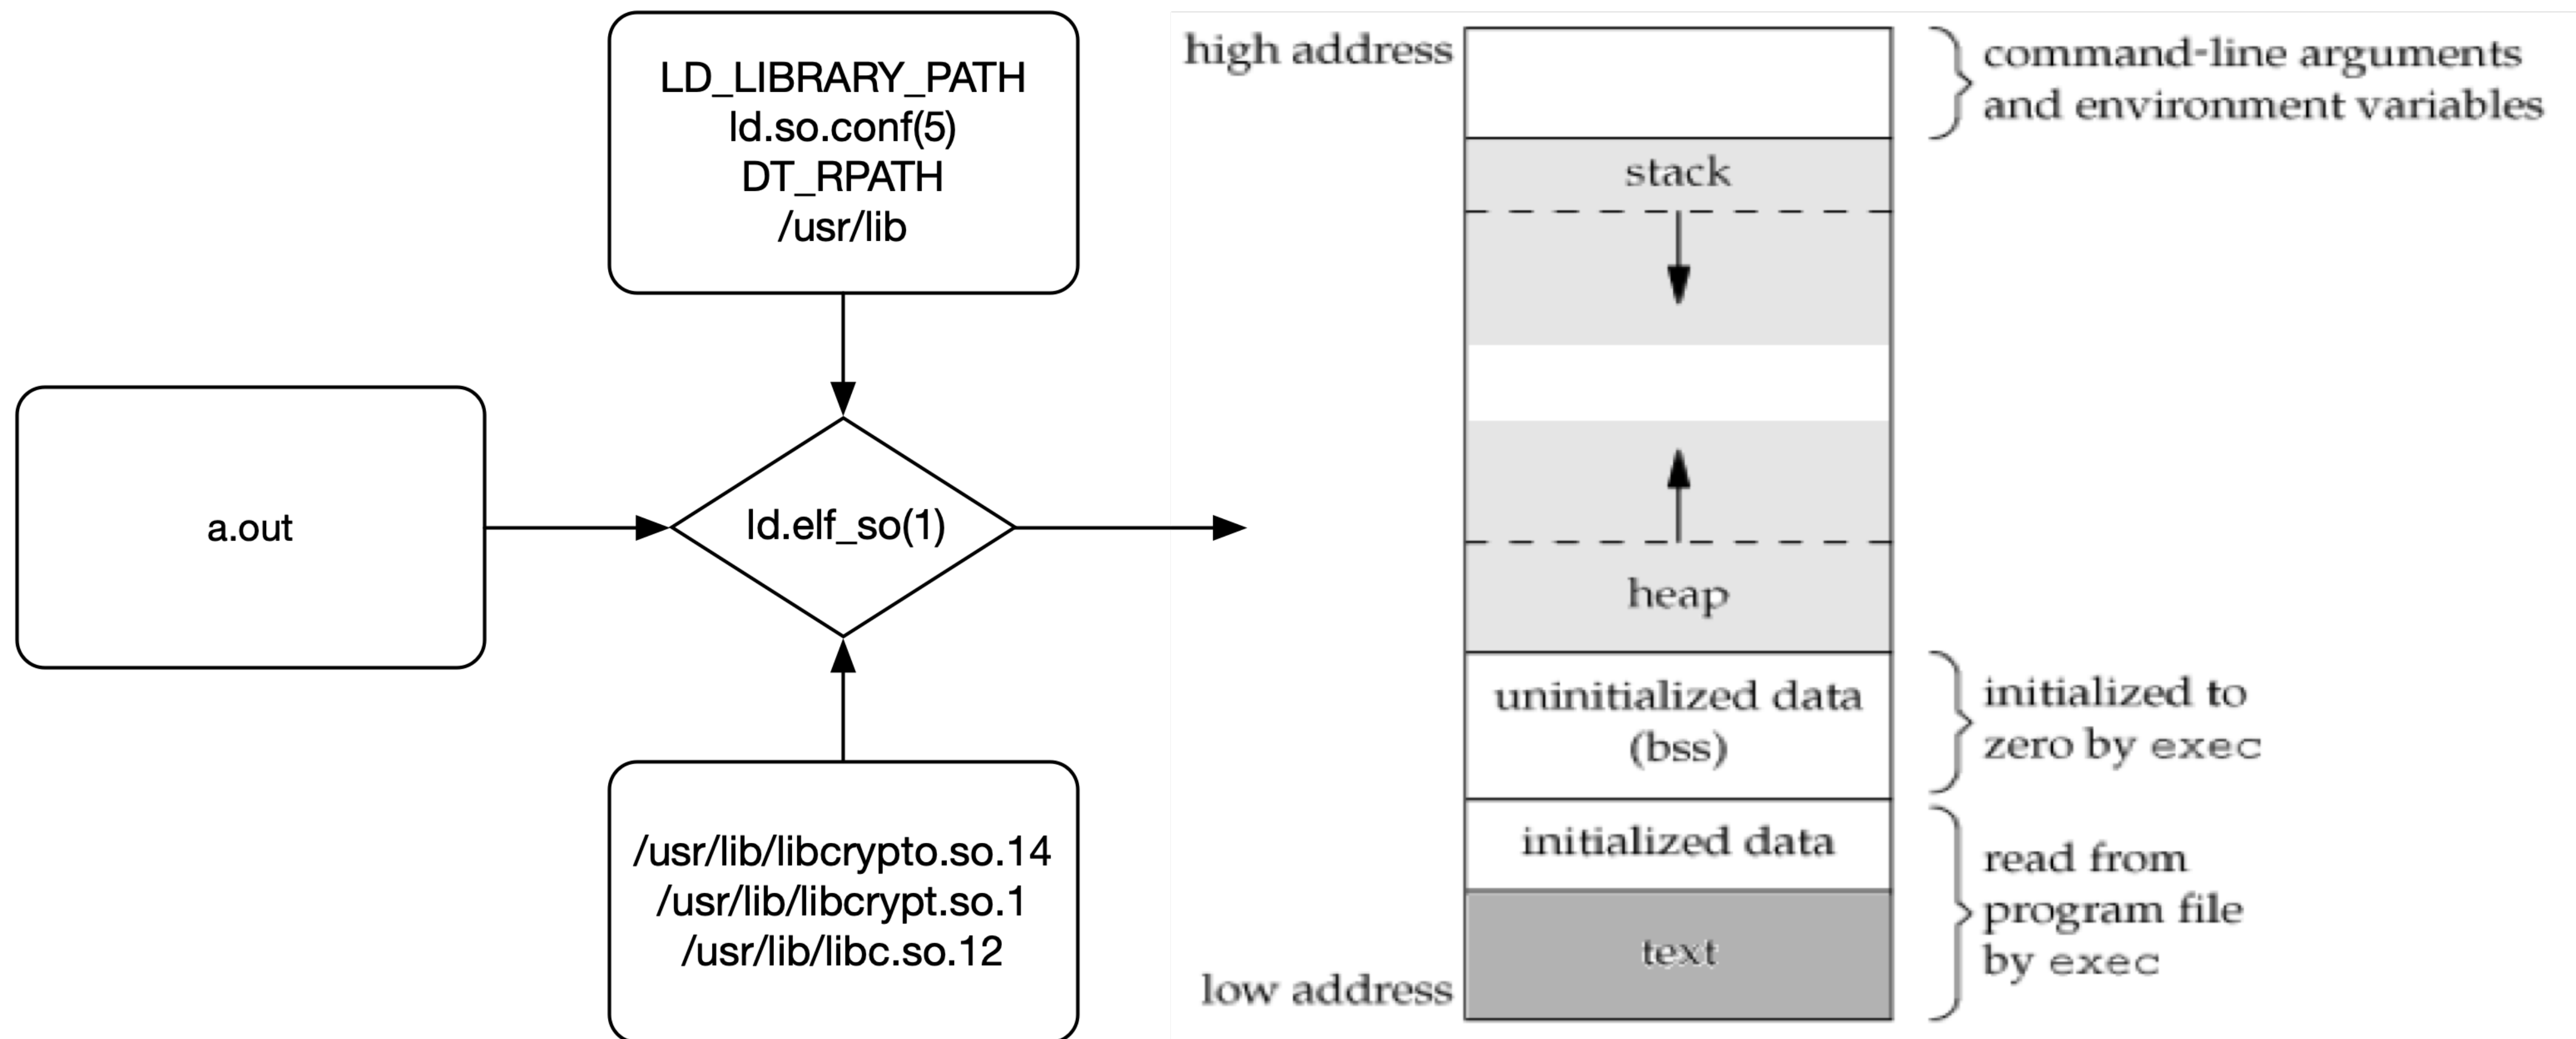
[jschauma@apue\$./a.out foo

\$1\$\$n1rTiFE0nRifwV/43bVon/









Summary

- A linker takes multiple *object files*, resolves symbols to e.g., addresses in *libraries* (possibly relocating them in the process), and produces an *executable*.
- A loader copies a program into main memory, possibly invoking the *dynamic linker* or *run-time link editor* to find the right libraries, resolve addresses of symbols, and relocate them.
- On some systems, the run-time link-editor is itself an *executable*, allowing you to invoke it directly and passing another executable file as an argument. Try it out!

```
$ ld -no-dynamic-linker ... file.o ...
```

```
$ /lib64/ld-linux-x86-64.so ./a.out
```