

Advanced Programming in the UNIX Environment

Week 12, Segment 4: Asynchronous and Memory Mapped I/O

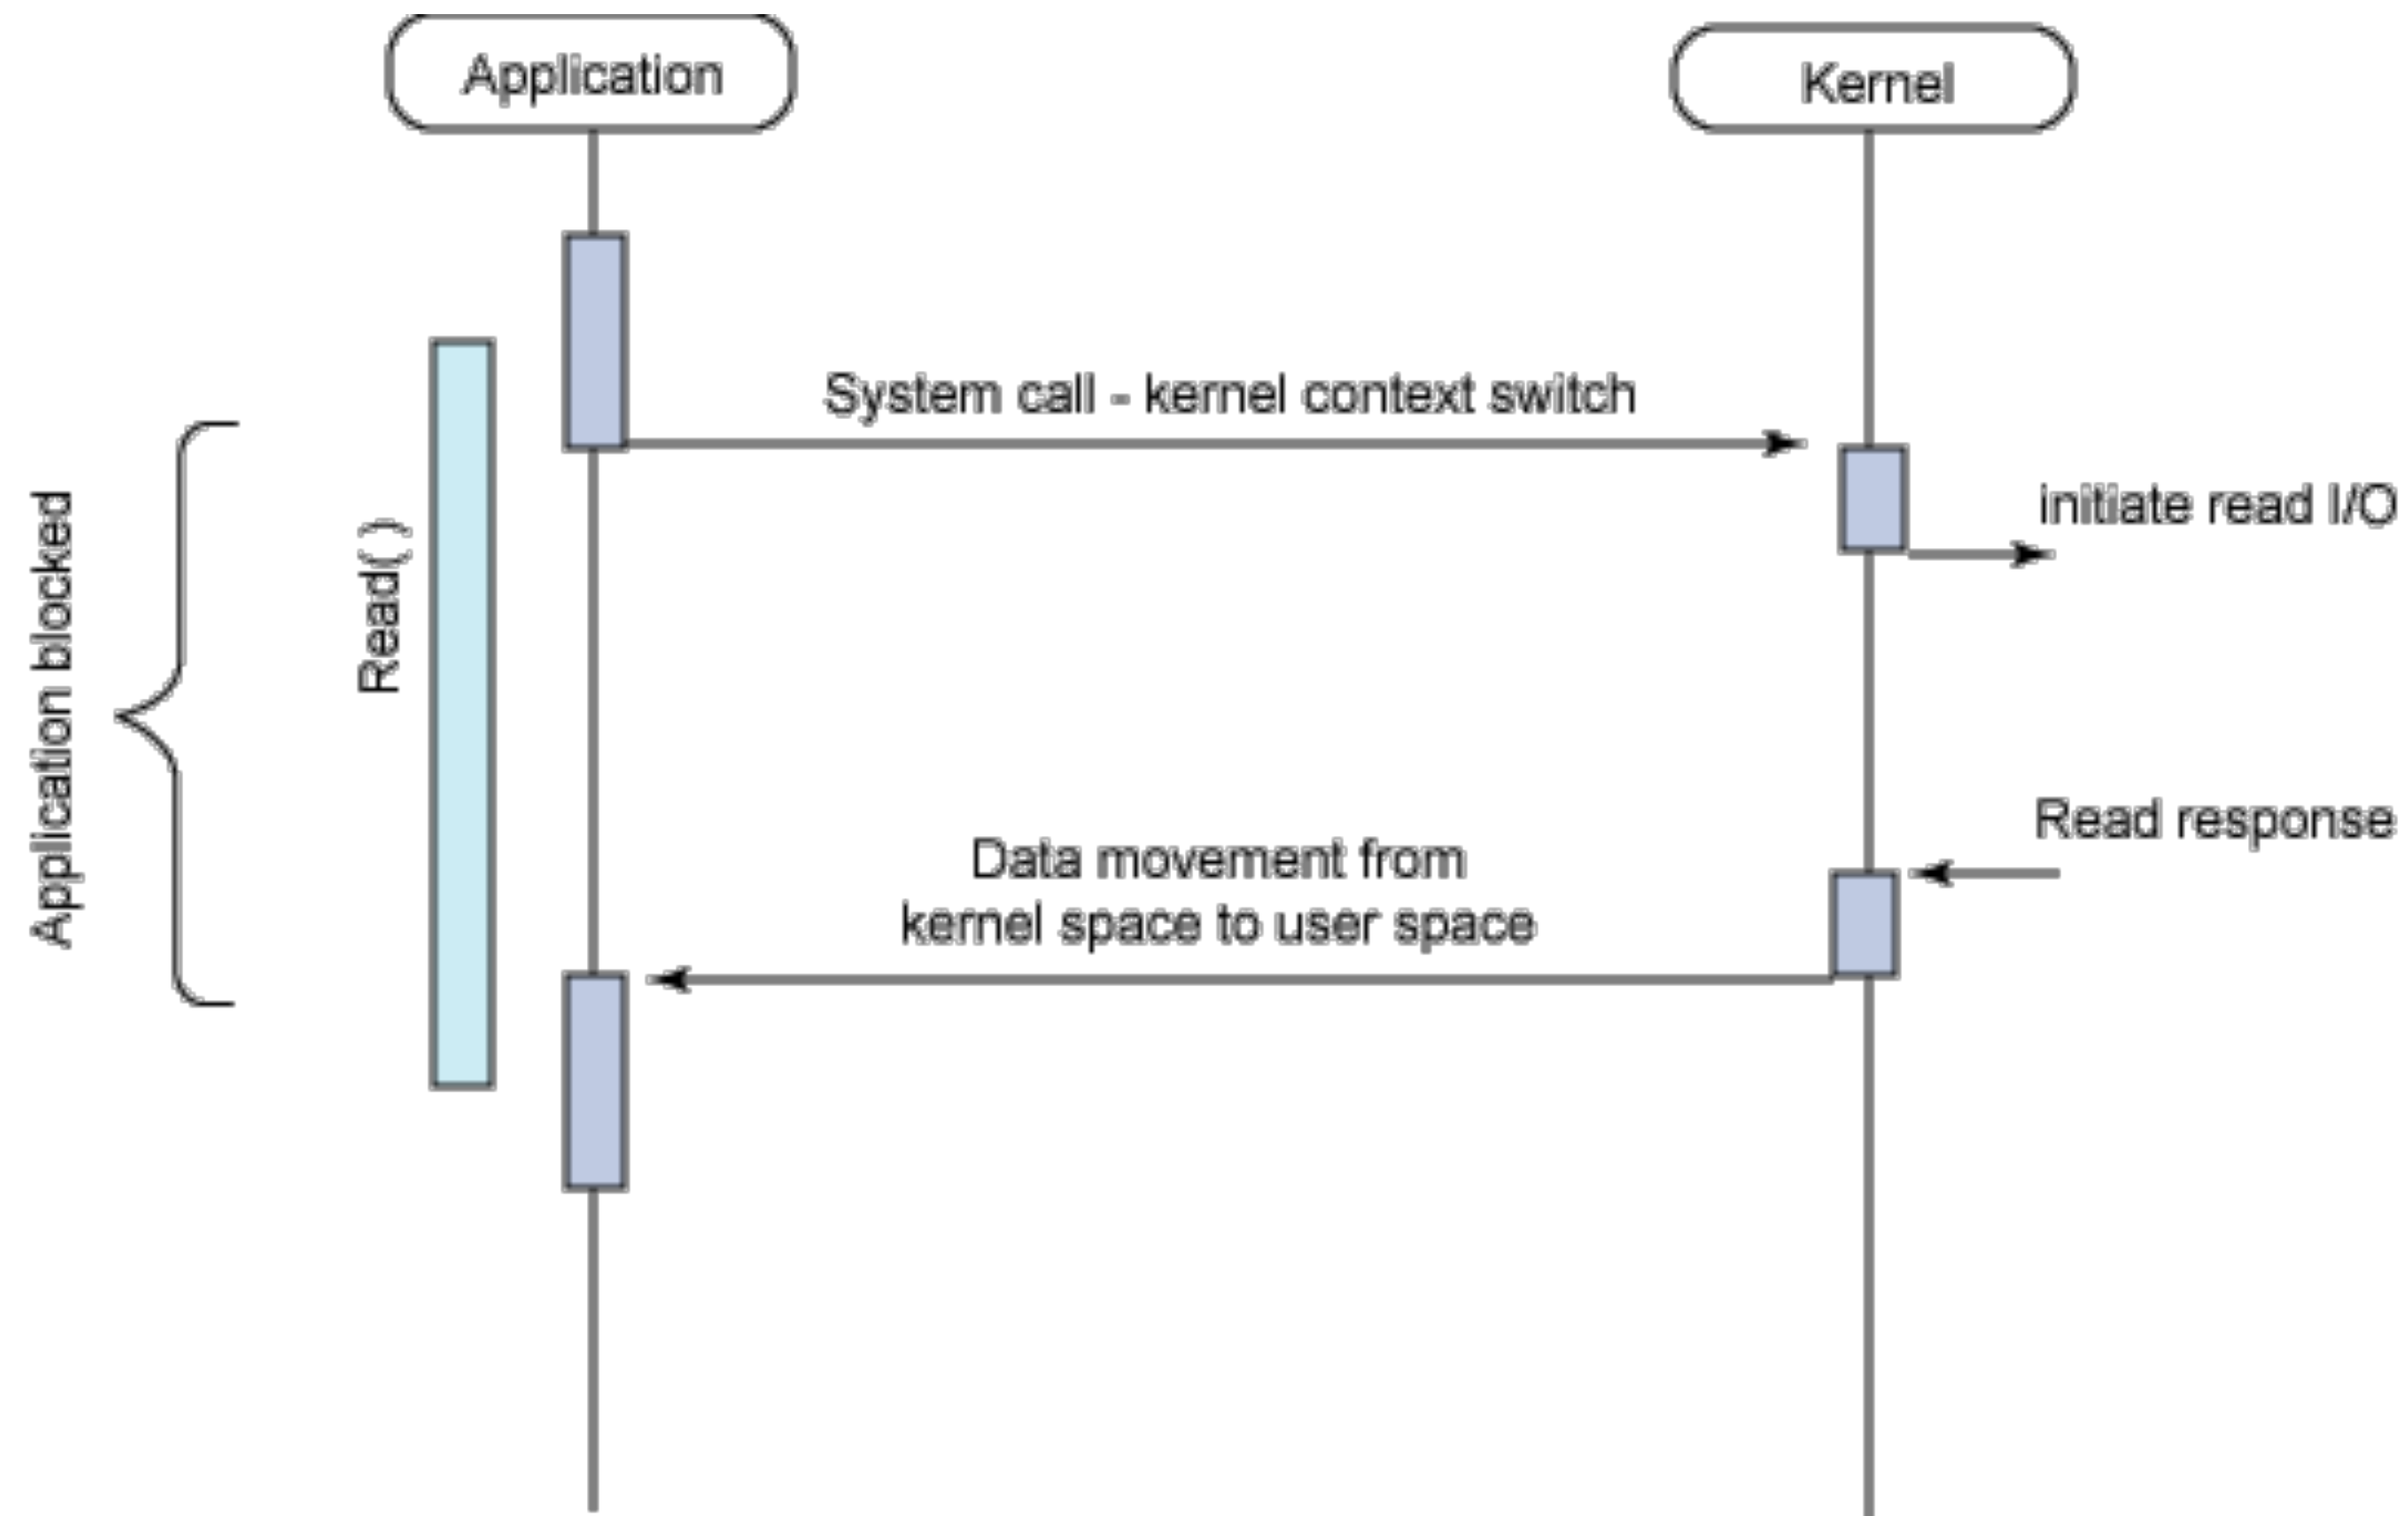
**Department of Computer Science
Stevens Institute of Technology**

Jan Schaumann

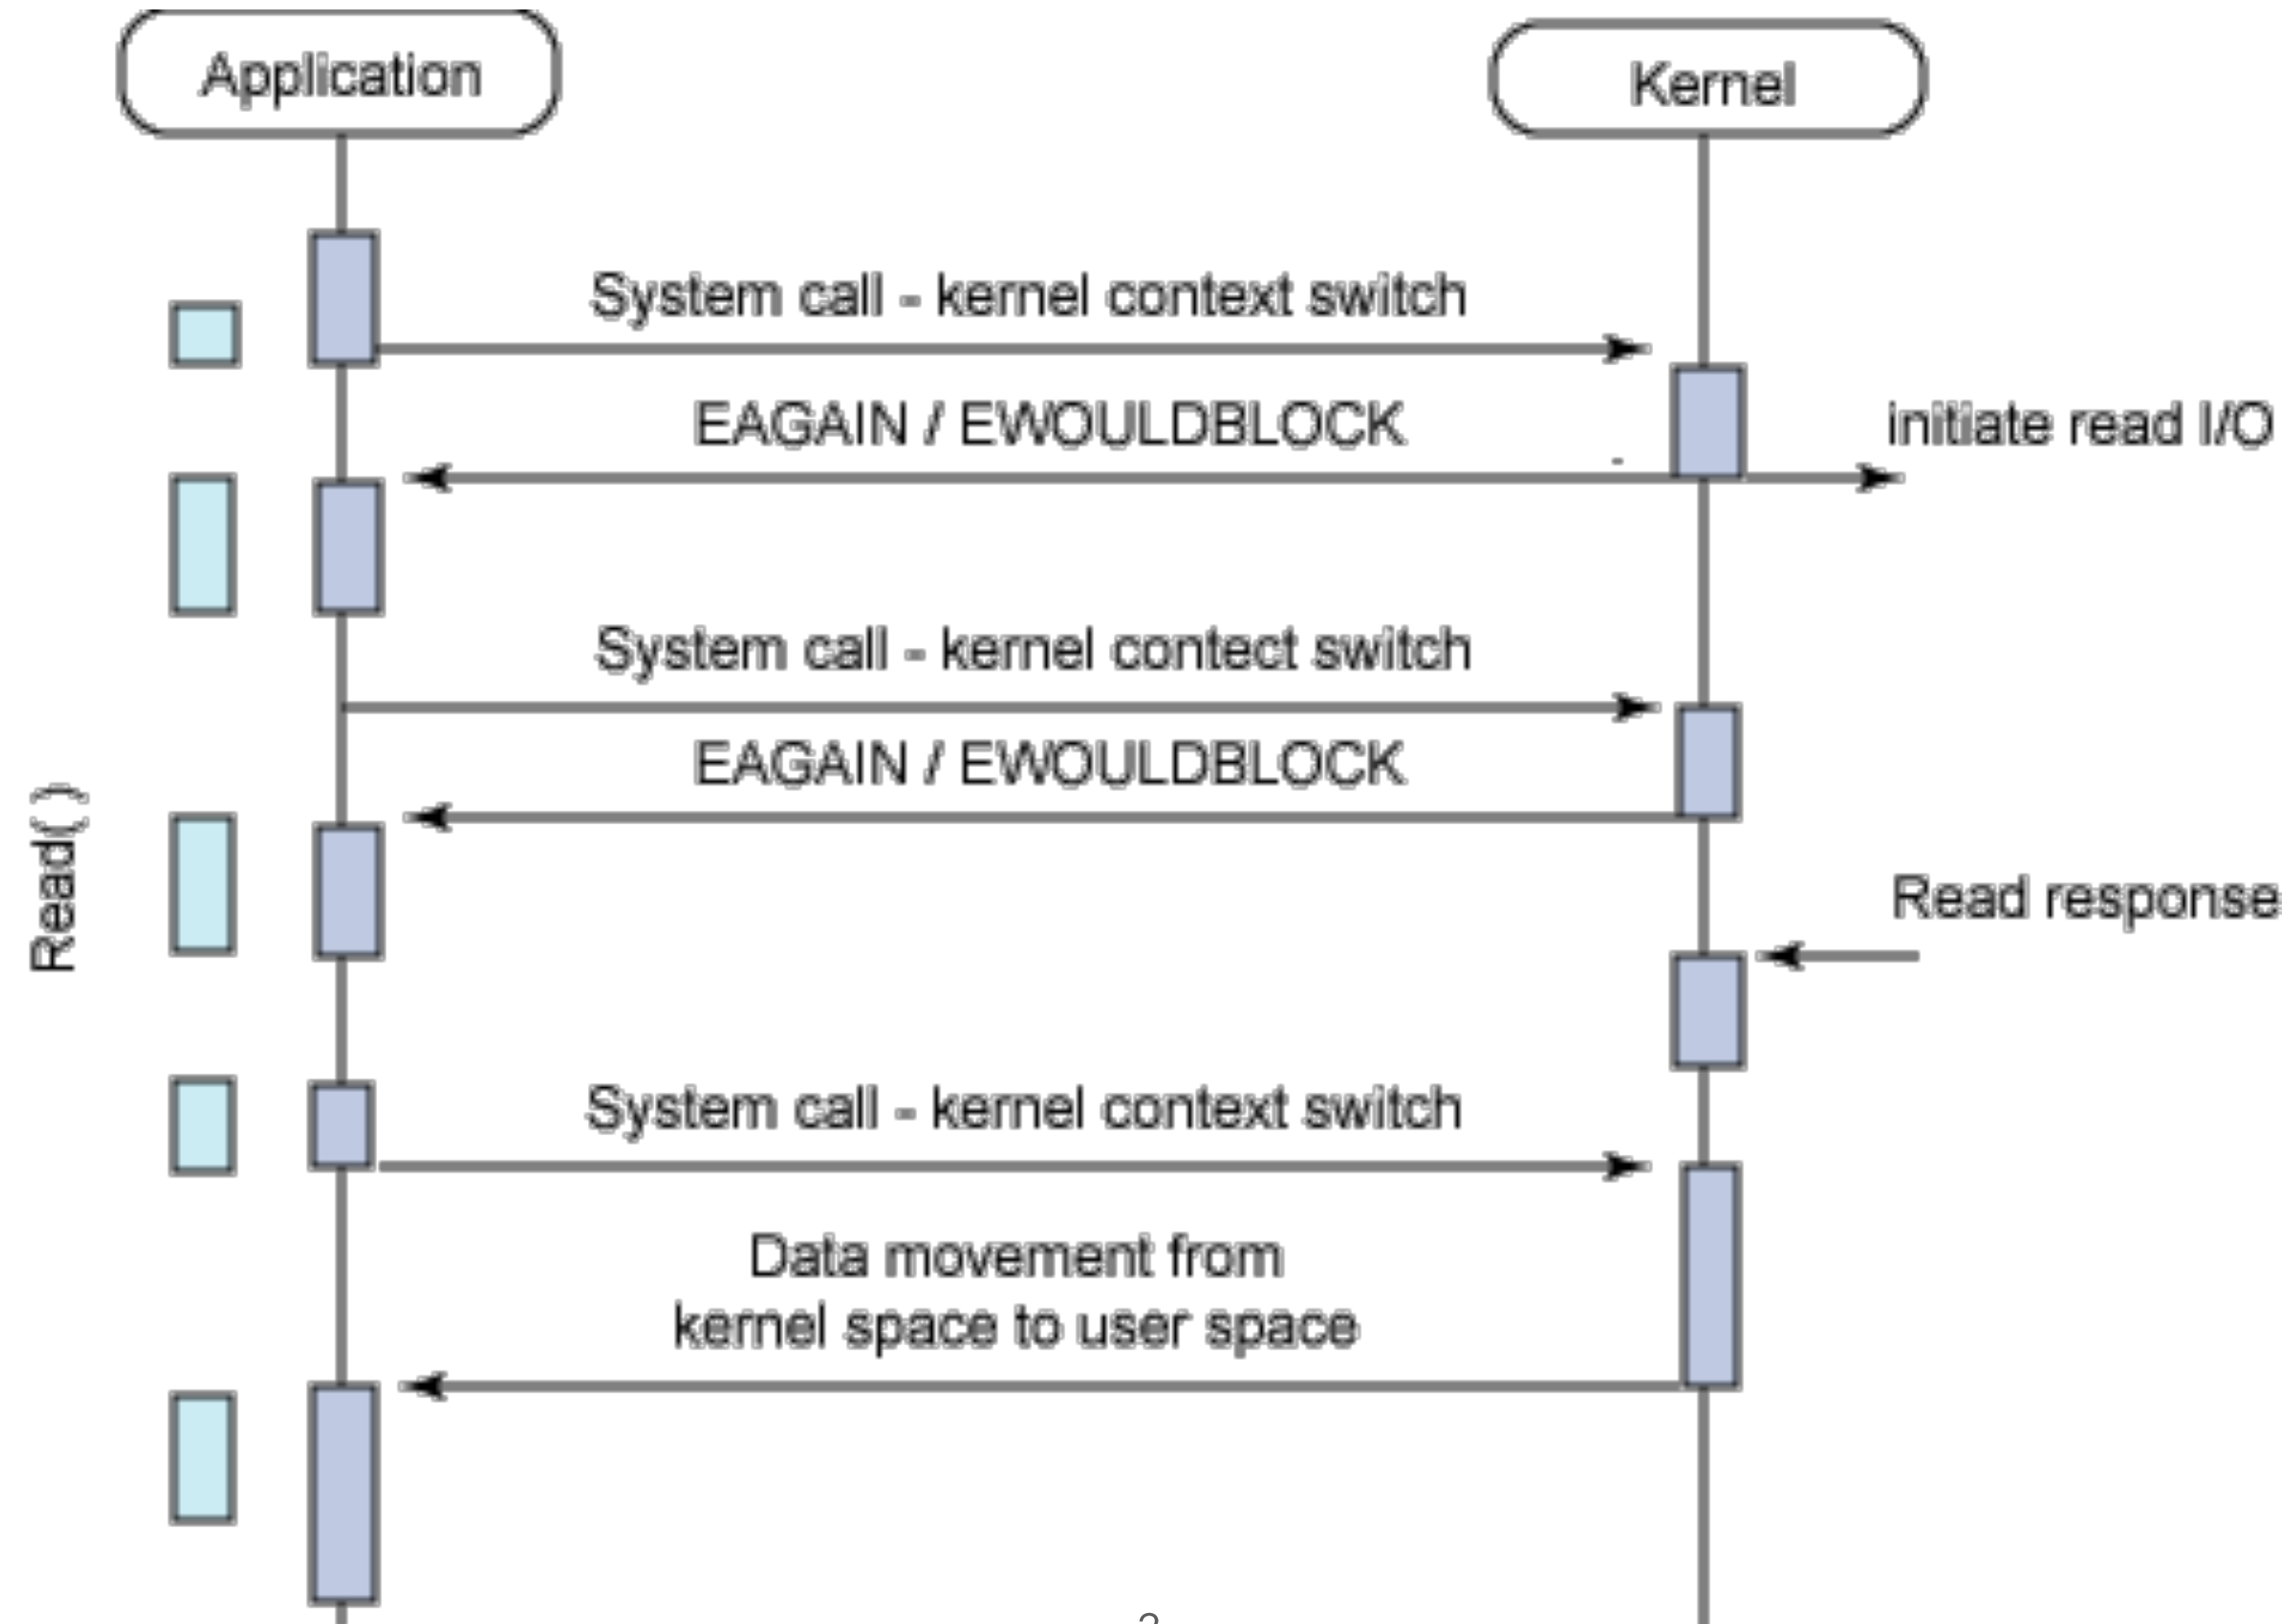
`jschauma@stevens.edu`

`https://stevens.netmeister.org/631/`

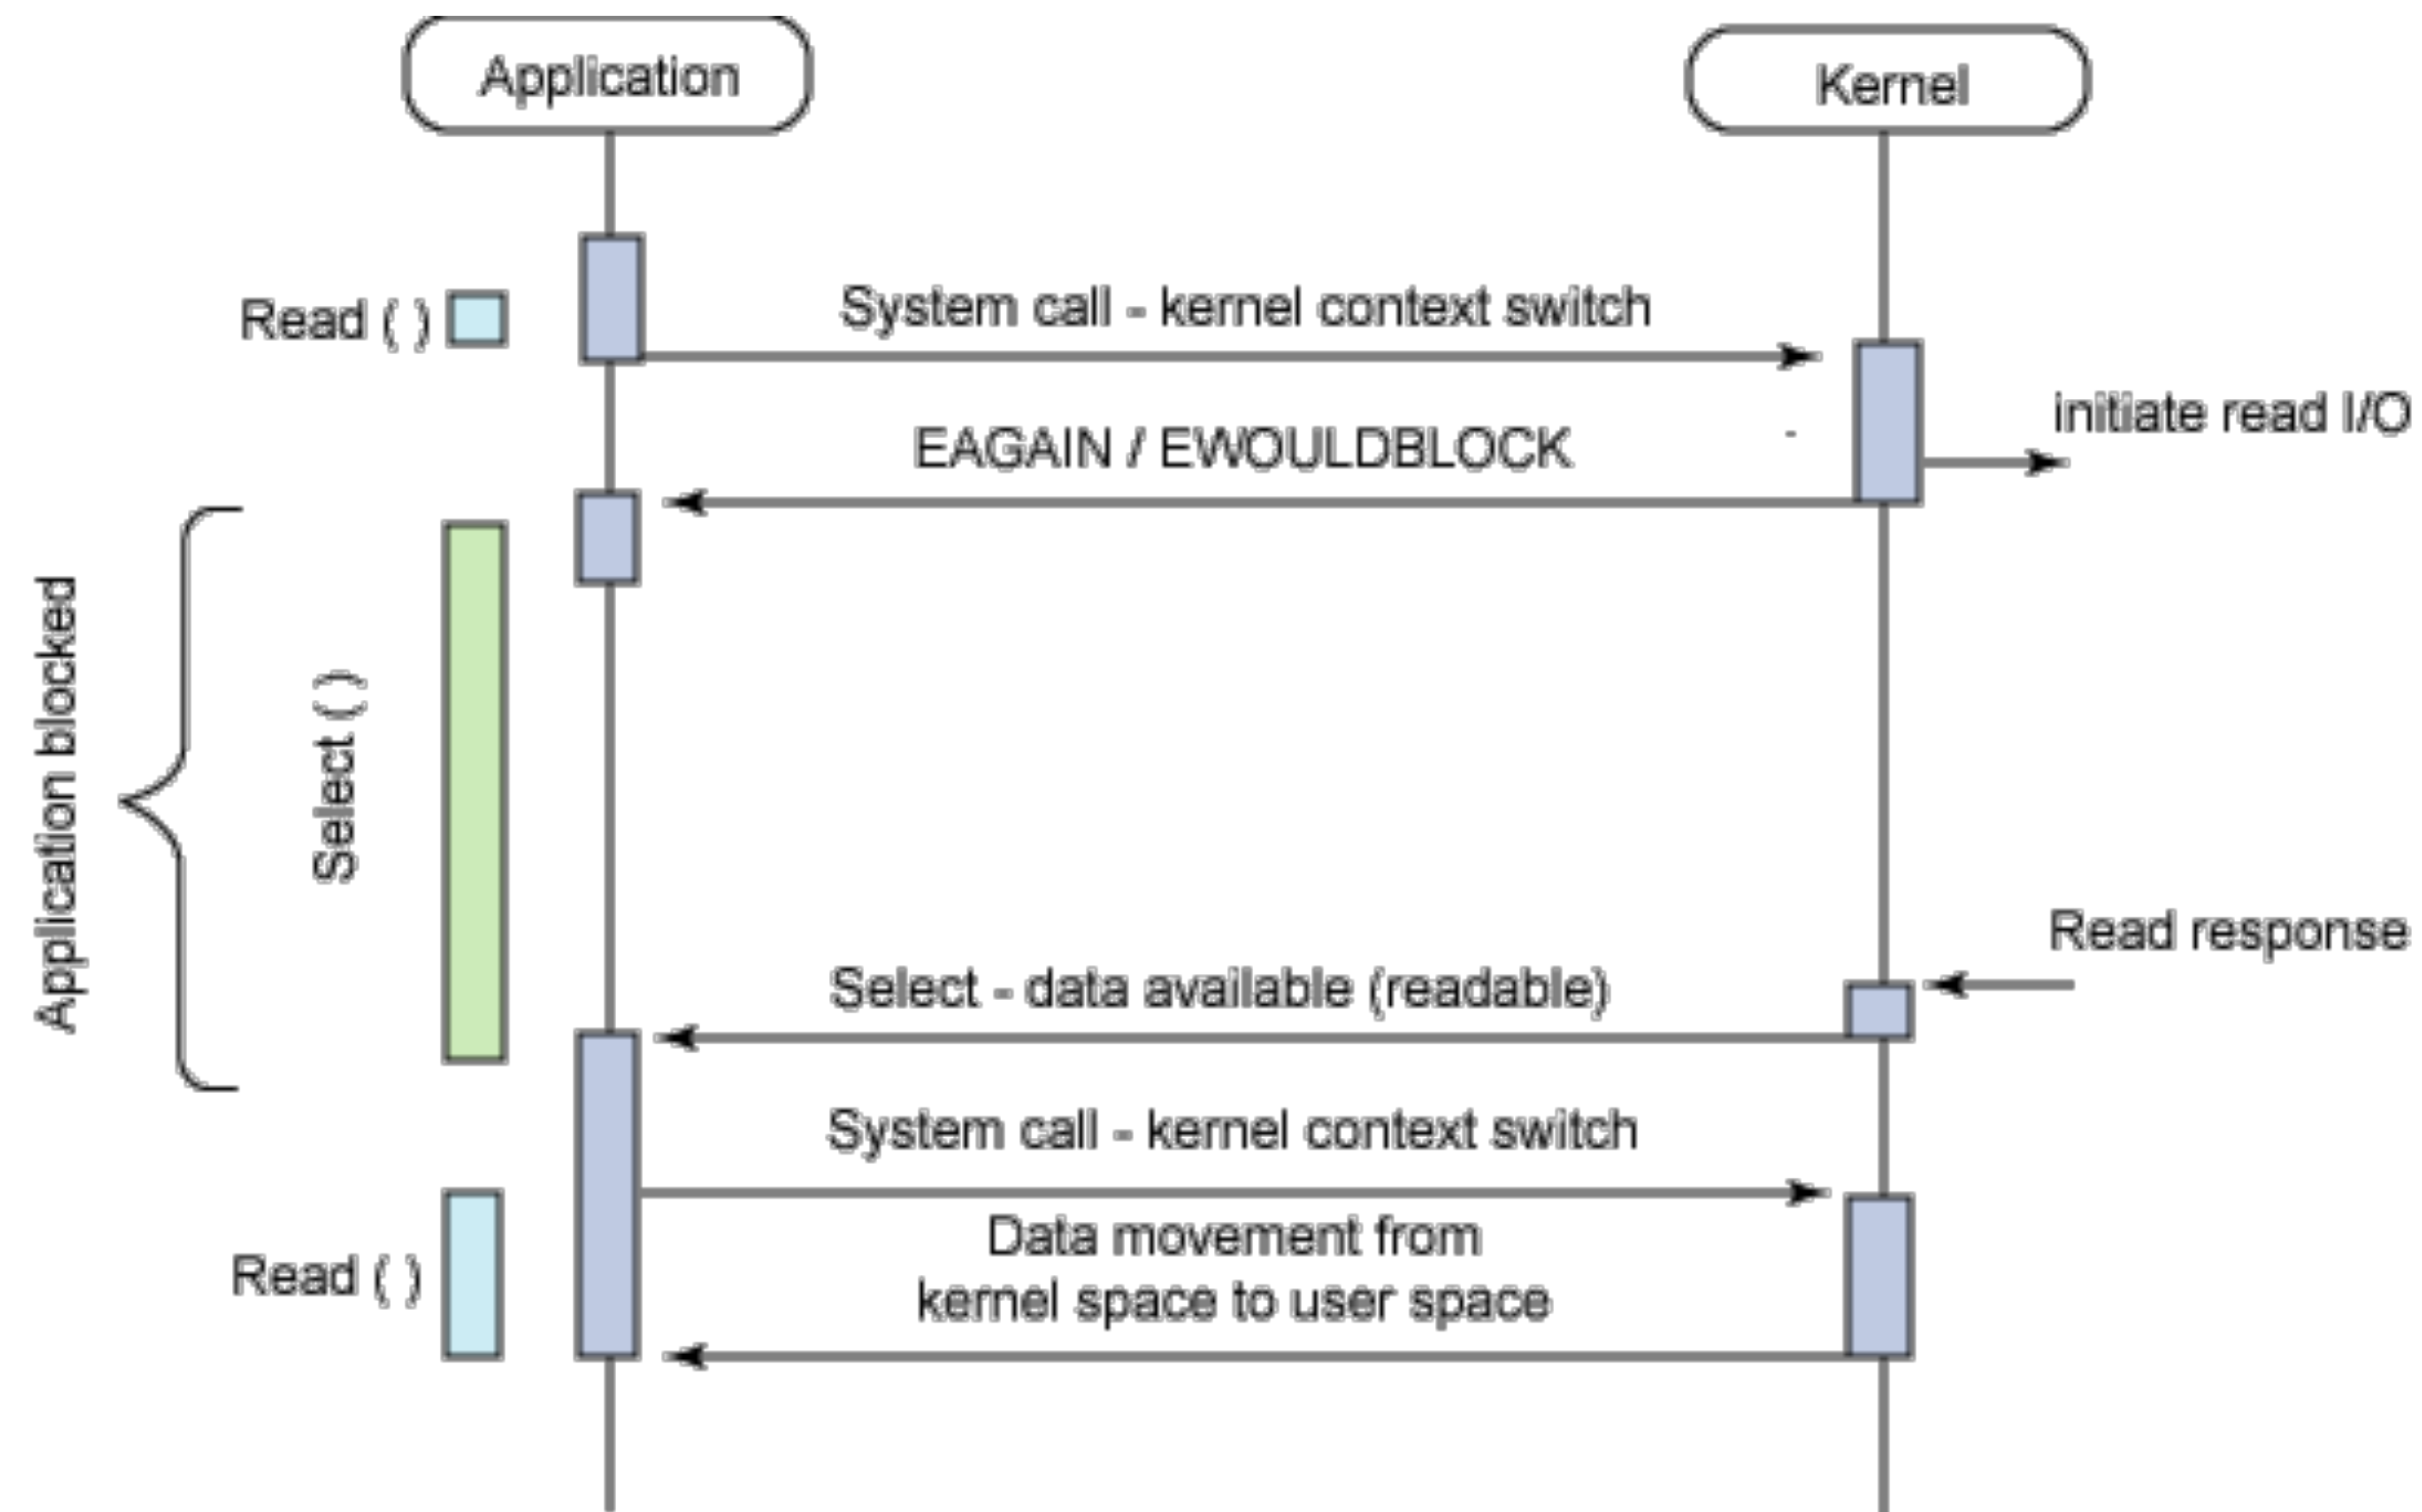
Synchronous, blocking I/O



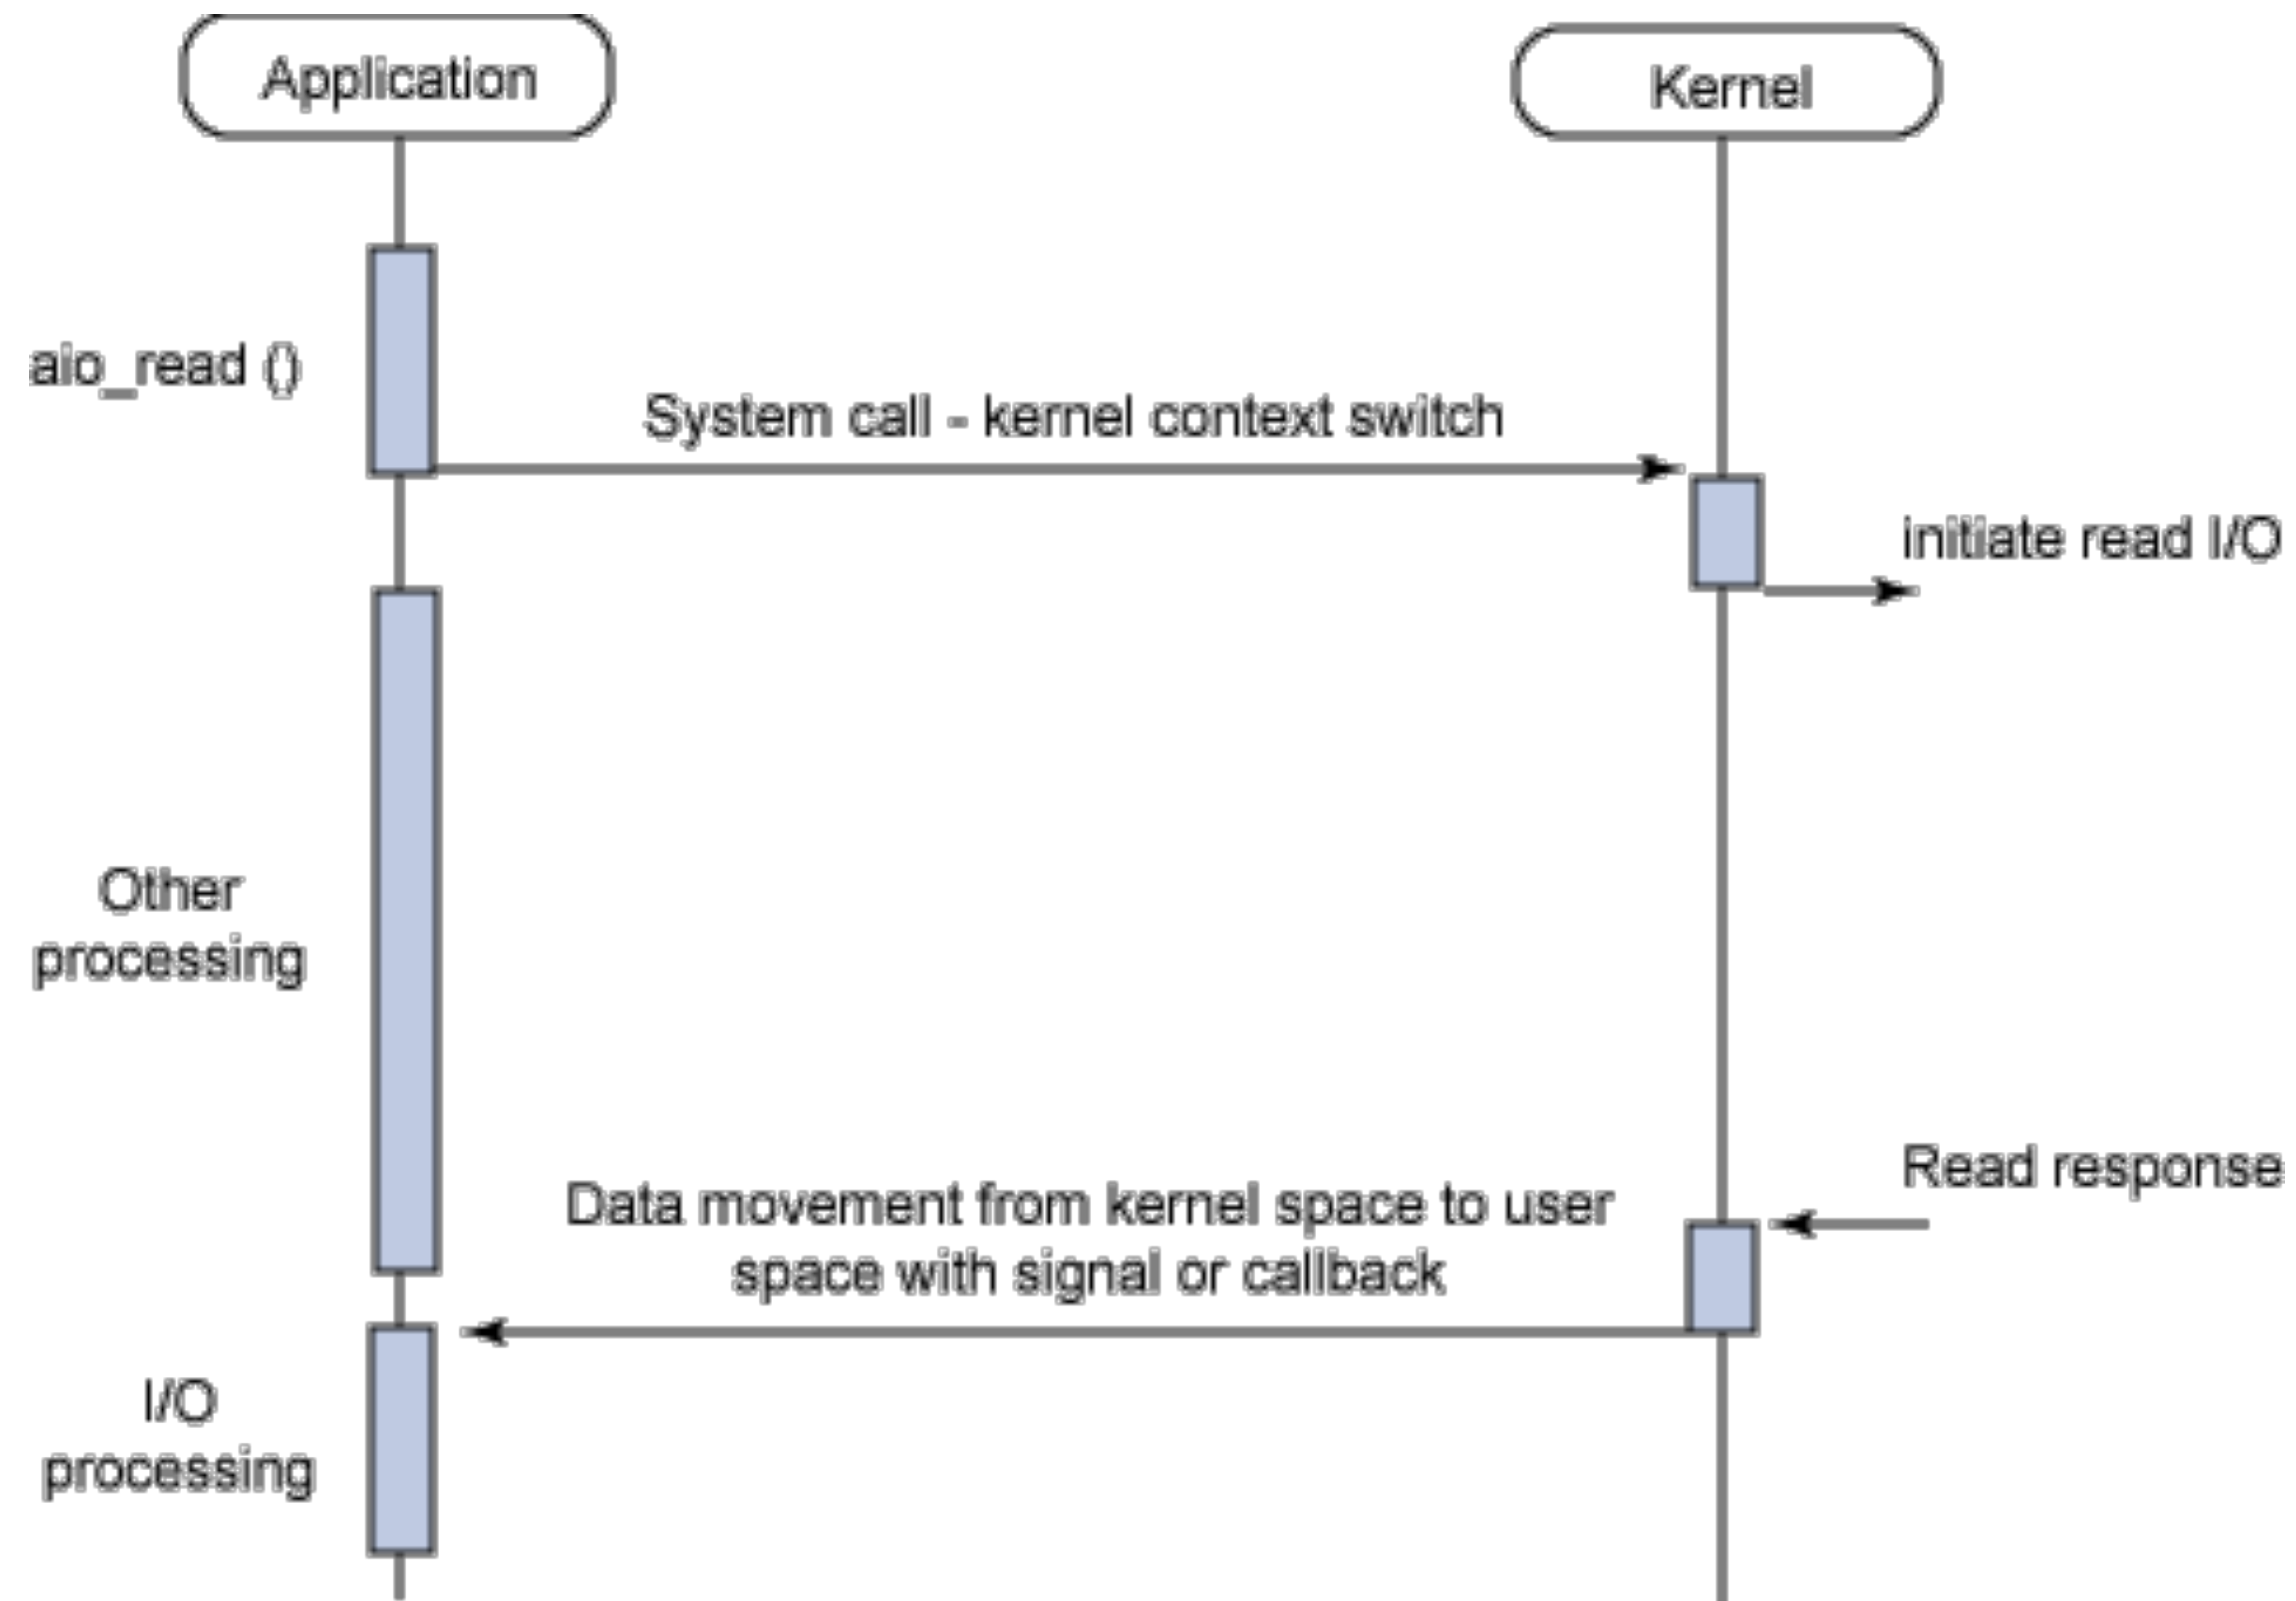
Asynchronous I/O



Asynchronous I/O



Asynchronous I/O



Asynchronous I/O

	Blocking	Non-blocking
Synchronous	read(2)/write(2)	read(2)/write(2) O_NONBLOCK
Asynchronous	I/O multiplexing (select(2)/poll(2))	AIO

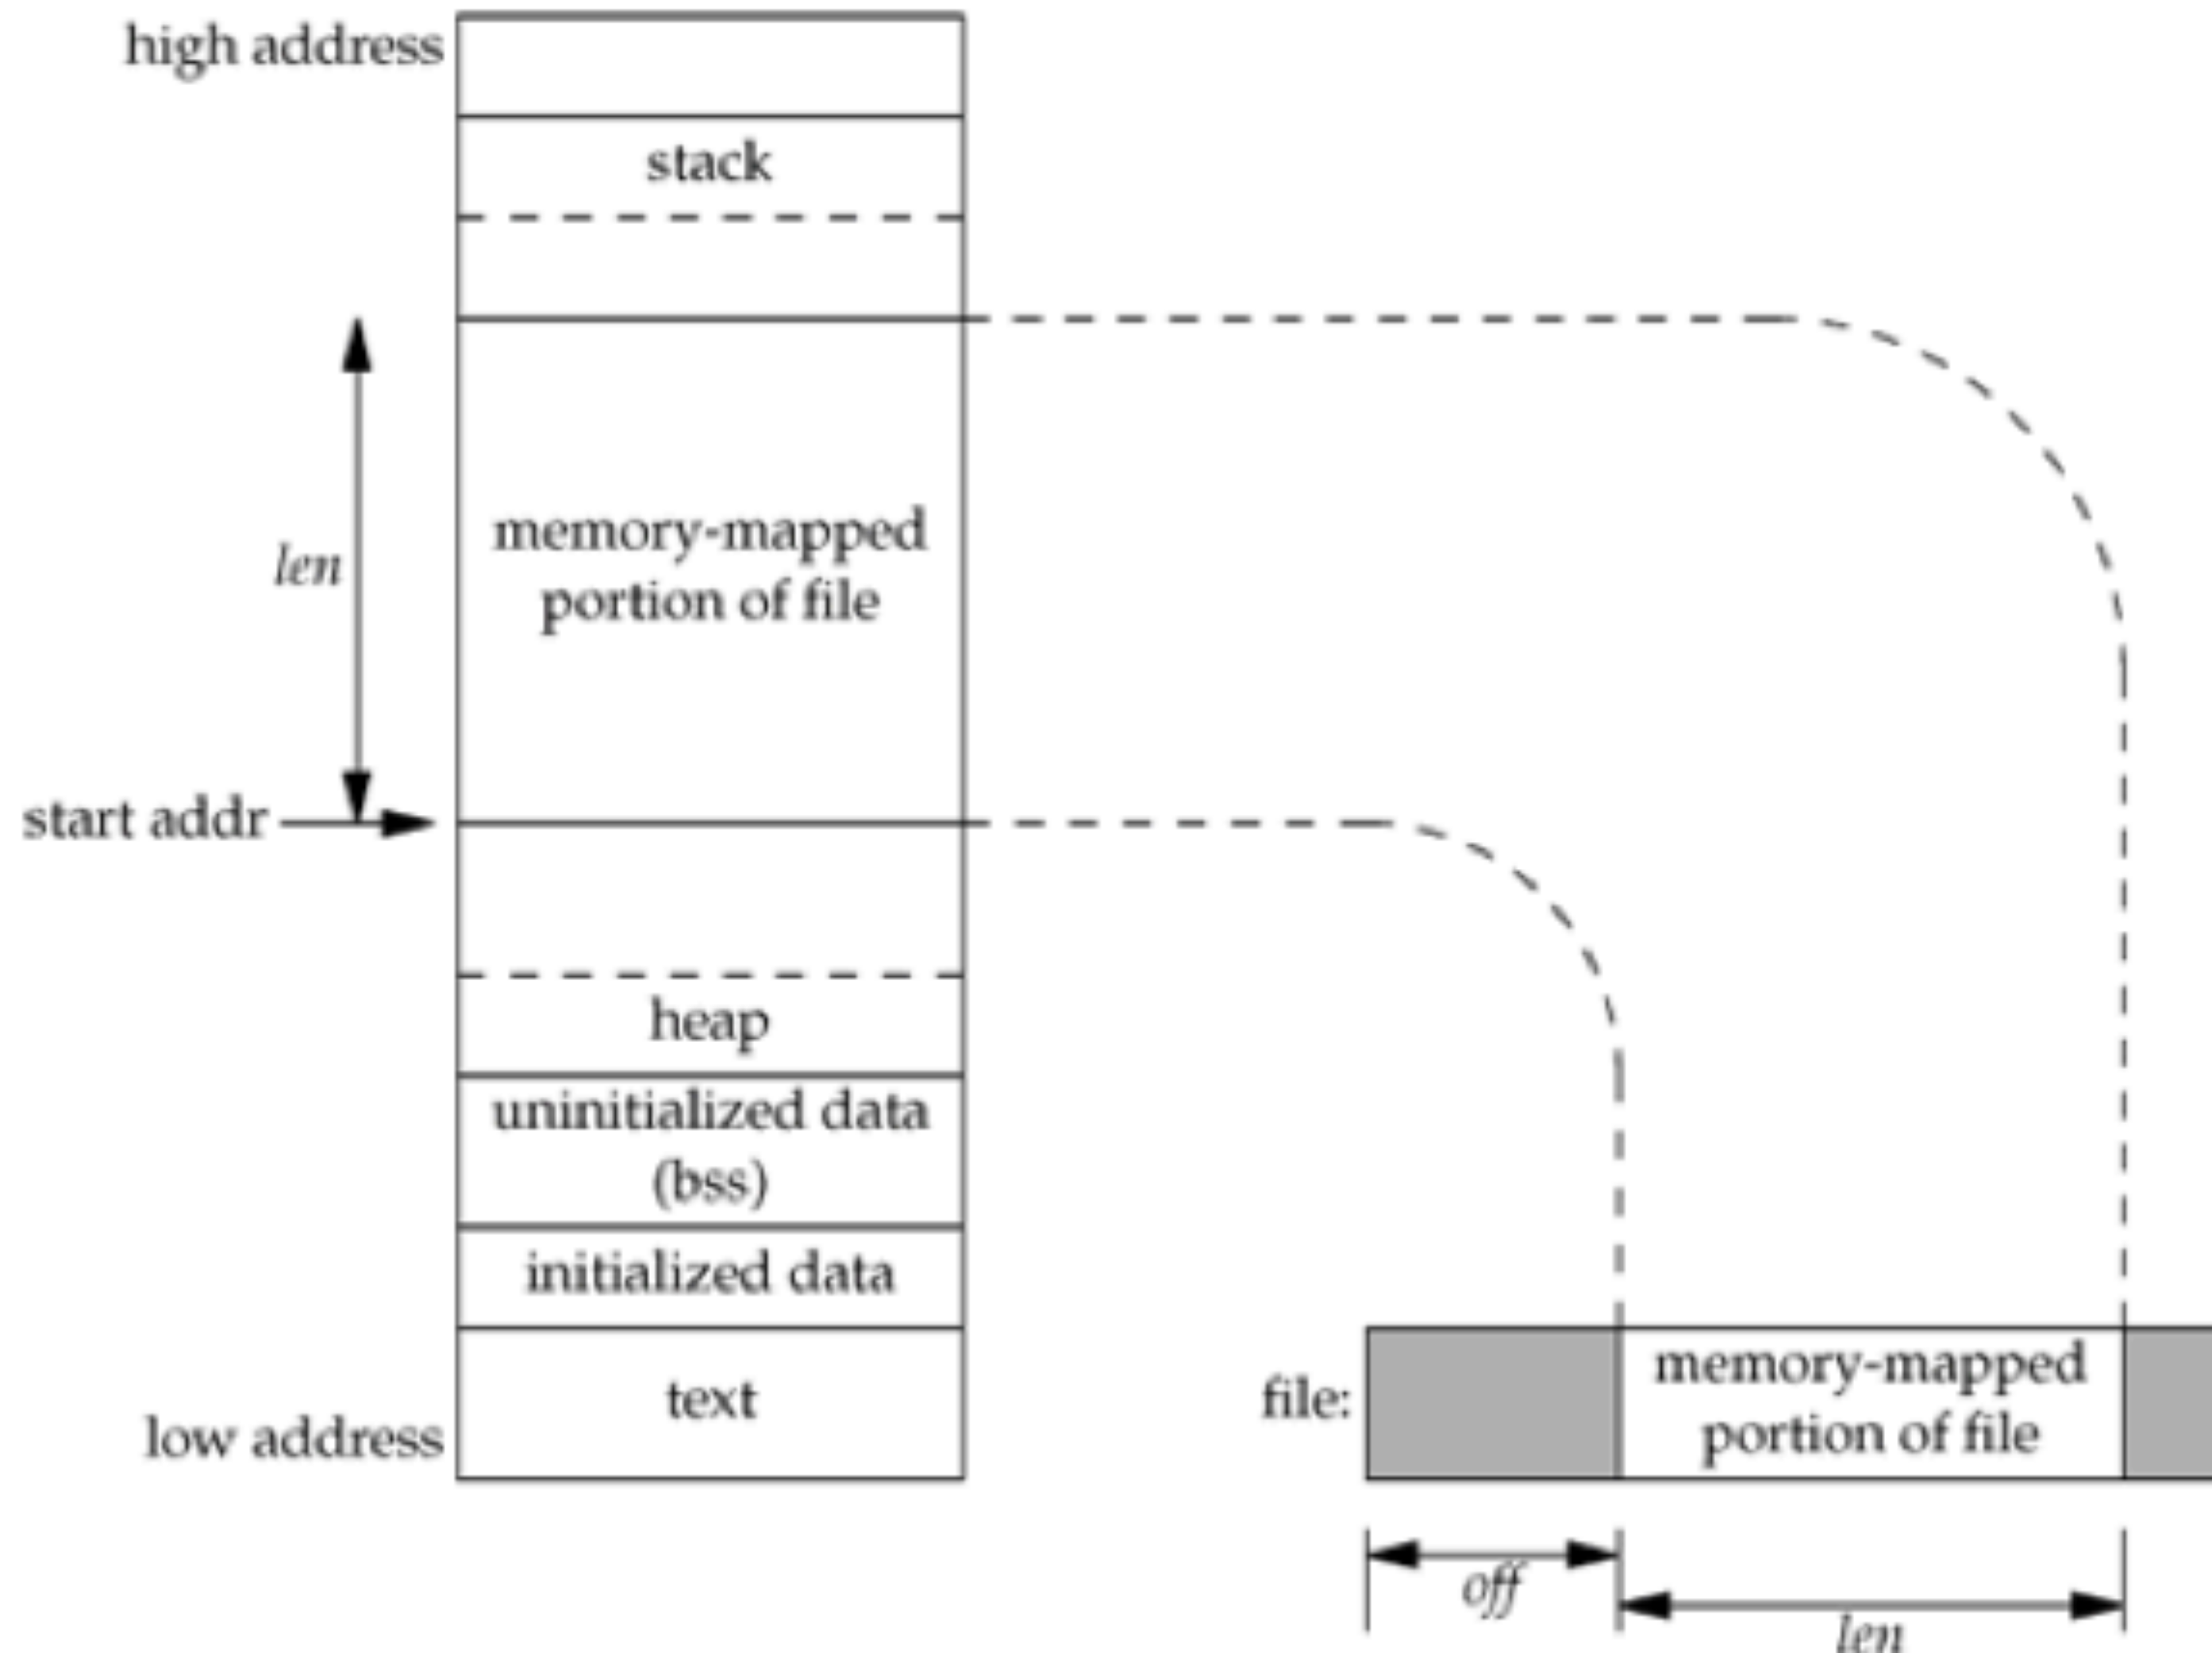
Asynchronous I/O

- semi-async I/O via `select(2)/poll(2)`
- System V derived async I/O
 - limited to STREAMS
 - enabled via `ioctl(2)`
 - uses `SIGPOLL`
- BSD derived async I/O
 - limited to terminals and networks
 - enabled via `open(2)/fcntl(2)` (`O_ASYNC`, `F_SETOWN`)
 - uses `SIGIO` and `SIGURG`

POSIX AIO

- see `aio(7)` on NetBSD
- kernel process manages queued I/O requests
- notification of calling process via signal or sigevent callback function
- calling process can still choose to block/wait
- Linux has multiple implementations:
 - glibc `aio(7)` - <https://is.gd/YZ5fuj>
 - libaio - <https://pagure.io/libaio>

Memory Mapped I/O



mmap(2)

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t len, int prot, int flags, int fd, off_t offset);
```

Returns: pointer to mapped region on success, MAP_FAILED on error

- protection specified for a region:
 - PROT_READ – region can be read
 - PROT_WRITE – region can be written
 - PROT_EXEC – region can be executed
 - PROT_NONE – region can not be accessed
- flag needs to be one of MAP_SHARED or MAP_PRIVATE, which may be OR'd with other flags (see mmap(2) for details).

mmap(2)

Operation	Linux 2.4.22 (Intel x86)			Solaris 9 (SPARC)		
	User	System	Clock	User	System	Clock
read/write	0.04	1.02	39.76	0.18	9.70	41.66
mmap/memcpy	0.64	1.31	24.26	1.68	7.94	28.53

Exercises

- The Linux `aio(7)` manual page includes a code example - can you port this to NetBSD?
- Rewrite your HW1 `cp` to use `mmap(2)/memcpy(2)` instead of `read(2)/write(2)`.
- Benchmark your two implementations on different operating- and file systems.
- Review the NetBSD source code for `cp(1)` - why/when is `mmap(2)` used here?
Why is it not used for all I/O?

<http://cvsweb.netbsd.org/bsdweb.cgi/src/bin/cp/utils.c?rev=HEAD>