

Advanced Programming in the UNIX Environment

Week 13, Segment 3: Restricted Shells, Chroots, and Jails

**Department of Computer Science
Stevens Institute of Technology**

Jan Schaumann

`jschauma@stevens.edu`

`https://stevens.netmeister.org/631/`

Restricted Shells

Another way of restricting what a user can do is to only allow them to execute specific commands, for example via a *restricted* shell:

- prohibit `cd`
- prohibit changing *e.g.*, `PATH` etc.
- prohibit use of commands containing a `'/'` (i.e., only commands found in the (fixed) `PATH` can be executed)
- prohibit redirecting output into files

Beware trivial break-outs via commands that allow invoking other commands!

```
# $NetBSD: dot.shrc,v 1.3 2007/11/24 11:14:42 pavel Exp $

if [ -f /etc/shrc ]; then
    . /etc/shrc
fi

case "$-" in *i*)
    # interactive mode settings go here
    ;;
esac

export PS1="$(whoami)@apue$ "

set -o vi
@apue$ exit
jschauma@apue$ su - fred
fred@apue$ echo $PATH
/home/fred/bin:/bin:/sbin:/usr/bin:/usr/sbin:/usr/X11R7/bin:/usr/pkg/bin:/usr/pkg/sbin:/usr/games:/usr/local/bin:/usr/local/sbin
fred@apue$ sh
fred@apue$ /tmp/.sh -p
fred@apue$ whoami
jschauma
fred@apue$
```

Restricted Shells

To properly restrict a user in this way:

- create a new directory, e.g., `/usr/local/rbin`
- carefully reviewed executables needed, then link them in there
- ensure those commands cannot shell out themselves
- set `PATH=/usr/local/rbin`
- mark user config files immutable via `chflags(1)`
- hope you didn't miss anything

NAME

chroot -- change root directory

SYNOPSIS

```
#include <unistd.h>
```

```
int
```

```
chroot(const char *dirname);
```

DESCRIPTION

Dirname is the address of the pathname of a directory, terminated by an ASCII NUL. **chroot()** causes dirname to become the root directory, that is, the starting point for path searches of pathnames beginning with `/'.

In order for a directory to become the root directory a process must have execute (search) access for that directory.

WARNINGS

There are ways for a root process to escape from the chroot jail.

HISTORY

The **chroot()** function call appeared in 4.2BSD.

4.2 Berkeley Distribution

June 4, 1993

4.2 Berkeley Distribution

(END)

```
uid=0 gid=0 groups=0,2,3,4,5,20,31,34
# cd /
# cd ../../
# cd ../../../../
# echo *
bin lib libexec usr
# cd usr/bin
# echo *
*
# ps
  PID TTY          STAT       TIME COMMAND
 1037 pts/2    0+       0:00.00 ps
 1090 pts/2    I        0:00.01 -csh -c chroot /var/chroot/apue /bin/sh
 1998 pts/2    S        0:00.00 /bin/sh
   607 ?        Is+      0:00.00 /usr/libexec/getty Pc constty
   691 ?        Is+      0:00.00 /usr/libexec/getty Pc ttyE1
   586 ?        Is+      0:00.00 /usr/libexec/getty Pc ttyE2
   532 ?        Is+      0:00.00 /usr/libexec/getty Pc ttyE3
# exit
jschauma@apue$ cd /var/chroot
jschauma@apue$ ls
apue          named          ntpd           rtadvd         tcpdump        unbound
ftp-proxy     nsd            pflogd         sshd           tftp-proxy
jschauma@apue$
```


Chroot

Expose a restricted copy or view of the filesystem to a process via `chroot(2)/chroot(8)`:

- restrict a process's view of the filesystem hierarchy
- restrict commands by only providing needed executables
- must provide full environment, shared libraries, config files, etc.
- combine with null mounts / mount options
- open file descriptors may be brought into the chroot
- processes outside the chroot are visible!

Jails

FreeBSD added the `jail(2)` system call and `jail(8)` utility around 2000. Jails...

- enforce a per-jail process view
- prohibit changing `sysctls` or `securelevels`
- prohibit mounting and unmounting filesystems
- can be bound to a specific network address
- prohibit modifying the network configuration
- disable raw sockets

Jails effectively implement a process sandbox environment, forming the first OS-level virtualization.

Summary

- Restricted shells run fully within the OS, with restrictions entirely enforced within the shell.
- A chroot(2) can create a severely restricted environment with a “changed root” filesystem:
 - present in most Unix versions since the 80s, but since removed from POSIX
 - requires root privileges
 - chroot escapes may be possible
 - process space outside of the chroot remains visible
- Jails were introduced in FreeBSD ~2000 as the first real version of OS-level virtualization and predecessor of true containers.