

# CS631 - Advanced Programming in the UNIX Environment

—

## Process Environment, Process Control

---

Department of Computer Science  
Stevens Institute of Technology  
Jan Schaumann

`jschauma@stevens.edu`

`http://www.cs.stevens.edu/~jschauma/631/`

## Shell Command-Line Processing

---


```
$ cc -Wall argv.c
$ ./a.out
$ ./a.out *.c
$ ./a.out *.none
$ ./a.out *. [1c]
$ ./a.out "*.c"
$ ./a.out $USER
$ ./a.out "$(echo *.1)"
$ ./a.out {foo,bar,baz}.whatever
$ ./a.out {1..5}
$ ./a.out {1..5}{a..f}
```

See also: <https://is.gd/kJXbpa> and <http://is.gd/iZa9rC>

# Memory Layout of a C Program

---


memory-layout.c



# Memory Layout of a C Program

---


memory-layout.c



# Memory Layout of a C Program

---

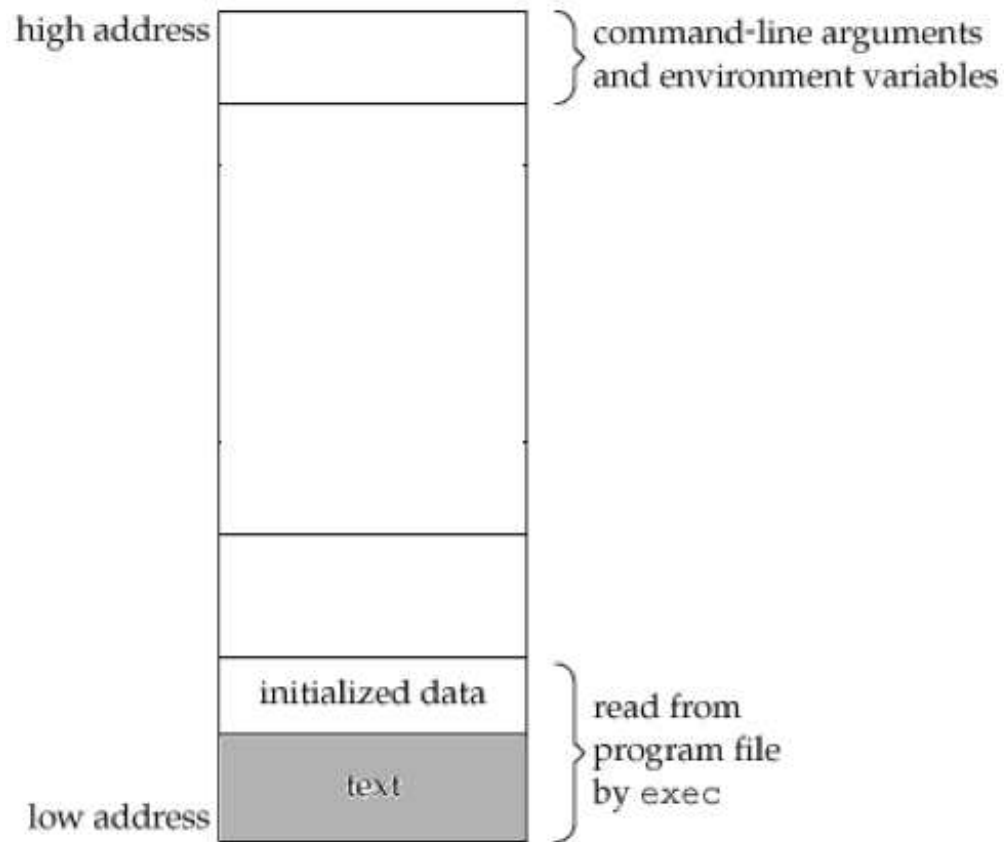
memory-layout.c



# Memory Layout of a C Program


---

memory-layout.c




# Memory Layout of a C Program

memory-layout.c



# Memory Layout of a C Program

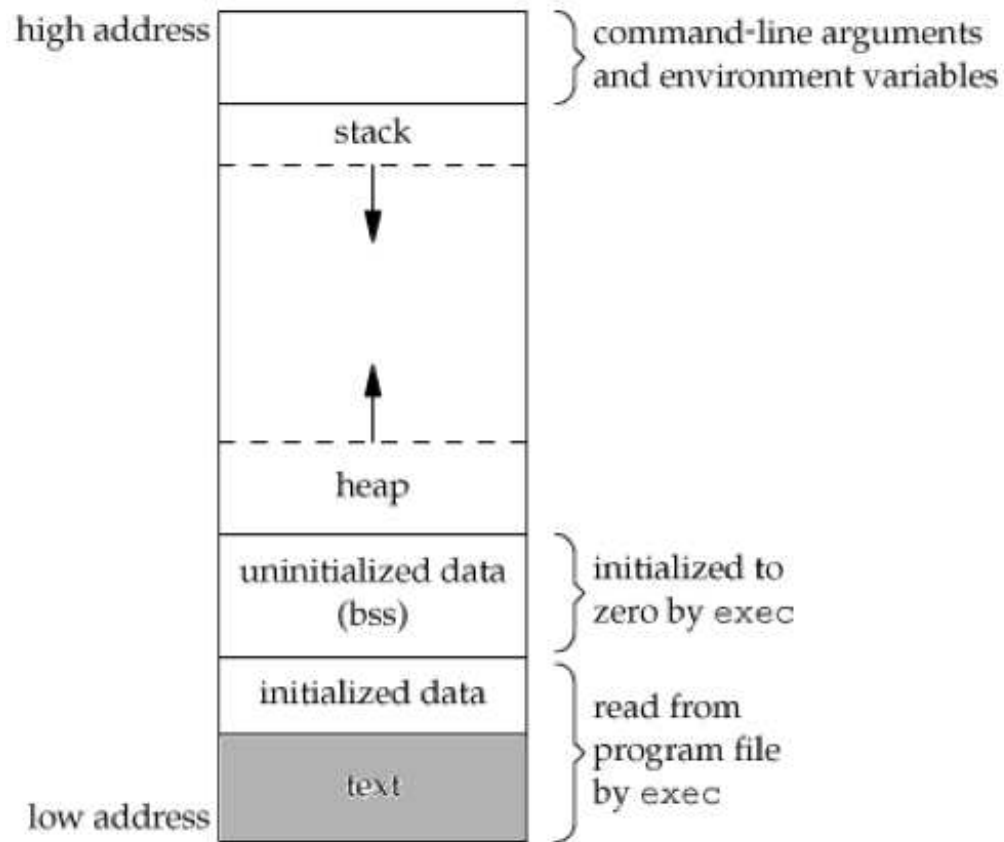
memory-layout.c





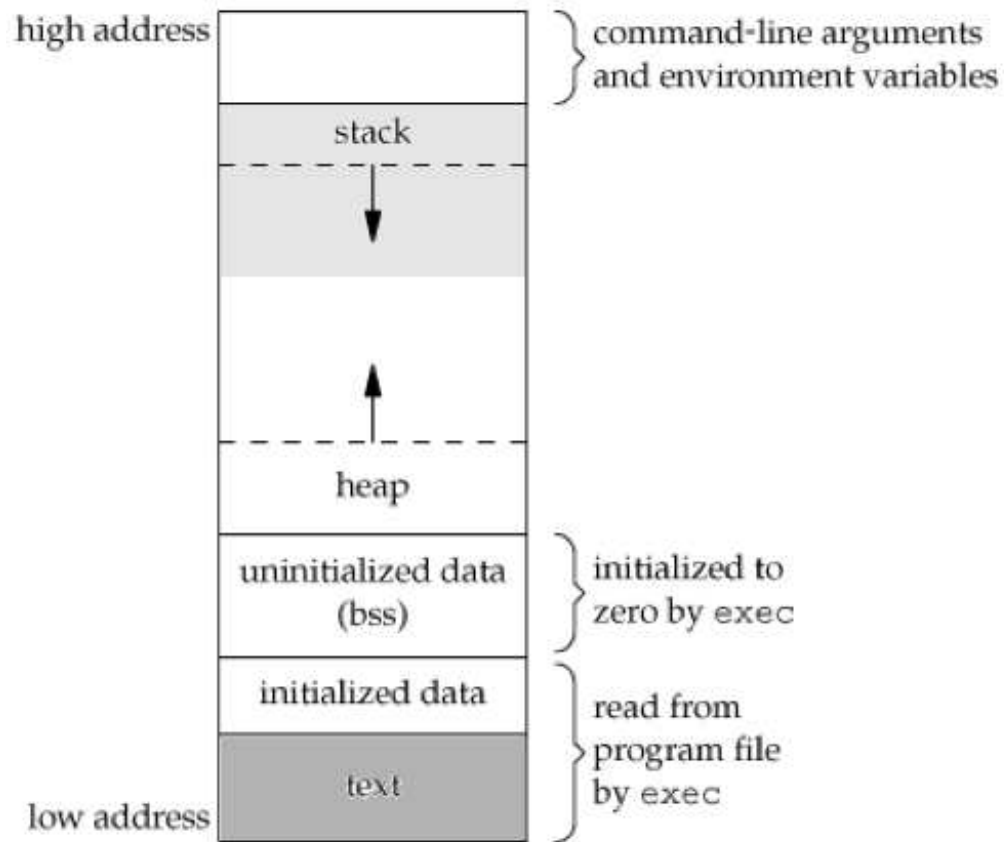
# Memory Layout of a C Program

memory-layout.c



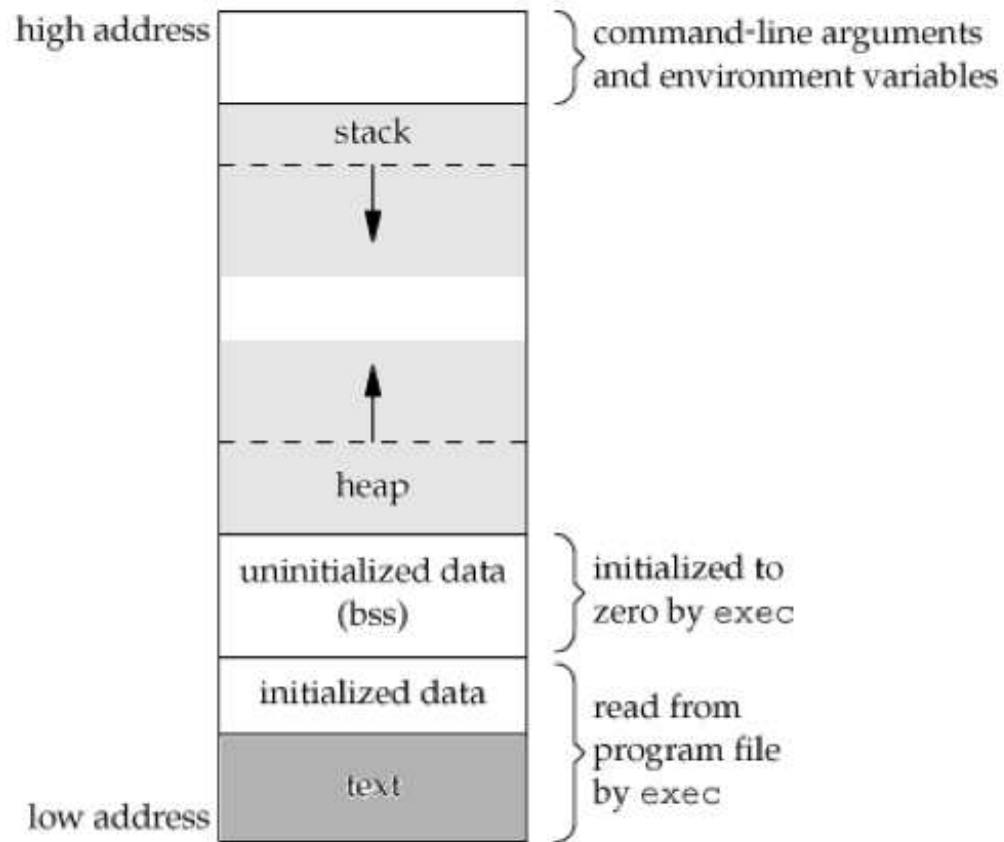
# Memory Layout of a C Program

memory-layout.c



# Memory Layout of a C Program

memory-layout.c



## Memory Layout of a C Program

---

Obligatory "Smashing The Stack For Fun And Profit" links:

<http://insecure.org/stf/smashstack.html>

## The `main` function

---

```
int main(int argc, char **argv);
```

## The `main` function

---

```
int main(int argc, char **argv);
```

- C program started by kernel (by one of the `exec` functions)
- special startup routine called by kernel which sets up things for `main` (or whatever entrypoint is defined)
- `argc` is a count of the number of command line arguments (including the command itself)
- `argv` is an array of pointers to the arguments
- it is guaranteed by both ANSI C and POSIX.1 that `argv[argc] == NULL`

## Process Creation

---

On Linux:

```
$ cc -Wall entry.c
```

```
$ readelf -h a.out | more
```

ELF Header:

```
[...]
```

Entry point address:	0x400460
Start of program headers:	64 (bytes into file)
Start of section headers:	4432 (bytes into file)

```
$ objdump -d a.out
```

```
[...]
```

```
0000000000400460 <_start>:
```

400460:	31 ed	xor	%ebp,%ebp
400462:	49 89 d1	mov	%rdx,%r9

```
[...]
```

```
$
```

## Process Creation

---

glibc/sysdeps/x86\_64/start.S

0000000000401058 <\_start>:

401058:	31 ed	xor	%ebp,%ebp
40105a:	49 89 d1	mov	%rdx,%r9
40105d:	5e	pop	%rsi
40105e:	48 89 e2	mov	%rsp,%rdx
401061:	48 83 e4 f0	and	\$0xfffffffffffffffff0,%rsp
401065:	50	push	%rax
401066:	54	push	%rsp
401067:	49 c7 c0 e0 1a 40 00	mov	\$0x401ae0,%r8
40106e:	48 c7 c1 50 1a 40 00	mov	\$0x401a50,%rcx
401075:	48 c7 c7 91 11 40 00	mov	\$0x401191,%rdi
40107c:	e8 2f 01 00 00	callq	4011b0 <__libc_start_main>
401081:	f4	hlt	
401082:	90	nop	
401083:	90	nop	



## Process Creation

---

glibc/csu/libc-start.c

```
STATIC int
LIBC_START_MAIN (int (*main) (int, char **, char ** MAIN_AUXVEC_DECL),
                 int argc, char **argv,
                 __typeof (main) init,
                 void (*fini) (void),
                 void (*rtld_fini) (void), void *stack_end)
{
    [...]
    result = main (argc, argv, __environ MAIN_AUXVEC_PARAM);

    exit (result);
}
```

## Process Creation

---

On Linux:

```
$ cc -Wall entry.c
```

```
$ readelf -h a.out | more
```

ELF Header:

```
[...]
```

Entry point address:	0x400460
Start of program headers:	64 (bytes into file)
Start of section headers:	4432 (bytes into file)

```
$ objdump -d a.out
```

```
[...]
```

```
0000000000400460 <_start>:
```

400460:	31 ed	xor	%ebp,%ebp
400462:	49 89 d1	mov	%rdx,%r9

```
[...]
```

```
$
```

<http://dbp-consulting.com/tutorials/debugging/linuxProgramStartup.html>

## Process Creation

---

On Linux:

```
$ cc -e foo entry.c
```

```
$ ./a.out
```

```
Foo for the win!
```

```
Memory fault
```

```
$ cc -e bar entry.c
```

```
$ ./a.out
```

```
bar rules!
```

```
$ echo $?
```

```
1
```

```
$ cc entry.c
```

```
$ ./a.out
```

```
Hooray main!
```

```
$ echo $?
```

```
13
```

```
$
```

## Process Termination

---

There are 8 ways for a process to terminate.

Normal termination:

- return from `main`
- calling `exit`
- calling `_exit` (or `_Exit`)
- return of last thread from its start routine
- calling `pthread_exit` from last thread

## Process Termination

---

There are 8 ways for a process to terminate.

Normal termination:

- return from `main`
- calling `exit`
- calling `_exit` (or `_Exit`)
- return of last thread from its start routine
- calling `pthread_exit` from last thread

Abnormal termination:

- calling `abort`
- terminated by a signal
- response of the last thread to a cancellation request

## exit(3) and \_exit(2)

---

```
#include <stdlib.h>
void exit(int status);
void _Exit(int status);

#include <unistd.h>
void _exit(int status);
```

- `_exit` and `_Exit`
  - return to the kernel immediately
  - `_exit` required by POSIX.1
  - `_Exit` required by ISO C99
  - synonymous on Unix
- `exit` does some cleanup and then returns
- both take integer argument, aka *exit status*

## atexit(3)

---

```
#include <stdlib.h>

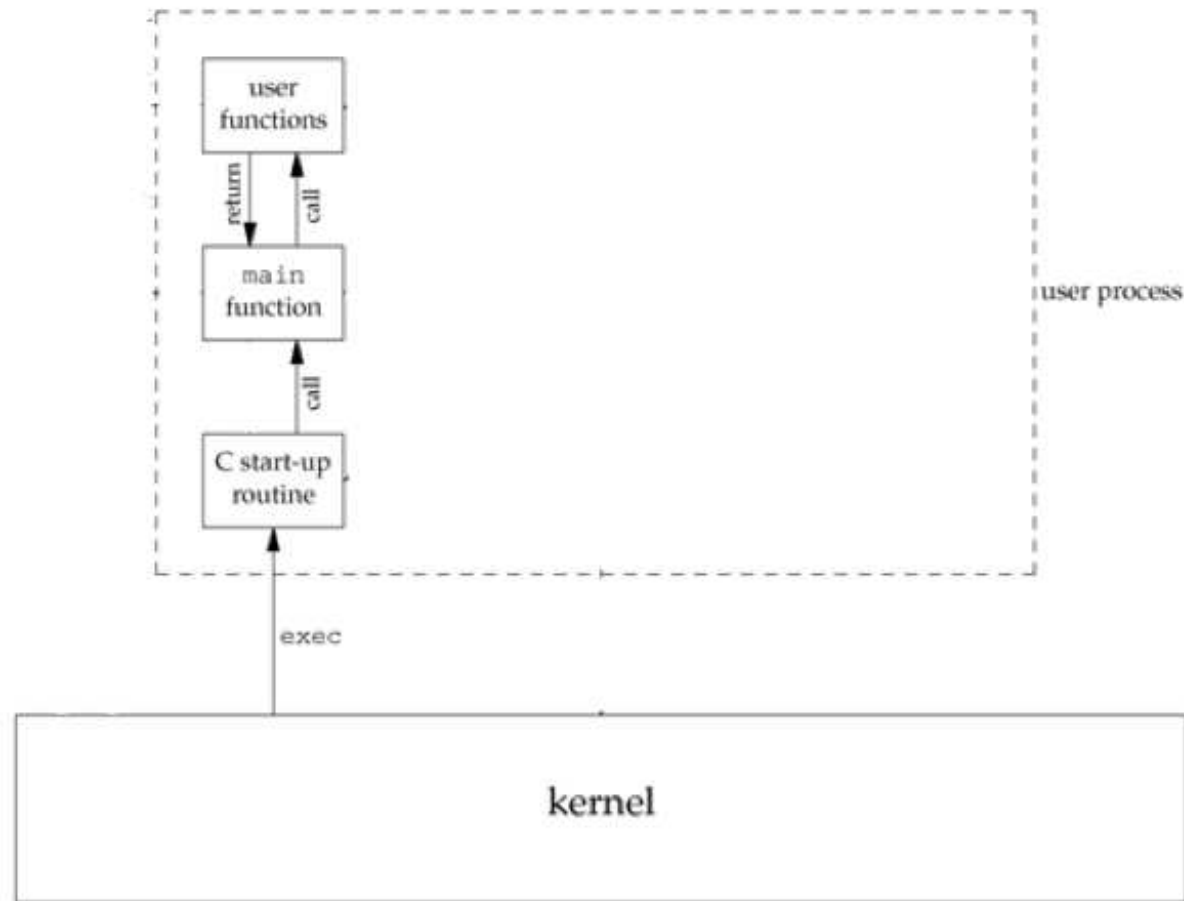
int atexit(void (*func)(void));
```

- Registers a function with a signature of `void funcname(void)` to be called at exit
- Functions invoked in reverse order of registration
- Same function can be registered more than once
- Extremely useful for cleaning up open files, freeing certain resources, etc.

`exit-handlers.c`

## Lifetime of a UNIX Process

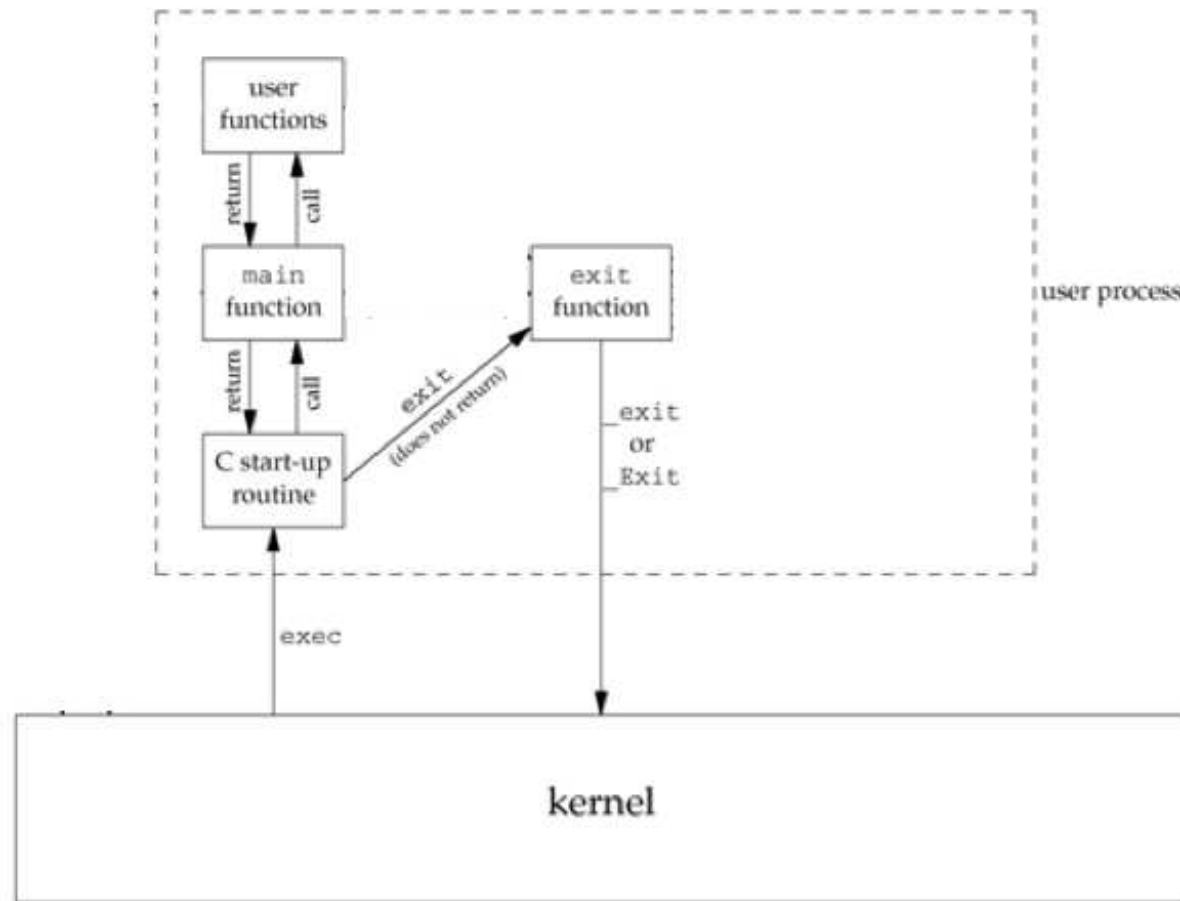
---






## Lifetime of a UNIX Process


---



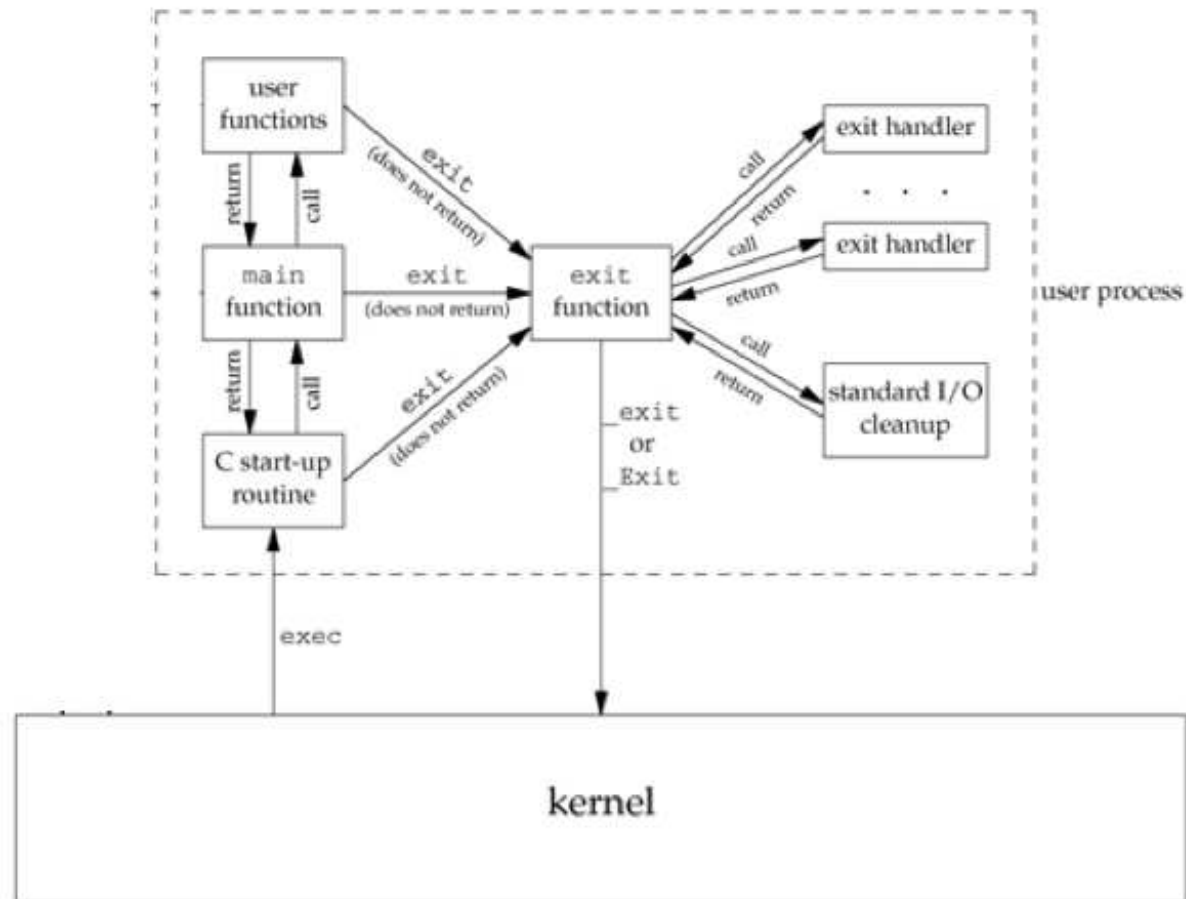
## Lifetime of a UNIX Process




## Lifetime of a UNIX Process



## Lifetime of a UNIX Process



## Lifetime of a UNIX Process



## Exit codes

---

```
$ cc -Wall hw.c
hw.c: In function 'main':
hw.c:7: warning: control reaches end of non-void function
$ ./a.out
Hello World!
$ echo $?
13
$
```

## Exit codes

---

```
$ cc -Wall hw.c
hw.c: In function 'main':
hw.c:7: warning: control reaches end of non-void function
$ ./a.out
Hello World!
$ echo $?
13
$ cc -Wall --std=c99 hw.c
$ ./a.out
Hello World!
$ echo $?
0
$
```

## Environment List

---

Environment variables are stored in a global array of pointers:

```
extern char **environ;
```

The list is `null` terminated.

These can also be accessed by:

```
#include <stdlib.h>

char *getenv(const char *name);
int putenv(const char *string);
int setenv(const char *name, const char *value, int rewrite);
void unsetenv(const char *name);
```



## Environment List

---

Environment variables are stored in a global array of pointers:

```
extern char **environ;
```

The list is `null` terminated.

These can also be accessed by:

```
#include <stdlib.h>

char *getenv(const char *name);
int putenv(const char *string);
int setenv(const char *name, const char *value, int rewrite);
void unsetenv(const char *name);
```

```
int main(int argc, char **argv, char **envp);
```

## Memory Allocation

---

```
#include <stdlib.h>

void *malloc(size_t size);
void *calloc(size_t nobj, size_t size);
void *realloc(void *ptr, size_t newsz);
void *alloca(size_t size);


void free(void *ptr);
```

- *malloc* – initial value is indeterminate.
- *calloc* – initial value set to all zeros.
- *realloc* – changes size of previously allocated area. Initial value of any additional space is indeterminate.
- *alloca* – allocates memory on stack

Now consider manipulation of the environment by your program...

# Memory Layout of a C Program

---



## Memory Layout of a C Program

---

On NetBSD:

```
$ cc hw.c
$ file a.out
a.out: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically
linked (uses shared libs), for NetBSD 5.0, not stripped
$ ldd a.out
a.out:
        -lc.12 => /usr/lib/libc.so.12
$ size a.out
   text    data     bss     dec     hex filename
   2301     552     120     2973     b9d a.out
$ objdump -d a.out > obj
$ wc -l obj
    271 obj
$
```

## Memory Layout of a C Program

---

On Mac OS X:

```
$ cc hw.c
$ file a.out
a.out: Mach-O 64-bit executable x86_64
$ otool -L a.out
a.out:
/usr/lib/libSystem.B.dylib (compatibility version 1.0.0,
current version 125.2.11)
$ size a.out
__TEXT __DATA __OBJC others dec hex
4096 4096 0 4294971392 4294979584 100003000
$ otool -t -v a.out > obj
$ wc -l obj
    32 obj
$
```

## Memory Layout of a C Program

---

On Linux:

```
$ cc hw.c
$ file a.out
a.out: ELF 32-bit LSB executable, Intel 80386, version 1 (SYSV),
dynamically linked (uses shared libs), for GNU/Linux 2.6.15, not stripped
$ ldd a.out
linux-gate.so.1 => (0x00c66000)
libc.so.6 => /lib/tls/i686/cmov/libc.so.6 (0x006b4000)
/lib/ld-linux.so.2 (0x005fe000)
$ size a.out
   text    data     bss     dec     hex filename
   918     264        8    1190    4a6 a.out
$ objdump -d a.out >obj
$ wc -l obj
225 obj
$
```

## Memory Layout of a C Program

---

On NetBSD:

```
$ cc -static hw.c
```

```
$ file a.out
```

```
a.out: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), statically  
linked, for NetBSD 5.0, not stripped
```

```
$ ldd a.out
```

```
ldd: a.out: unrecognized file format2 [2 != 1]
```

```
$ size a.out
```

text	data	bss	dec	hex	filename
151877	4416	16384	172677	2a285	a.out

```
$ size a.out.dyn
```

text	data	bss	dec	hex	filename
2301	552	120	2973	b9d	a.out

```
$ objdump -d a.out > obj
```

```
$ wc -l obj
```

```
35029 obj
```

```
$
```

## Memory Layout of a C Program

---

On Mac OS X:

```
$ cc -static hw.c  
ld: library not found for -lcrt0.o  
collect2: ld returned 1 exit status  
$
```



## Memory Layout of a C Program

---

On Linux:

```
$ cc -static hw.c
```

```
$ file a.out
```

```
a.out: ELF 32-bit LSB executable, Intel 80386, version 1 (SYSV),  
statically linked, for GNU/Linux 2.6.15, not stripped
```

```
$ ldd a.out
```

```
/usr/bin/ldd: line 161: /lib64/ld-linux-x86-64.so.2: cannot execute binary file  
not a dynamic executable
```

```
$ size a.out
```

text	data	bss	dec	hex	filename
510786	1928	7052	519766	7ee56	a.out

```
$ objdump -d a.out >obj
```

```
$ wc -l obj
```

```
114420 obj
```

```
$
```

## Process limits

---

```
$ ulimit -a
time(cpu-seconds)      unlimited
file(blocks)           unlimited
coredump(blocks)       unlimited
data(kbytes)           262144
stack(kbytes)          2048
lockedmem(kbytes)      249913
memory(kbytes)         749740
nofiles(descriptors)   128
processes              160
vmemory(kbytes)        unlimited
sbsize(bytes)          unlimited
$
```

## getrlimit(2) and setrlimit(2)

---

```
#include <sys/resource.h>

int getrlimit(int resource, struct rlimit *rlp);
int setrlimit(int resource, const struct rlimit *rlp);
```

Changing resource limits follows these rules:

- a *soft limit* can be changed by any process to a value less than or equal to its hard limit
- any process can lower its *hard limit* greater than or equal to its soft limit
- only superuser can raise *hard limits*
- changes are per process only

## getrlimit(2) and setrlimit(2)

---

```
#include <sys/resource.h>

int getrlimit(int resource, struct rlimit *rlp);
int setrlimit(int resource, const struct rlimit *rlp);
```

Changing resource limits follows these rules:

- a *soft limit* can be changed by any process to a value less than or equal to its hard limit
- any process can lower its *hard limit* greater than or equal to its soft limit
- only superuser can raise *hard limits*
- changes are per process only (which is why `ulimit` is a shell built-in)

# Process Control

## Review from our first class, the world's simplest shell:

```
int
main(int argc, char **argv)
{
    char buf[1024];
    pid_t pid;
    int status;

    while (getinput(buf, sizeof(buf))) {
        buf[strlen(buf) - 1] = '\0';

        if((pid=fork()) == -1) {
            fprintf(stderr, "shell: can't fork: %s\n",
                    strerror(errno));
            continue;
        } else if (pid == 0) {
            /* child */
            execlp(buf, buf, (char *)0);
            fprintf(stderr, "shell: couldn't exec %s: %s\n", buf,
                    strerror(errno));
            exit(EX_DATAERR);
        }

        if ((pid=waitpid(pid, &status, 0)) < 0)
            fprintf(stderr, "shell: waitpid error: %s\n",
                    strerror(errno));
    }

    exit(EX_OK);
}
```

## Process Identifiers

---

```
#include <unistd.h>

pid_t getpid(void);
pid_t getppid(void);
```

*Process ID's* are guaranteed to be unique and identify a particular executing process with a non-negative integer.

Certain processes have fixed, special identifiers. They are:

- *swapper*, process ID 0 – responsible for scheduling
- *init*, process ID 1 – bootstraps a Unix system, owns orphaned processes
- *pagedaemon*, process ID 2 – responsible for the VM system (some Unix systems)

## fork(2)

---

```
#include <unistd.h>


pid_t fork(void);
```

fork(2) causes creation of a new process. The new process (child process) is an exact copy of the calling process (parent process) except for the following:

- The child process has a unique process ID.
- The child process has a different parent process ID (i.e., the process ID of the parent process).
- The child process has its own copy of the parent's descriptors.
- The child process' resource utilizations are set to 0.


**Note:** no order of execution between child and parent is guaranteed!

## fork(2)





## fork(2)



## fork(2)

---

```
$ cc -Wall forkflush.c
$ ./a.out
a write to stdout
before fork
pid = 12149, glob = 7, var = 89
pid = 12148, glob = 6, var = 88
$ ./a.out | cat
a write to stdout
before fork
pid = 12153, glob = 7, var = 89
before fork
pid = 12151, glob = 6, var = 88
$
```

## The `exec(3)` functions

---

```
#include <unistd.h>

int execl(const char *pathname, const char *arg0, ... /* (char *) 0 */);
int execv(const char *pathname, char * const argv[]);
int execlp(const char *pathname, const char *arg0, ... /* (char *) 0, char *const envp[] */ );
int execve(const char *pathname, char * const argv[], char * const envp[]);
int execlp(const char *filename, const char *arg0, ... /* (char *) 0 */);
int execvp(const char *filename, char *const argv[]);
```

The `exec()` family of functions are used to completely replace a running process with a new executable.

- if it has a `v` in its name, `argv`'s are a vector: `const * char argv[]`
- if it has an `l` in its name, `argv`'s are a list: `const char *arg0, ... /* (char *) 0 */`
- if it has an `e` in its name, it takes a `char * const envp[]` array of environment variables
- if it has a `p` in its name, it uses the `PATH` environment variable to search for the file

## wait(2) and waitpid(2)

---

```
#include <sys/types.h>
#include <sys/wait.h>

pid_t wait(int *status);
pid_t waitpid(pid_t wpid, int *status, int options);
pid_t wait3(int *status, int options, struct rusage *rusage);
pid_t wait4(pid_t wpid, int *status, int options, struct rusage *rusage);
```

A parent that calls wait(2) or waitpid(2) can:

- block (if all of its children are still running)
- return immediately with the termination status of a child
- return immediately with an error

## wait(2) and waitpid(2)

---

Differences between `wait(2)`, `wait3(2)`, `wait4(2)` and `waitpid(2)`:

- `wait(2)` will block until the process terminates, `waitpid(2)` has an option to prevent it from blocking
- `waitpid(2)` can wait for a specific process to finish
- `wait3(2)` and `wait4(2)` allow you to get detailed resource utilization statistics
- `wait3(2)` is the same as `wait4(2)` with a *wpid* value of `-1`

## wait(2) and waitpid(2)

---

Once we get a termination status back in `status`, we'd like to be able to determine how a child died. We do this with the following macros:

- `WIFEXITED(status)` – true if the child terminated normally. Then execute `WEXITSTATUS(status)` to get the exit status.
- `WIFSIGNALED(status)` – true if child terminated abnormally (by receiving a signal it didn't catch). Then we call:
  - `WTERMSIG(status)` to retrieve the signal number
  - `WCOREDUMP(status)` to see if the child left a core image
- `WIFSTOPPED(status)` – true if the child is currently stopped. Call `WSTOPSIG(status)` to determine the signal that caused this.

Additionally, `waitpid`'s behavior can be modified by supplying `WNOHANG` as an option, which says that if the requested pid has not terminated, return immediately instead of blocking.

## What if we don't wait(2)?

---

## What if we don't wait (2)?

---





## What if we don't wait(2)?

---

```
$ cc -Wall zombies.c
$ ./a.out
Let's create some zombies!
====
15603 s003 S+      0:00.00 ./a.out
15604 s003 Z+      0:00.00 (a.out)
====
15603 s003 S+      0:00.00 ./a.out
15604 s003 Z+      0:00.00 (a.out)
15608 s003 Z+      0:00.00 (a.out)
====
15603 s003 S+      0:00.00 ./a.out
15604 s003 Z+      0:00.00 (a.out)
15612 s003 Z+      0:00.00 (a.out)
====
15603 s003 S+      0:00.00 ./a.out
15604 s003 Z+      0:00.00 (a.out)
15612 s003 Z+      0:00.00 (a.out)
15616 s003 Z+      0:00.00 (a.out)
```

## Notes and Homework

---

Reading:

- Stevens, Chapter 7 and 8

Other:

- work on your midterm project!